

Lecture 11: Bayes Nets: Inference

Shuai Li

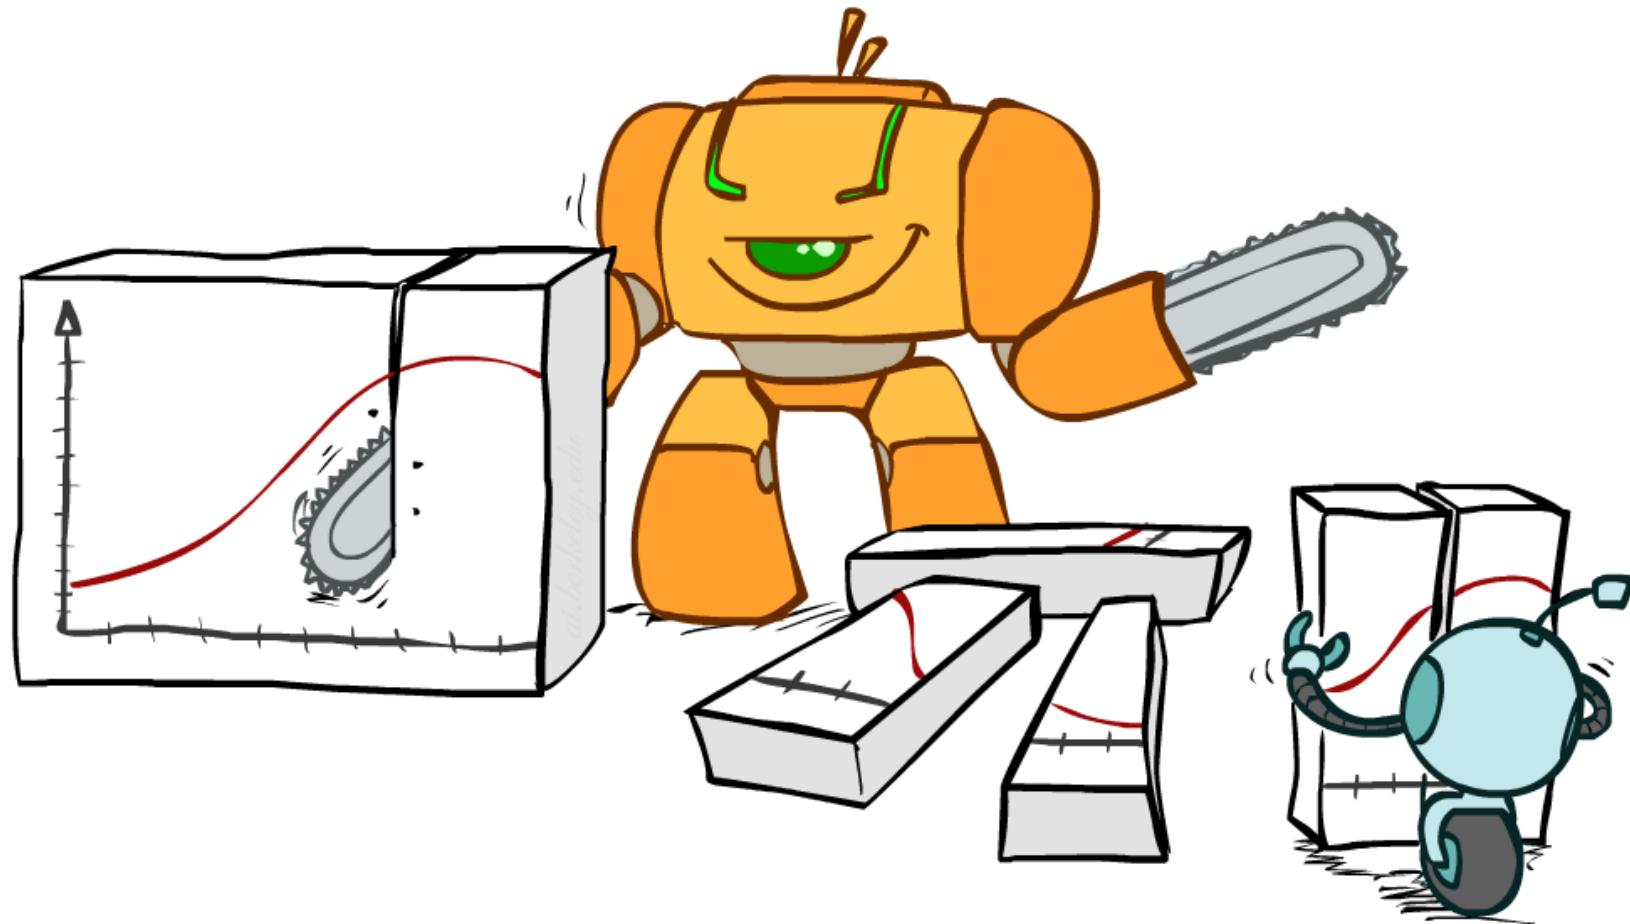
John Hopcroft Center, Shanghai Jiao Tong University

<https://shuaili8.github.io>

<https://shuaili8.github.io/Teaching/CS3317/index.html>

Part of slide credits: CMU AI & <http://ai.berkeley.edu>

Bayes Rule



Bayes' Rule

- Two ways to factor a joint distribution over two variables:

$$P(x, y) = P(x|y)P(y) = P(y|x)P(x)$$

- Dividing, we get:

$$P(x|y) = \frac{P(y|x)}{P(y)}P(x)$$

- Why is this at all helpful?

- Lets us build one conditional from its reverse
- Often one conditional is tricky but the other one is simple
- Foundation of many systems (e.g. ASR, MT)

- In the running for most important AI equation!



Inference with Bayes' Rule

- Example: Diagnostic probability from causal probability:

$$P(\text{cause}|\text{effect}) = \frac{P(\text{effect}|\text{cause})P(\text{cause})}{P(\text{effect})}$$

- Example:

- M: meningitis, S: stiff neck

$$\left. \begin{array}{l} P(+m) = 0.0001 \\ P(+s|m) = 0.8 \\ P(+s|-m) = 0.01 \end{array} \right\} \text{Example givens}$$

$$P(+m|s) = \frac{P(+s|m)P(+m)}{P(+s)} = \frac{P(+s|m)P(+m)}{P(+s|m)P(+m) + P(+s|-m)P(-m)} = \frac{0.8 \times 0.0001}{0.8 \times 0.0001 + 0.01 \times 0.999}$$

- Note: posterior probability of meningitis still very small
 - Note: you should still get stiff necks checked out! Why?

Quiz: Bayes' Rule

- Given: $P(W)$ $P(D|W)$

R	P
sun	0.8
rain	0.2

D	W	P
wet	sun	0.1
dry	sun	0.9
wet	rain	0.7
dry	rain	0.3

- What is $P(W | \text{dry})$?

Quiz: Bayes' Rule 2

- Given:

R	P
sun	0.8
rain	0.2

D	W	P
wet	sun	0.1
dry	sun	0.9
wet	rain	0.7
dry	rain	0.3
- What is $P(W | \text{dry})$?

$$P(\text{sun} | \text{dry}) \propto P(\text{dry} | \text{sun})P(\text{sun}) = .9 * .8 = .72$$

$$P(\text{rain} | \text{dry}) \propto P(\text{dry} | \text{rain})P(\text{rain}) = .3 * .2 = .06$$

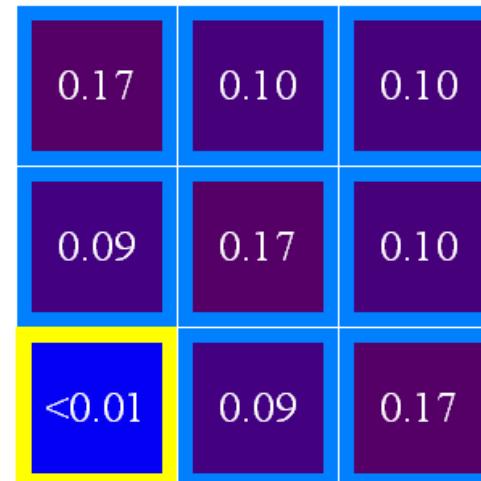
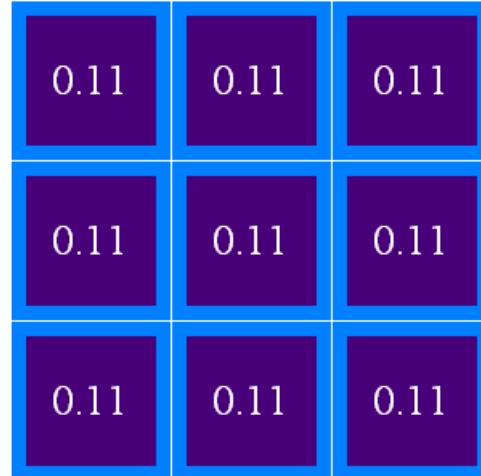
$$P(\text{sun} | \text{dry}) = 12/13$$

$$P(\text{rain} | \text{dry}) = 1/13$$

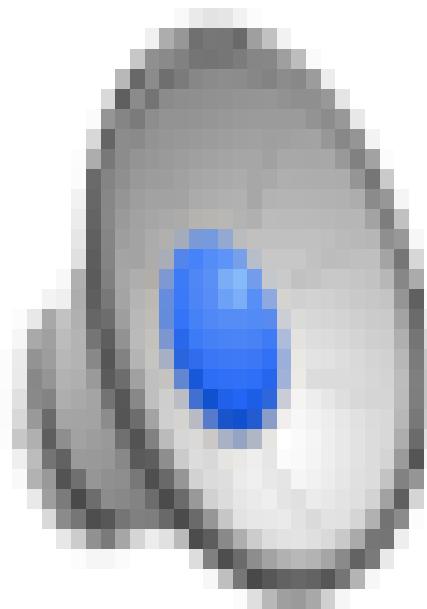
Ghostbusters, Revisited

- Let's say we have two distributions:
 - Prior distribution over ghost location: $P(G)$
 - Let's say this is uniform
 - Sensor reading model: $P(R | G)$
 - Given: we know what our sensors do
 - R = reading color measured at (1,1)
 - E.g. $P(R = \text{yellow} | G=(1,1)) = 0.1$
- We can calculate the posterior distribution $P(G|r)$ over ghost locations given a reading using Bayes' rule:

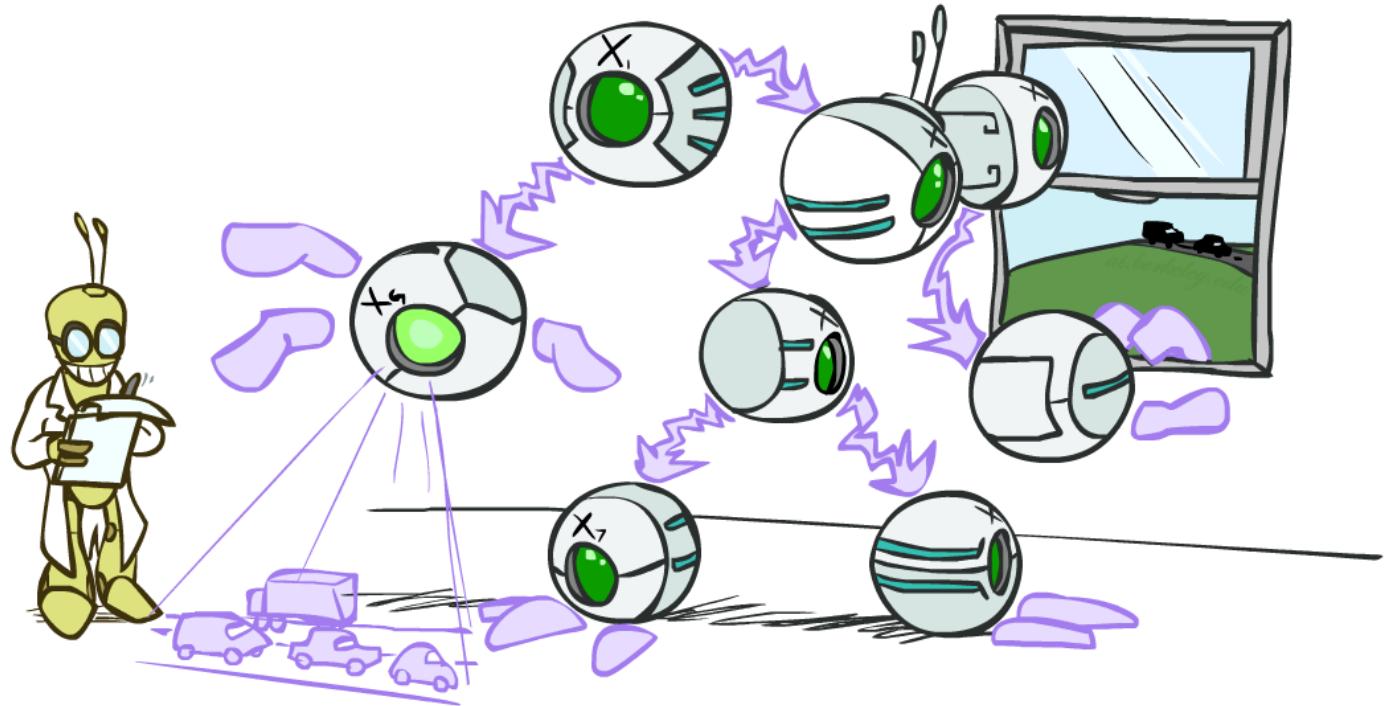
$$P(g|r) \propto P(r|g)P(g)$$



Video of Demo Ghostbusters with Probability



Inference



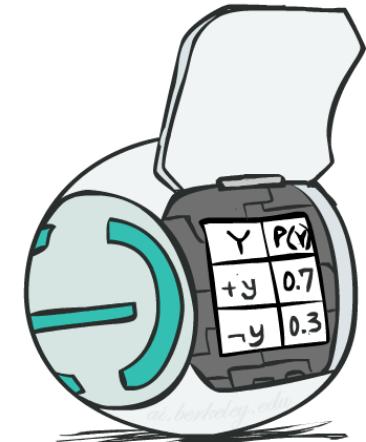
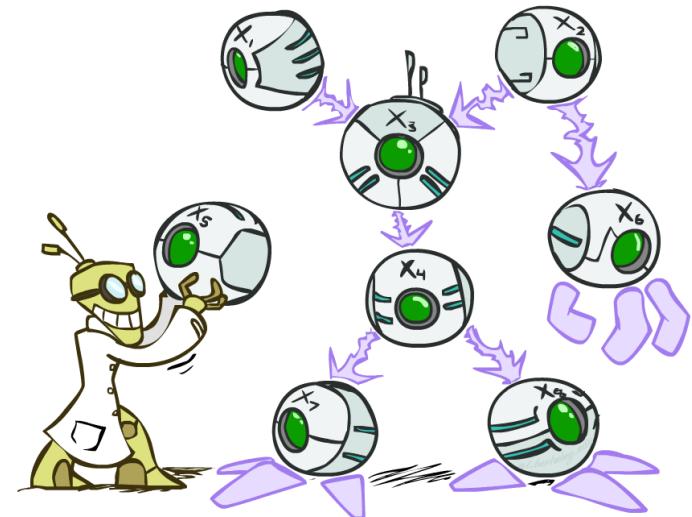
Recall: Bayes' Net Representation

- A directed, acyclic graph, one node per random variable
- A conditional probability table (CPT) for each node
 - A collection of distributions over X, one for each combination of parents' values

$$P(X|a_1 \dots a_n)$$

- Bayes' nets implicitly encode joint distributions
 - As a product of local conditional distributions
 - To see what probability a BN gives to a full assignment, multiply all the relevant conditionals together:

$$P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i | \text{parents}(X_i))$$



Inference

- Inference: calculating some useful quantity from a joint probability distribution

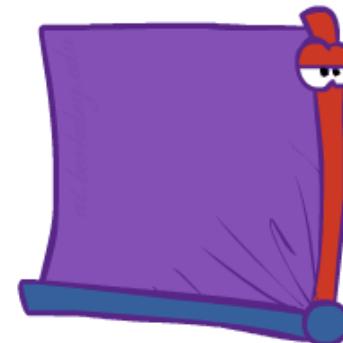
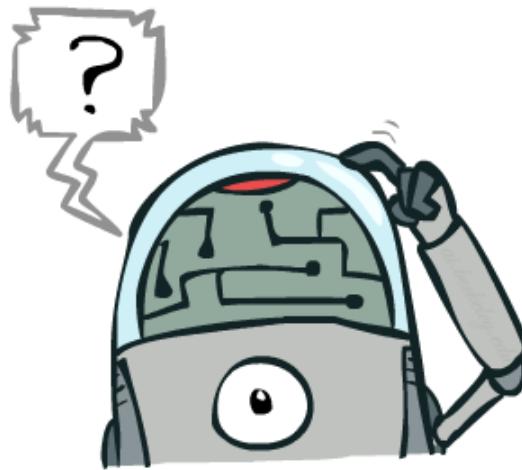
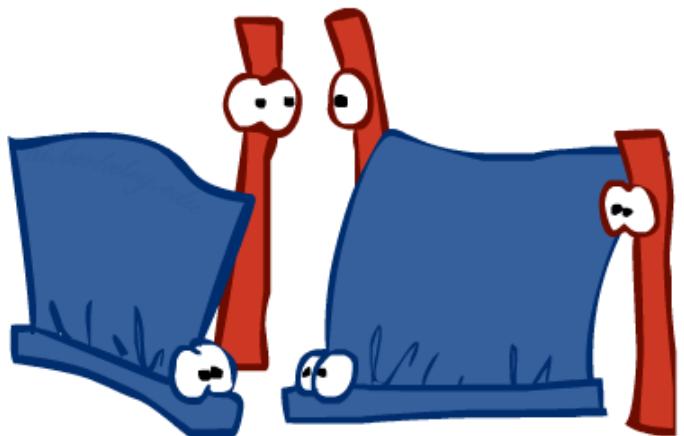
- Examples:

- Posterior probability

$$P(Q|E_1 = e_1, \dots, E_k = e_k)$$

- Most likely explanation:

$$\operatorname{argmax}_q P(Q = q|E_1 = e_1 \dots)$$



Queries

- What is the probability of *this* given what I know?

$$P(q | e) = \frac{P(q, e)}{P(e)} = \frac{\sum_{h_1} \sum_{h_2} P(q, h_1, h_2, e)}{P(e)}$$

- What are the probabilities of all the possible outcomes (given what I know)?

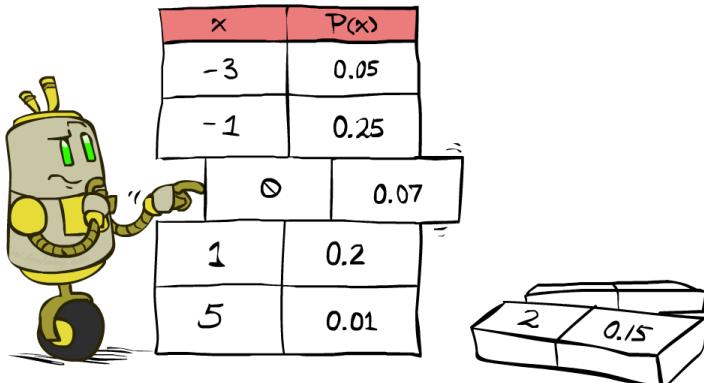
$$P(Q | e) = \frac{P(Q, e)}{P(e)} = \frac{\sum_{h_1} \sum_{h_2} P(Q, h_1, h_2, e)}{P(e)}$$

- Which outcome is the most likely outcome (given what I know)?

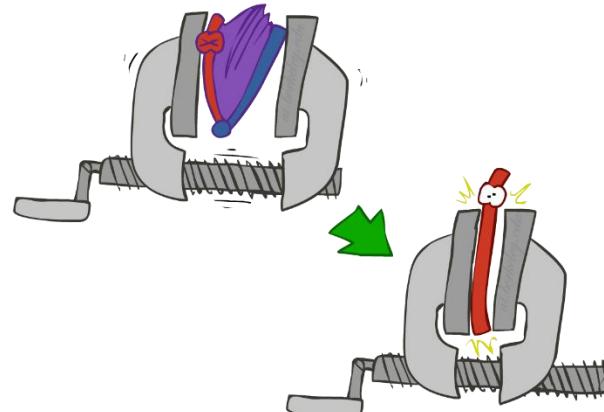
$$\begin{aligned} \operatorname{argmax}_{q \in Q} P(q | e) &= \operatorname{argmax}_{q \in Q} \frac{P(q, e)}{P(e)} \\ &= \operatorname{argmax}_{q \in Q} \frac{\sum_{h_1} \sum_{h_2} P(q, h_1, h_2, e)}{P(e)} \end{aligned}$$

Inference by Enumeration in Joint Distributions

- General case:
 - Evidence variables: $E_1 \dots E_k = e_1 \dots e_k$
 - Query* variable: Q
 - Hidden variables: $H_1 \dots H_r$
- Step 1: Select the entries consistent with the evidence



- We want: $P(Q|e_1 \dots e_k)$
- Step 2: Sum out H to get joint of Query and evidence
- Step 3: Normalize



$$P(Q, e_1 \dots e_k) = \sum_{h_1 \dots h_r} \underbrace{P(Q, h_1 \dots h_r, e_1 \dots e_k)}_{X_1, X_2, \dots, X_n}$$

$$Z = \sum_q P(Q, e_1 \dots e_k)$$

$$P(Q|e_1 \dots e_k) = \frac{1}{Z} P(Q, e_1 \dots e_k)$$

* Works fine with multiple query variables, too

$$\times \frac{1}{Z}$$

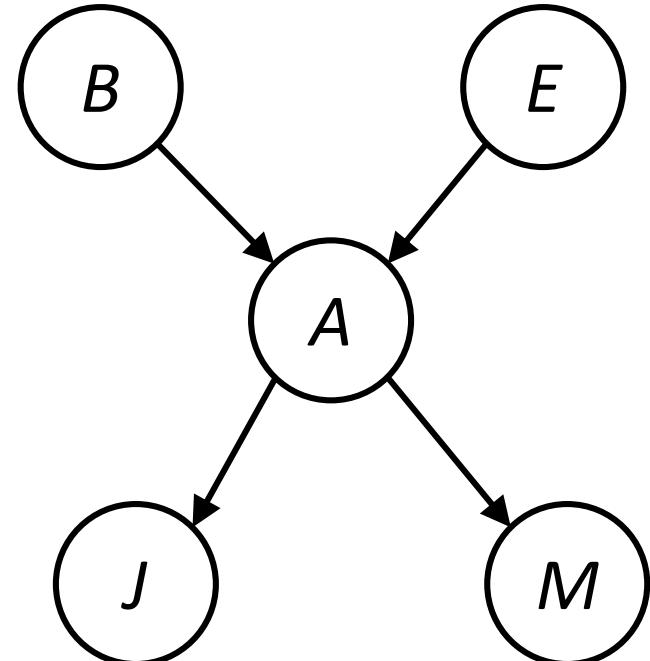
Inference by Enumeration in Bayes' Net

- Given unlimited time, inference in BNs is easy

$$P(B \mid +j, +m) \propto_B P(B, +j, +m)$$

$$= \sum_{e,a} P(B, e, a, +j, +m)$$

$$= \sum_{e,a} P(B)P(e)P(a|B, e)P(+j|a)P(+m|a)$$



$$= P(B)P(+e)P(+a|B, +e)P(+j|+a)P(+m|+a) + P(B)P(+e)P(-a|B, +e)P(+j|-a)P(+m|-a)$$

$$P(B)P(-e)P(+a|B, -e)P(+j|+a)P(+m|+a) + P(B)P(-e)P(-a|B, -e)P(+j|-a)P(+m|-a)$$

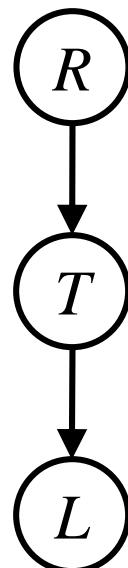
Example: Traffic Domain

- Random Variables
 - R: Raining
 - T: Traffic
 - L: Late for class!

$$P(L) = ?$$

$$= \sum_{r,t} P(r,t,L)$$

$$= \sum_{r,t} P(r)P(t|r)P(L|t)$$



$$P(R)$$

+r	0.1
-r	0.9

$$P(T|R)$$

+r	+t	0.8
+r	-t	0.2
-r	+t	0.1
-r	-t	0.9

$$P(L|T)$$

+t	+l	0.3
+t	-l	0.7
-t	+l	0.1
-t	-l	0.9

Inference by Enumeration: Procedural Outline

- Track objects called **factors**
- Initial factors are local CPTs (one per node)

$$P(R)$$

+r	0.1
-r	0.9

$$P(T|R)$$

+r	+t	0.8
+r	-t	0.2
-r	+t	0.1
-r	-t	0.9

$$P(L|T)$$

+t	+l	0.3
+t	-l	0.7
-t	+l	0.1
-t	-l	0.9

- Any known values are selected

- E.g. if we know $L = +\ell$, the initial factors are

$$P(R)$$

+r	0.1
-r	0.9

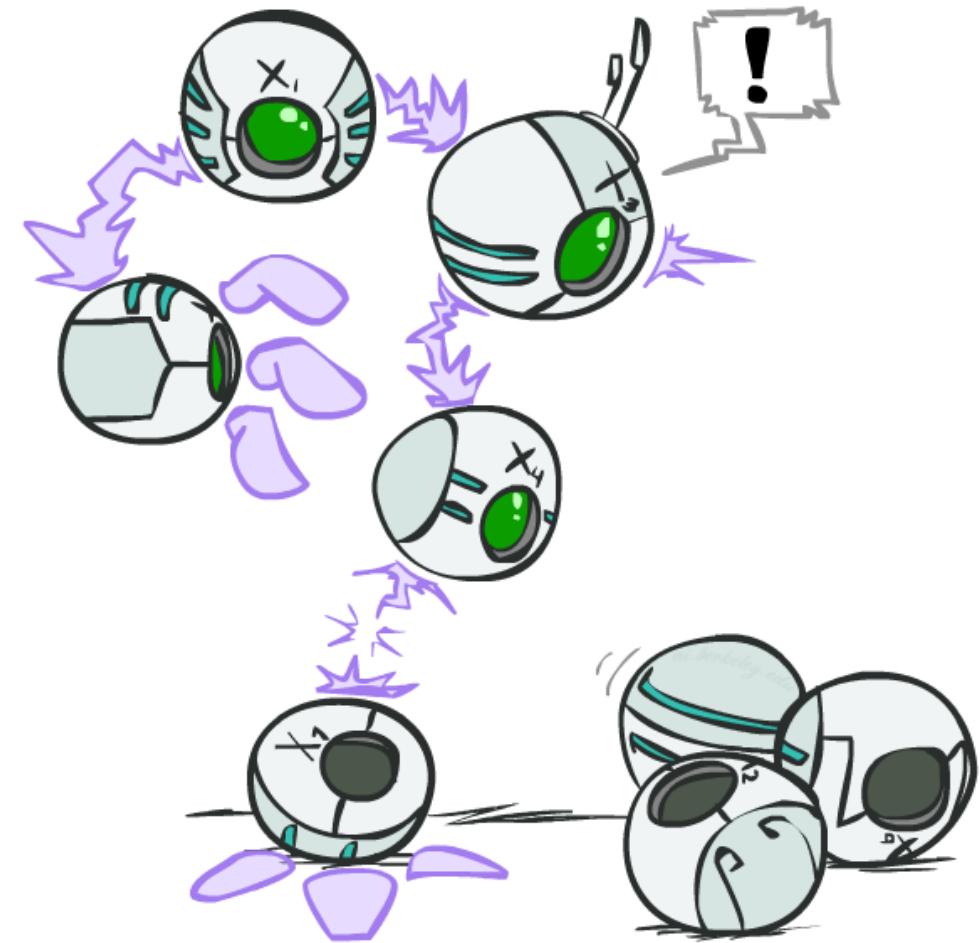
$$P(T|R)$$

+r	+t	0.8
+r	-t	0.2
-r	+t	0.1
-r	-t	0.9

$$P(+\ell|T)$$

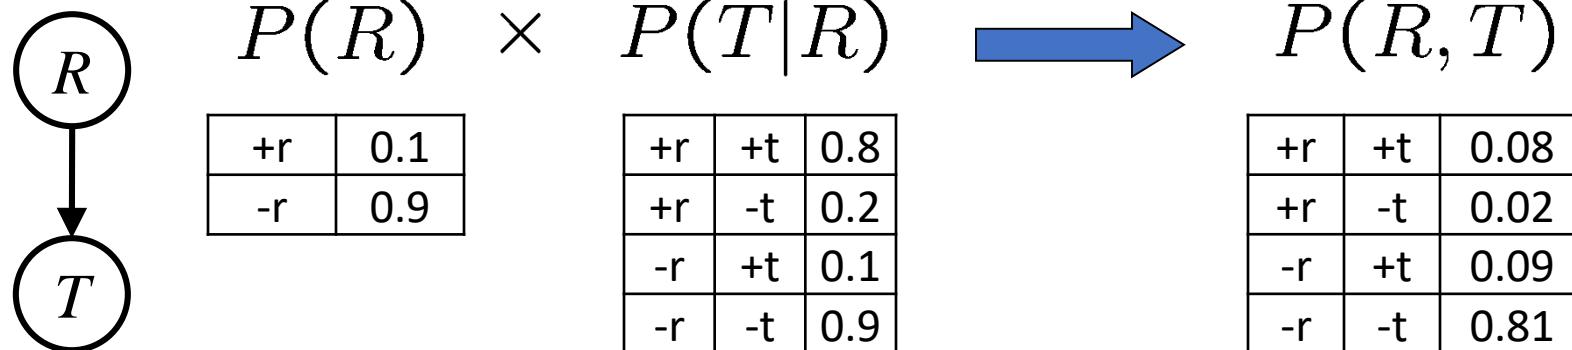
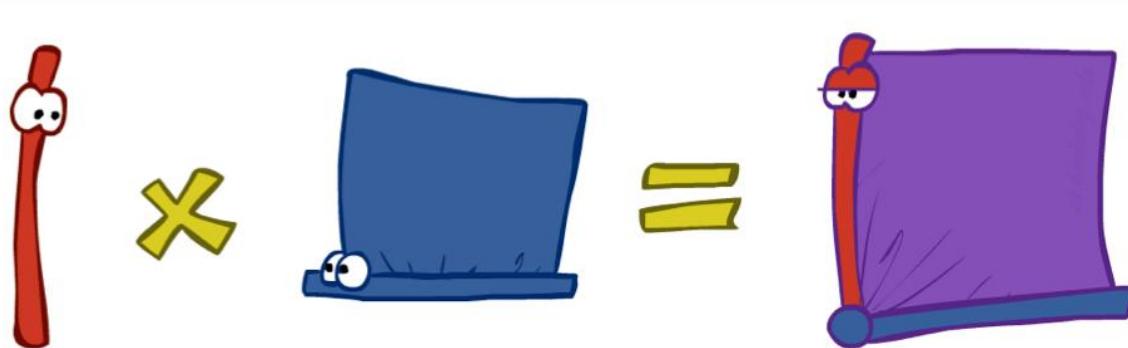
+t	+l	0.3
-t	+l	0.1

- Procedure: Join all factors, then sum out all hidden variables



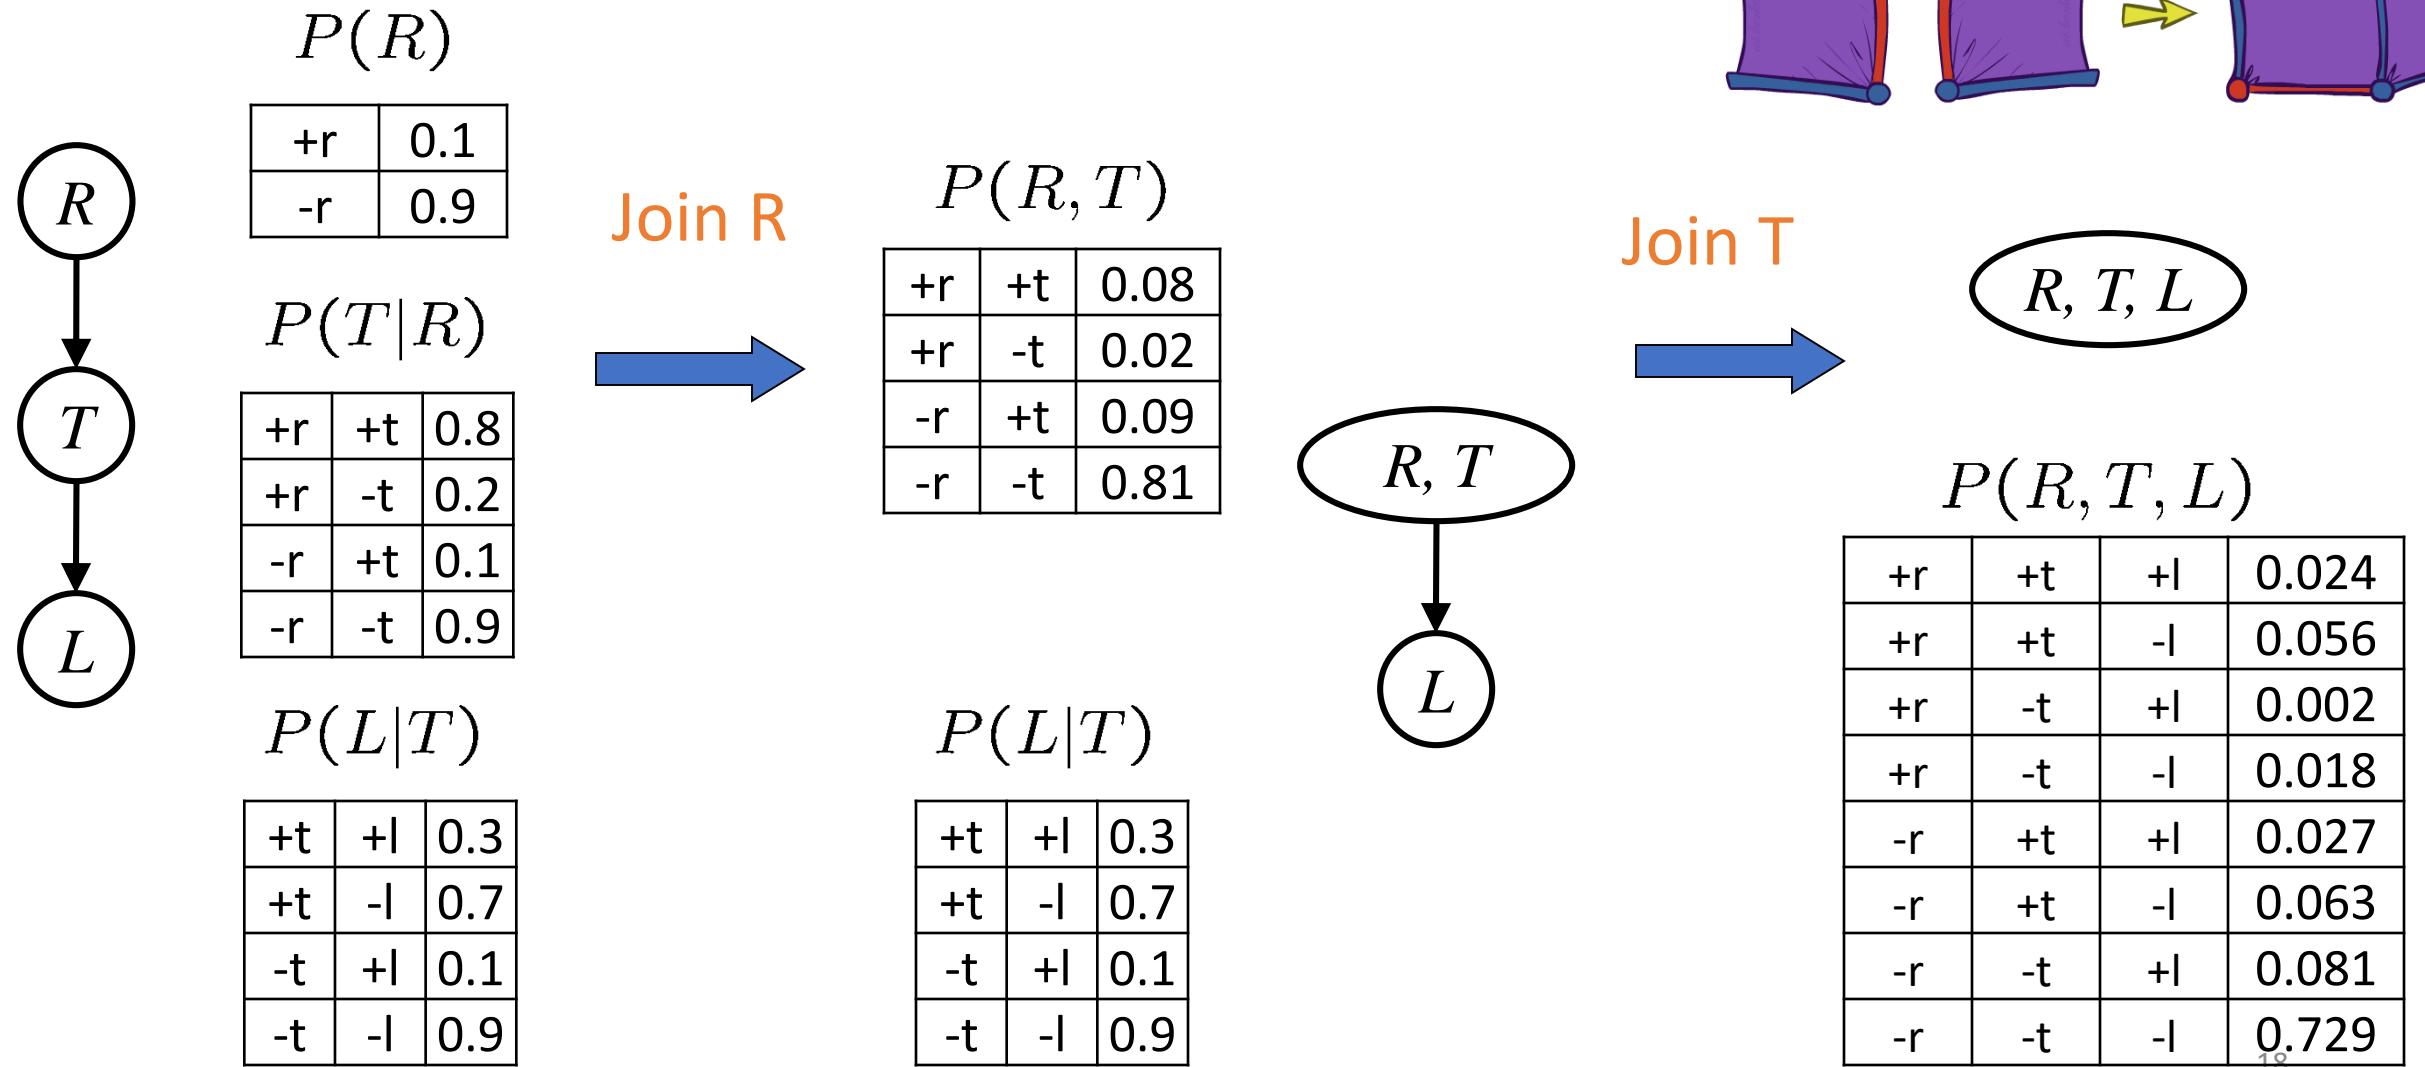
Operation 1: Join Factors

- First basic operation: **joining factors**
- Combining factors:
 - Just like a database join
 - Get all factors over the joining variable
 - Build a new factor over the union of the variables involved
- Example: Join on R

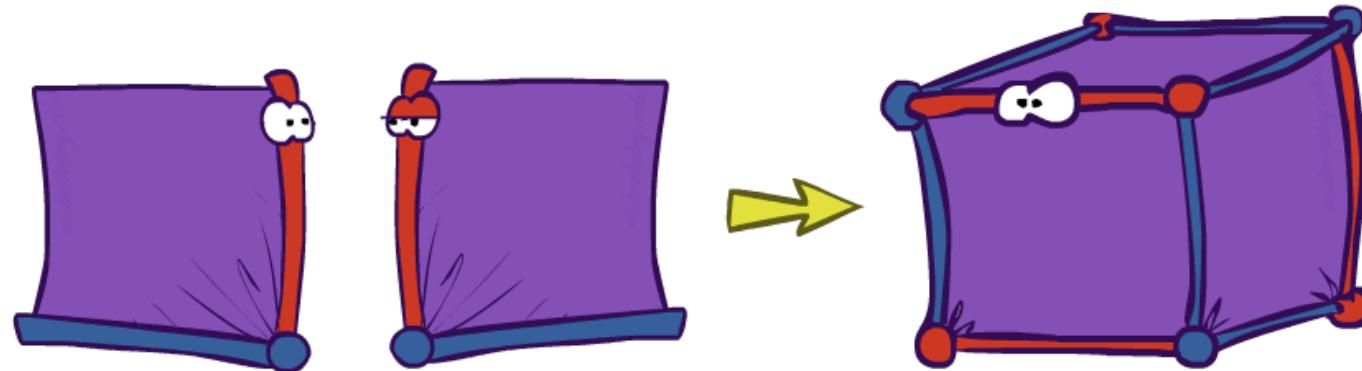


- Computation for each entry: pointwise products $\forall r, t : P(r, t) = P(r) \cdot P(t|r)$

Example: Multiple Joins



Example: Joining two conditional factors



- Example: $P(J|A) \times P(M|A) = P(J,M|A)$

$P(J A)$		
$A \setminus J$	true	false
true	0.99	0.01
false	0.145	0.855

\times

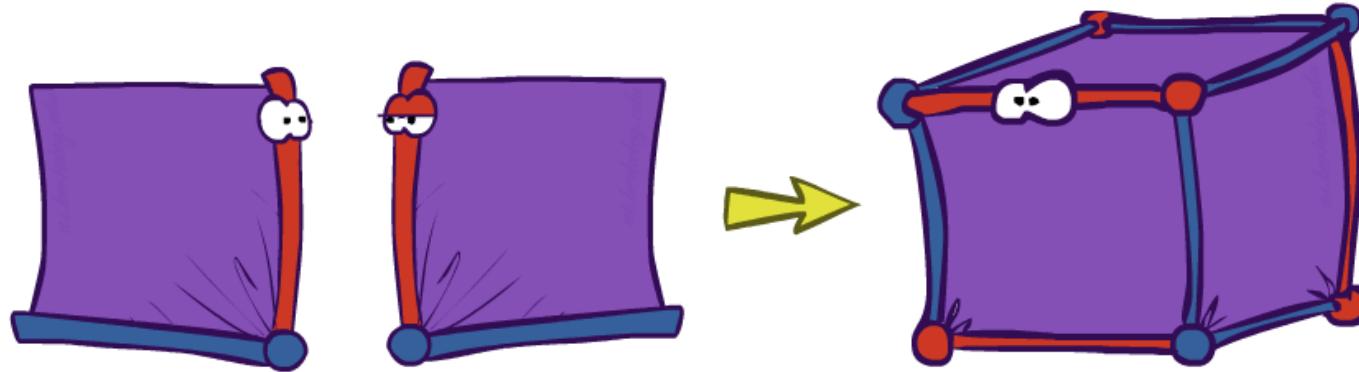
$P(M A)$		
$A \setminus M$	true	false
true	0.97	0.03
false	0.019	0.891

=

$P(J,M A)$		
$J \setminus M$	true	false
true		
false		.0003

18 A=false
 A=true

Example: Making larger factors



- Example: $f_1(U,V) \times f_2(V,W) \times f_3(W,X) = f_4(U,V,W,X)$
- Sizes: $[10,10] \times [10,10] \times [10,10] = [10,10,10,10]$
- i.e., 300 numbers blows up to 10,000 numbers!
- Factor blowup can make joining very expensive

Operation 2: Eliminate

- Second basic operation: **marginalization**
- Take a factor and sum out a variable
 - Shrinks a factor to a smaller one
 - A **projection** operation
- Example:

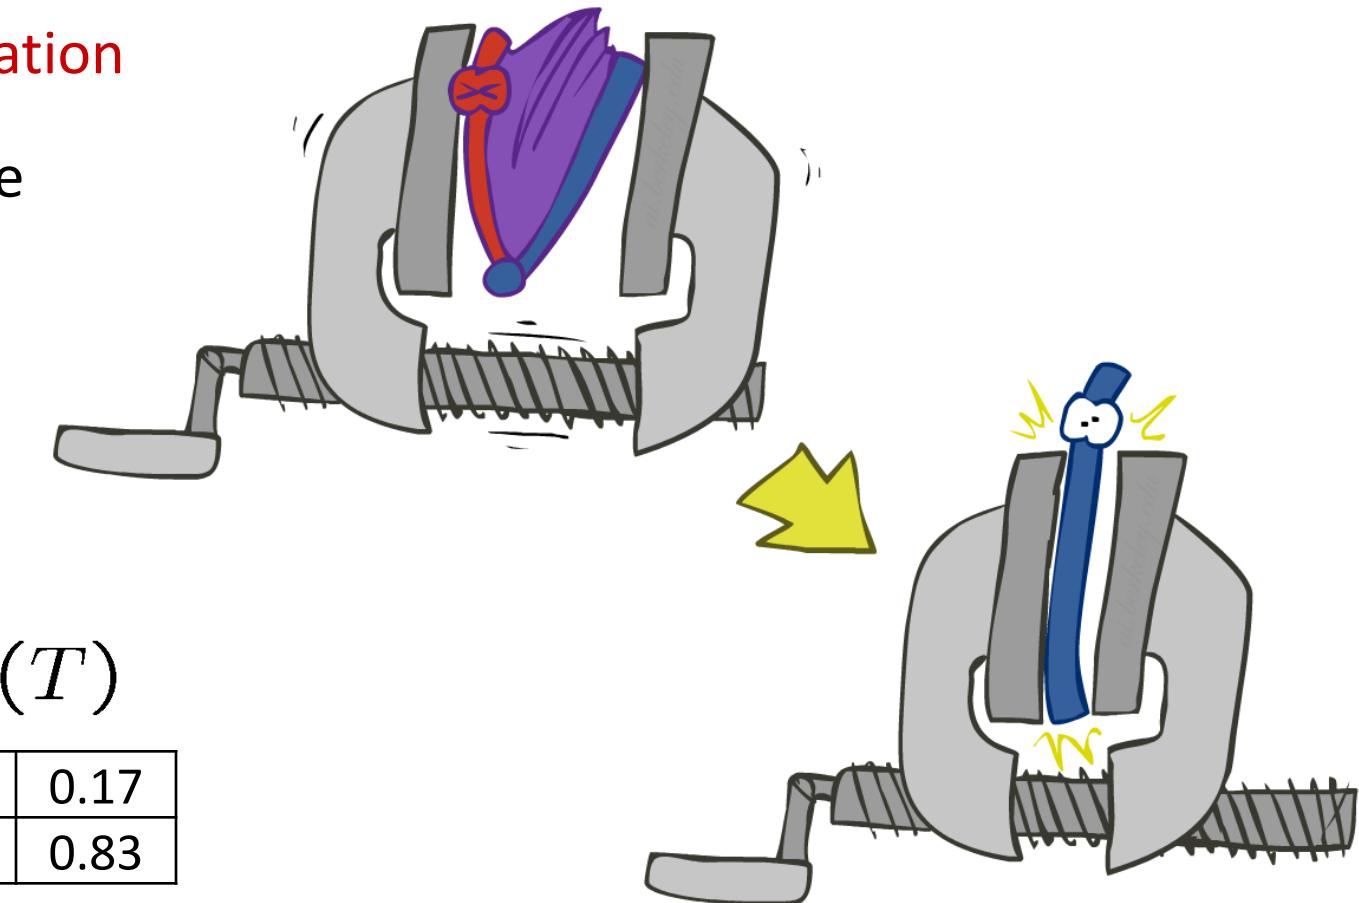
$$P(R, T)$$

+r	+t	0.08
+r	-t	0.02
-r	+t	0.09
-r	-t	0.81

sum R

$$P(T)$$

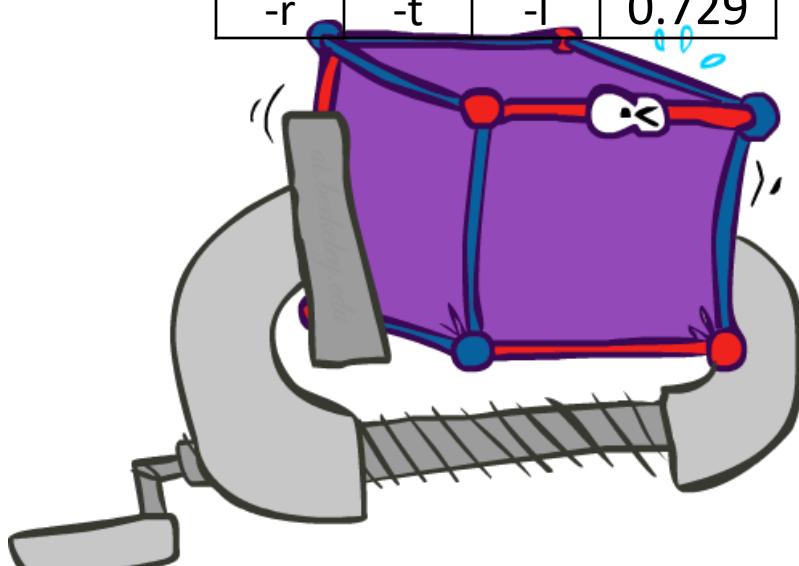
+t	0.17
-t	0.83



Multiple Elimination

$P(R, T, L)$

+r	+t	+l	0.024
+r	+t	-l	0.056
+r	-t	+l	0.002
+r	-t	-l	0.018
-r	+t	+l	0.027
-r	+t	-l	0.063
-r	-t	+l	0.081
-r	-t	-l	0.729



Sum
out R

T, L

$P(T, L)$

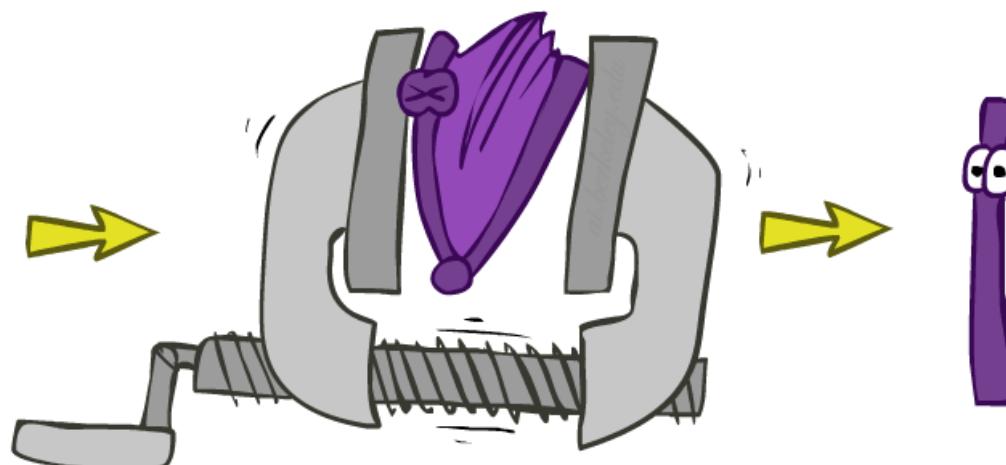
+t	+l	0.051
+t	-l	0.119
-t	+l	0.083
-t	-l	0.747

Sum
out T

L

$P(L)$

+l	0.134
-l	0.866



Thus Far: Multiple Join, Multiple Eliminate (= Inference by Enumeration)

$$P(R)$$

$$P(T|R)$$



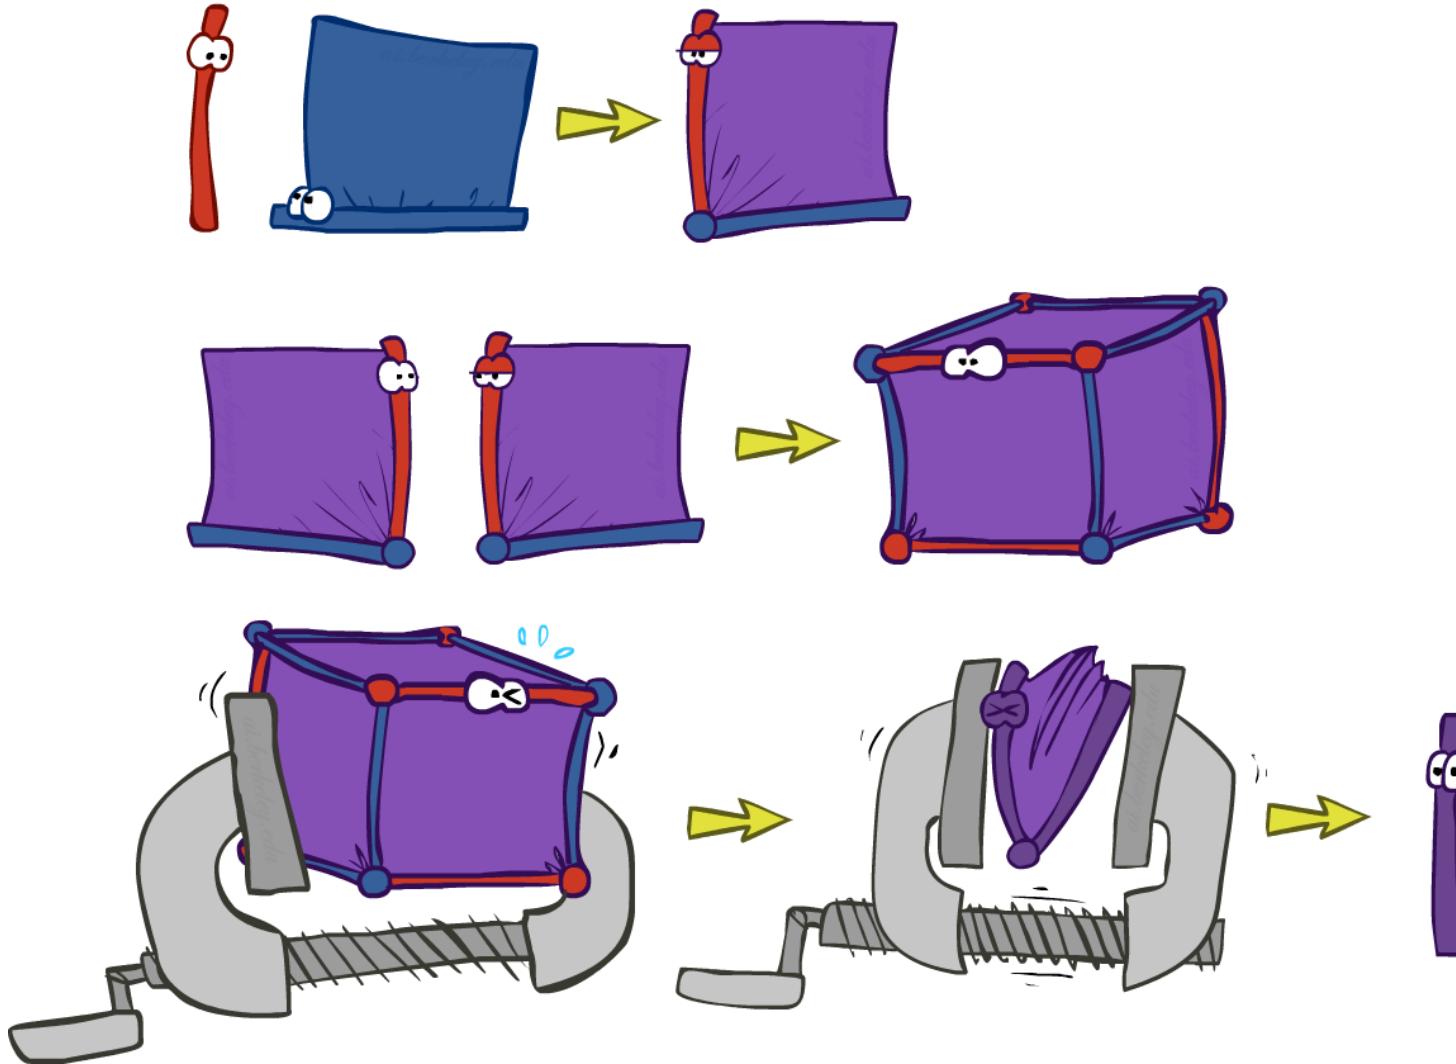
$$P(R, T, L)$$



$$P(L)$$

$$P(L|T)$$

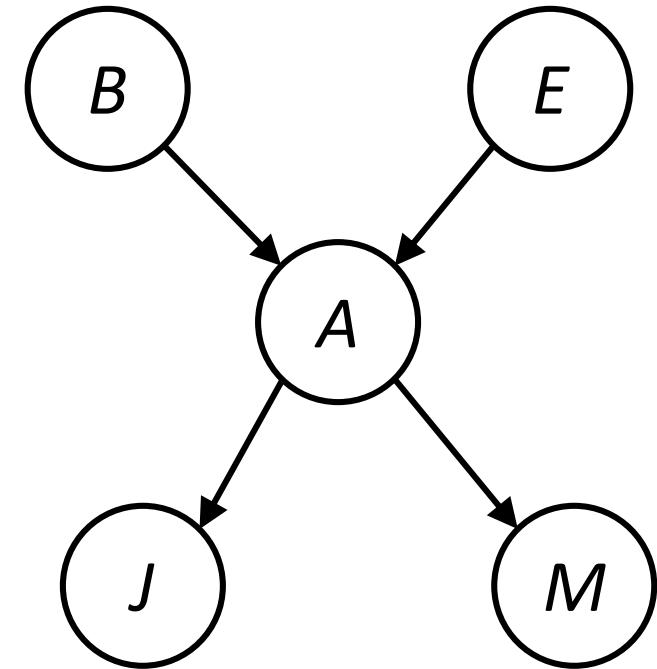
Thus Far: Multiple Join, Multiple Eliminate (= Inference by Enumeration)



Inference by Enumeration in Bayes Net

- Reminder of inference by enumeration:
 - Any probability of interest can be computed by summing entries from the joint distribution
 - Entries from the joint distribution can be obtained from a BN by multiplying the corresponding conditional probabilities

$$\begin{aligned} P(B \mid j, m) &= \alpha P(B, j, m) \\ &= \alpha \sum_{e,a} P(B, e, a, j, m) \\ &= \alpha \sum_{e,a} P(B) P(e) P(a|B,e) P(j|a) P(m|a) \end{aligned}$$



- So inference in Bayes nets means computing sums of products of numbers: sounds easy!!
- Problem: sums of *exponentially many* products!

Can we do better?

- Consider
 - $x_1y_1z_1 + x_1y_1z_2 + x_1y_2z_1 + x_1y_2z_2 + x_2y_1z_1 + x_2y_1z_2 + x_2y_2z_1 + x_2y_2z_2$
 - 16 multiplies, 7 adds
 - Lots of repeated subexpressions!
- Rewrite as
 - $(x_1 + x_2)(y_1 + y_2)(z_1 + z_2)$
 - 2 multiplies, 3 adds

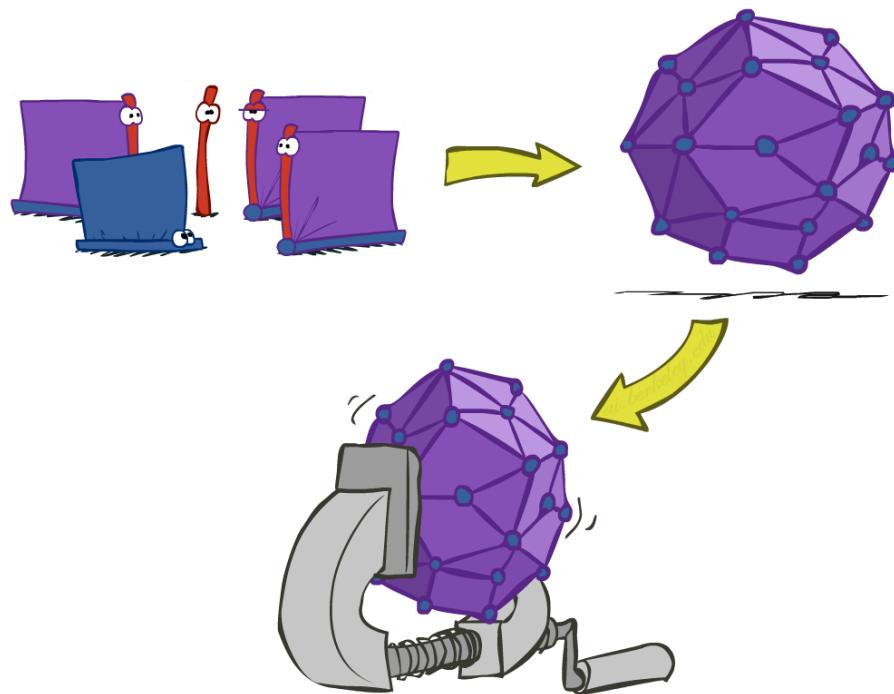
$$\begin{aligned} \sum_e \sum_a & P(B) P(e) P(a | B, e) P(j | a) P(m | a) \\ &= P(B) P(+e) P(+a | B, +e) P(j | +a) P(m | +a) \\ &\quad + P(B) P(-e) P(+a | B, -e) P(j | +a) P(m | +a) \\ &\quad + P(B) P(+e) P(-a | B, +e) P(j | -a) P(m | -a) \\ &\quad + P(B) P(-e) P(-a | B, -e) P(j | -a) P(m | -a) \end{aligned}$$

- Lots of repeated subexpressions!

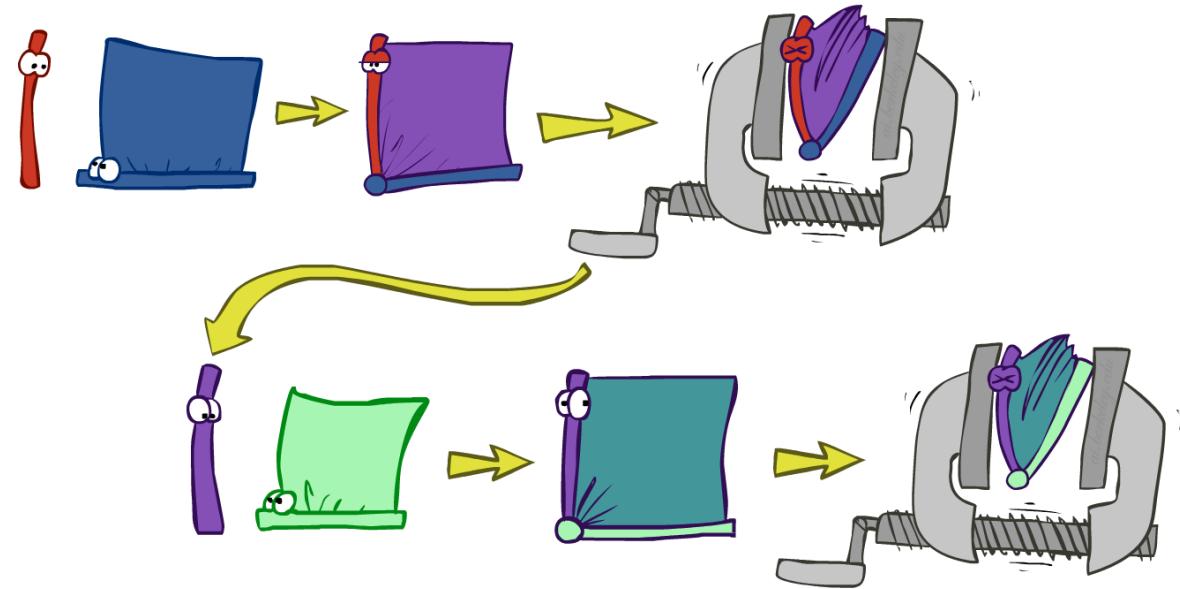
Variable Elimination

Inference by Enumeration vs. Variable Elimination

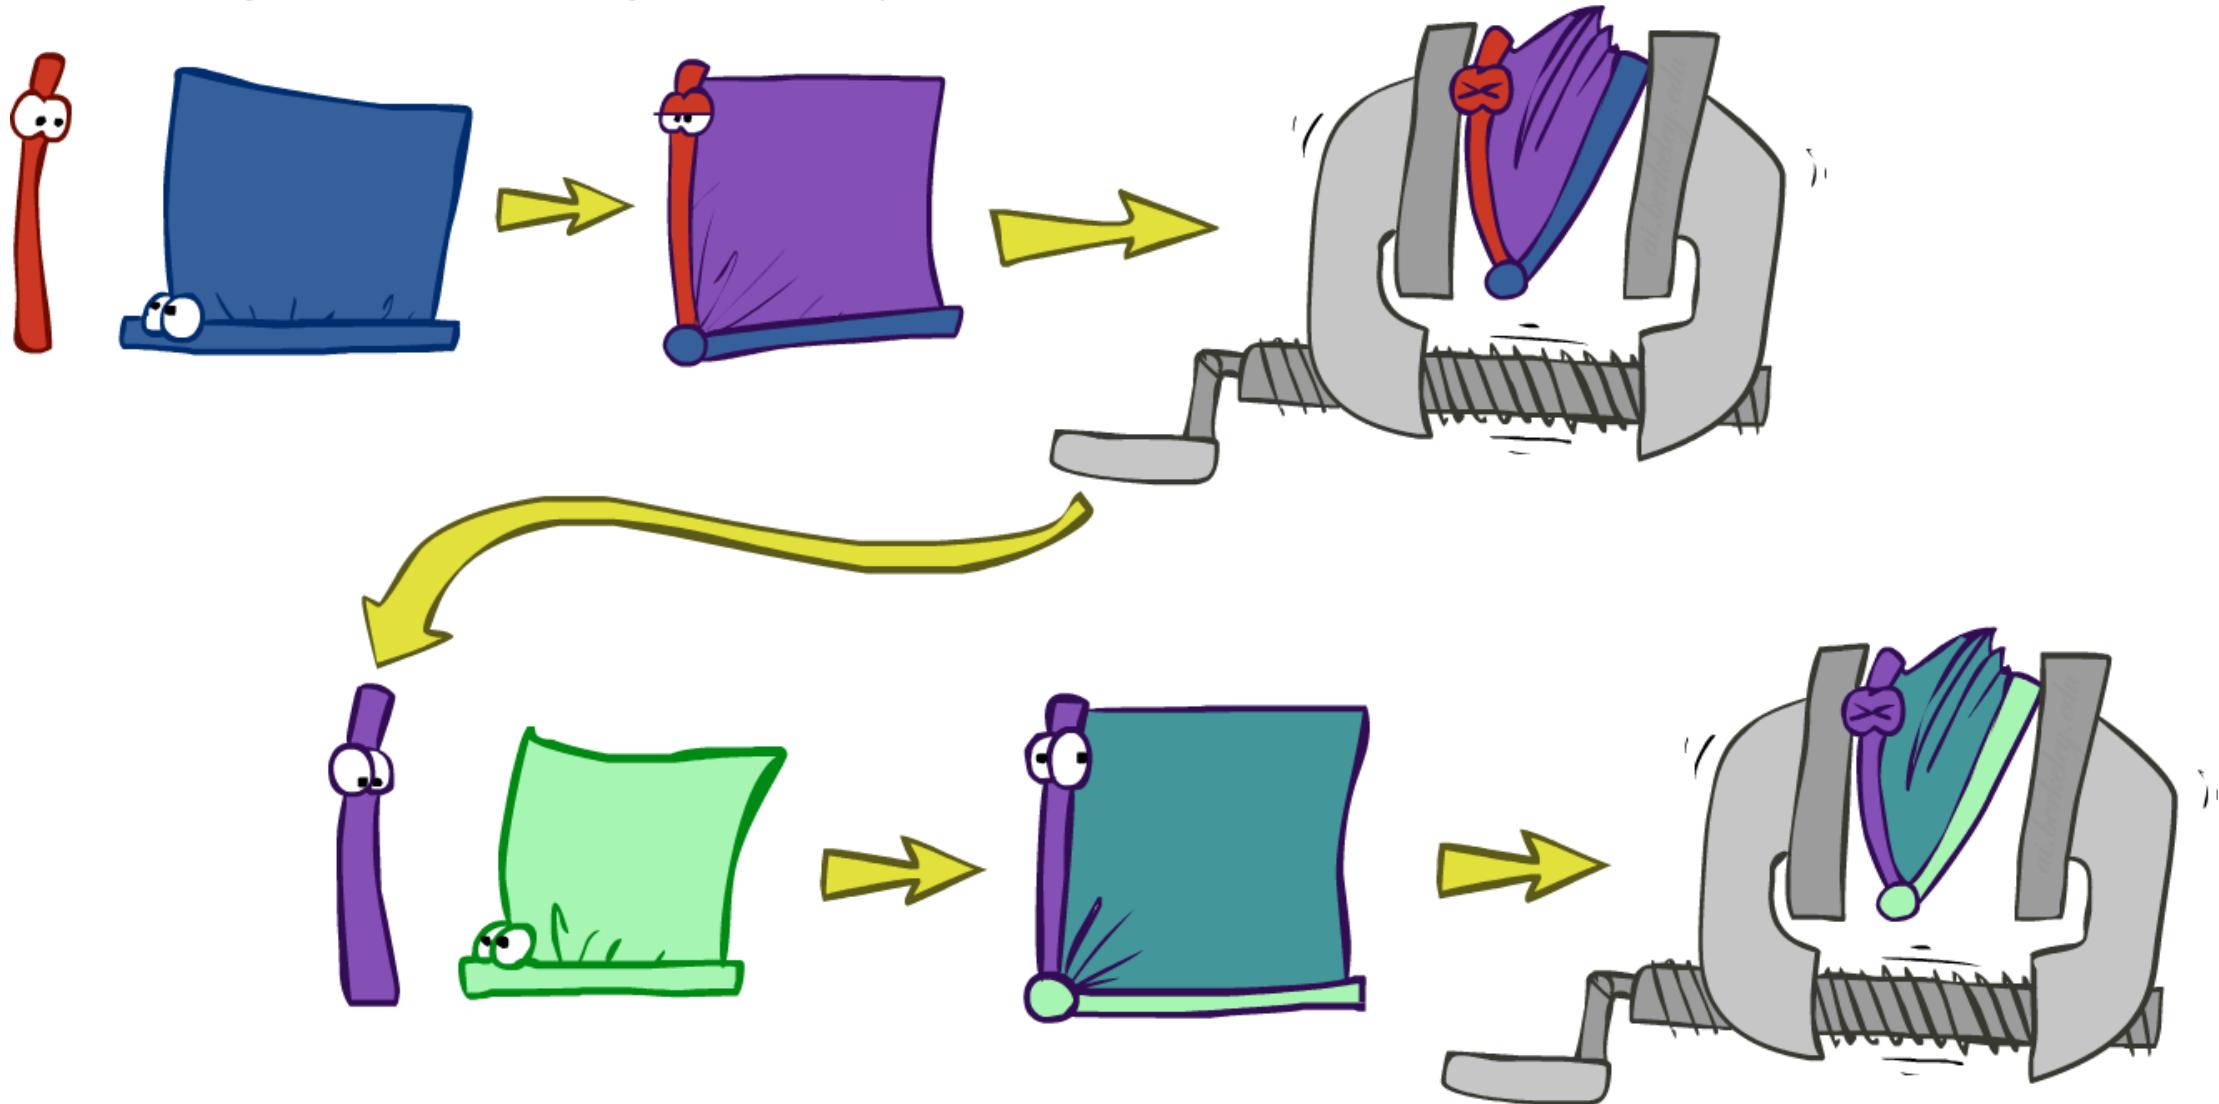
- Why is inference by enumeration so slow?
 - You join up the whole joint distribution before you sum out the hidden variables



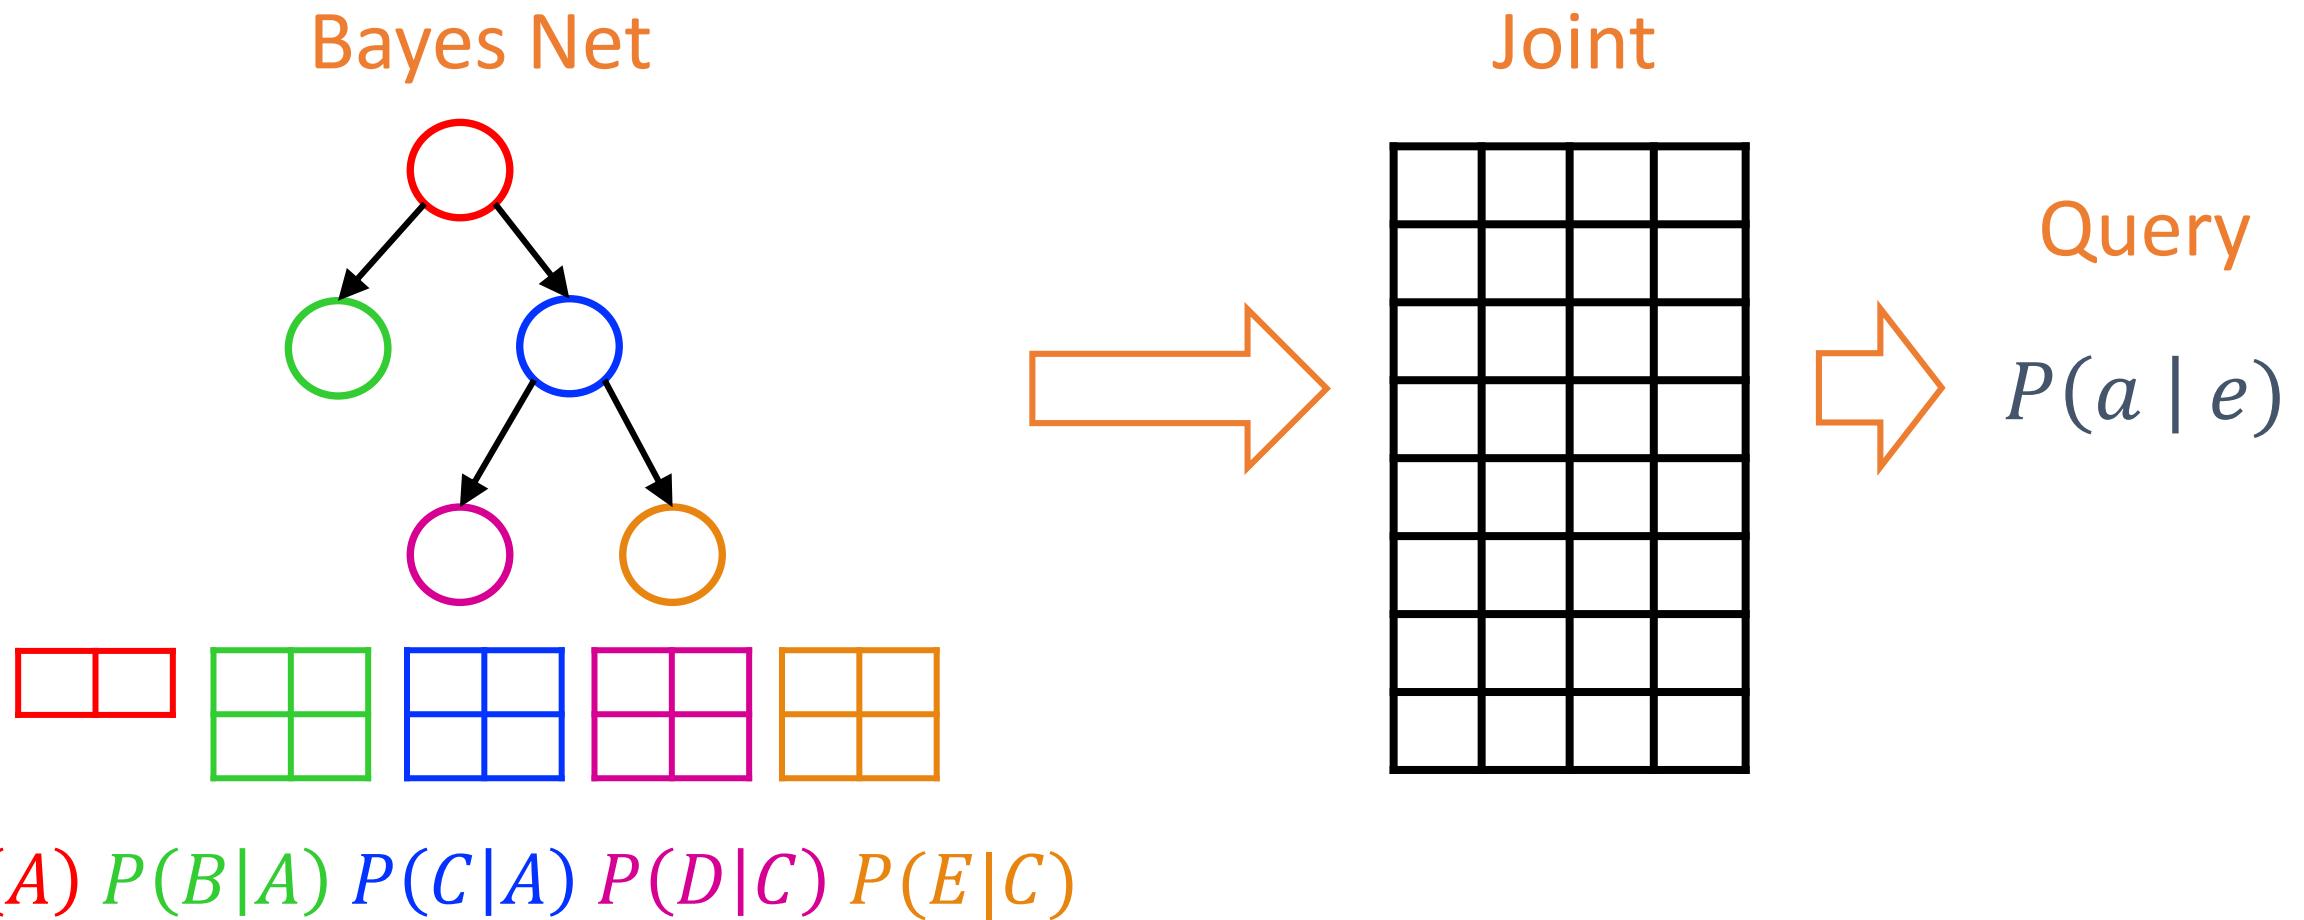
- Idea: interleave joining and marginalizing!
 - Called “Variable Elimination”
 - Still NP-hard, but usually much faster than inference by enumeration



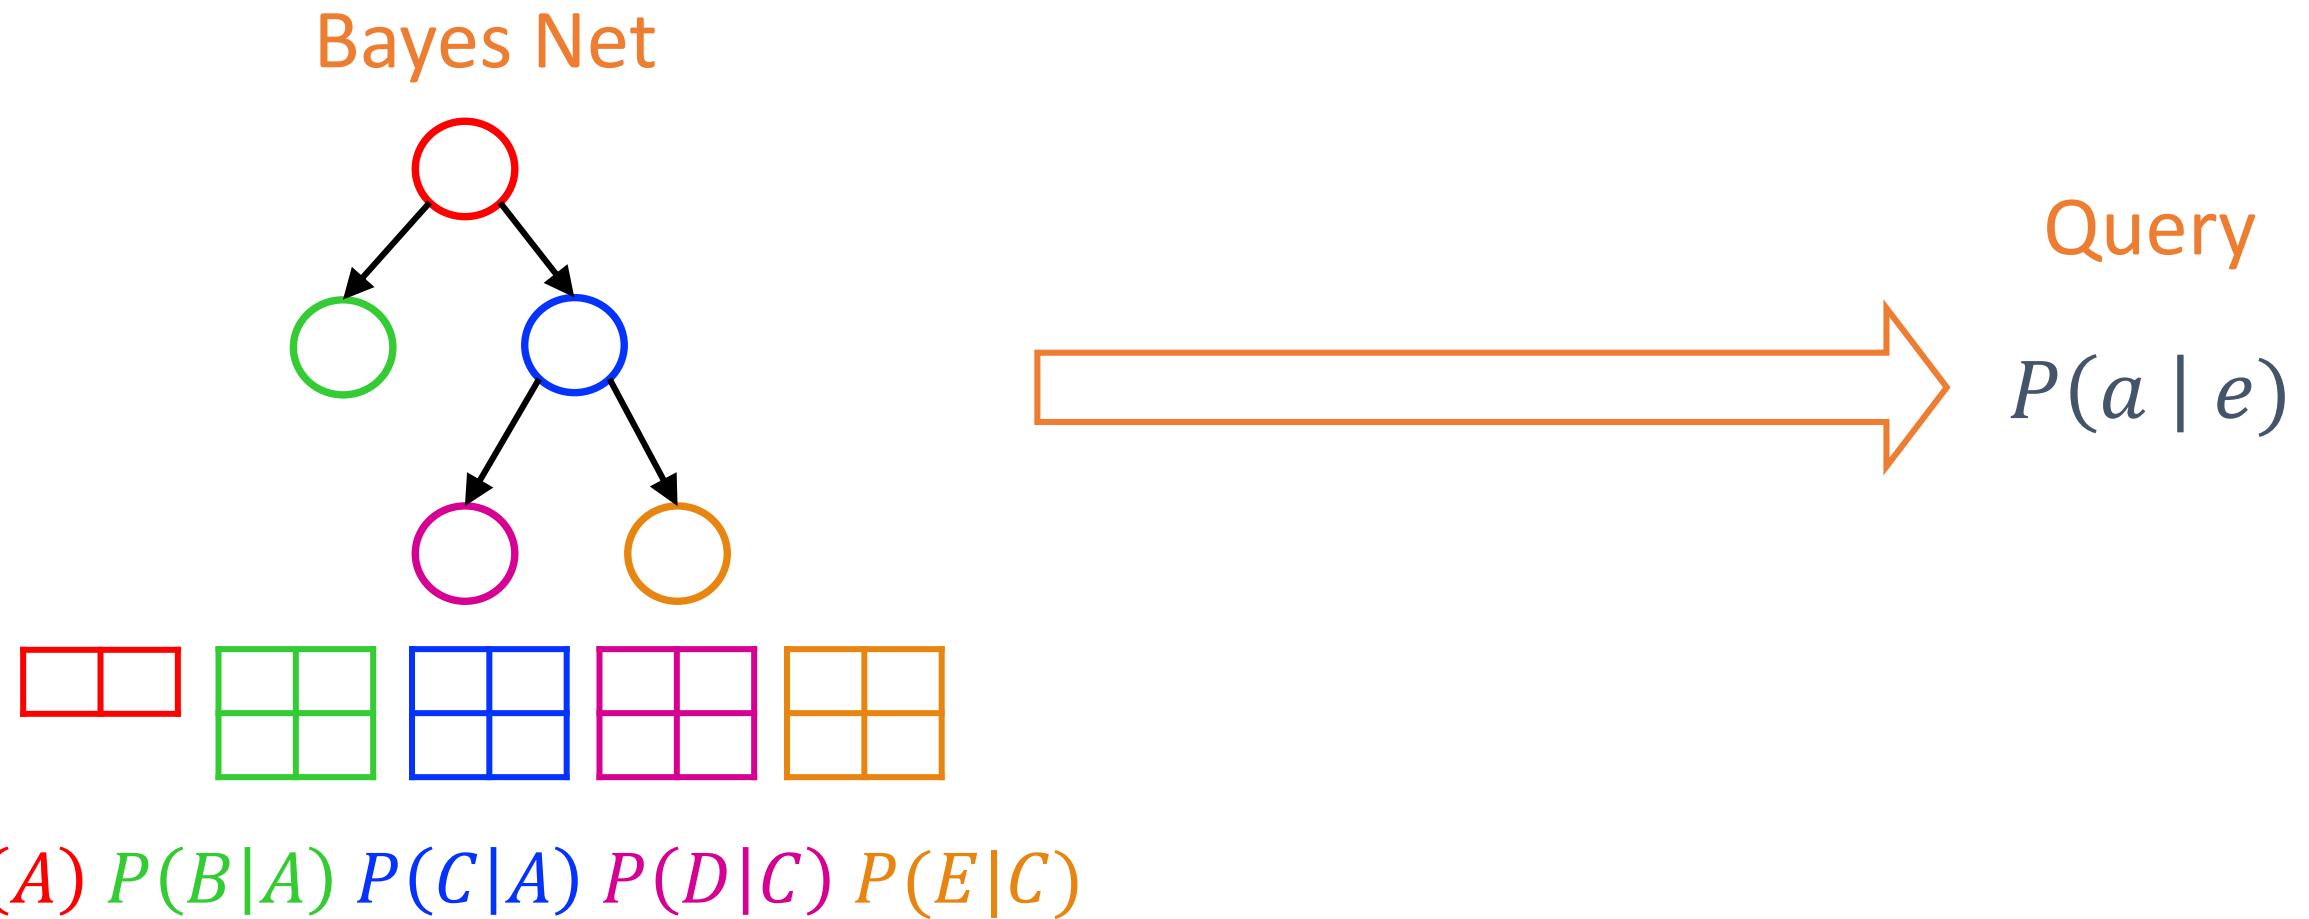
Marginalizing Early (= Variable Elimination)



Answer Any Query from Bayes Net (Previous)

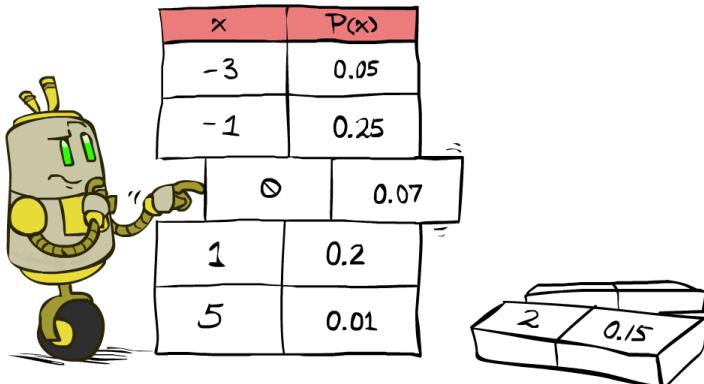


Next: Answer Any Query from Bayes Net



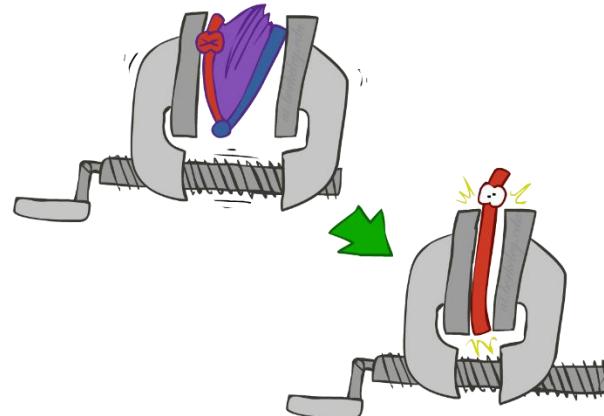
Inference by Enumeration

- General case:
 - Evidence variables: $E_1 \dots E_k = e_1 \dots e_k$
 - Query* variable: Q
 - Hidden variables: $H_1 \dots H_r$
- Step 1: Select the entries consistent with the evidence



$$P(Q, e_1 \dots e_k) = \sum_{h_1 \dots h_r} \underbrace{P(Q, h_1 \dots h_r, e_1 \dots e_k)}_{\text{Compute joint}}$$

- Sum out hidden variables X_1, X_2, \dots, X_n



- We want:

* Works fine with multiple query variables, too

$$P(Q|e_1 \dots e_k)$$

- Step 3: Normalize

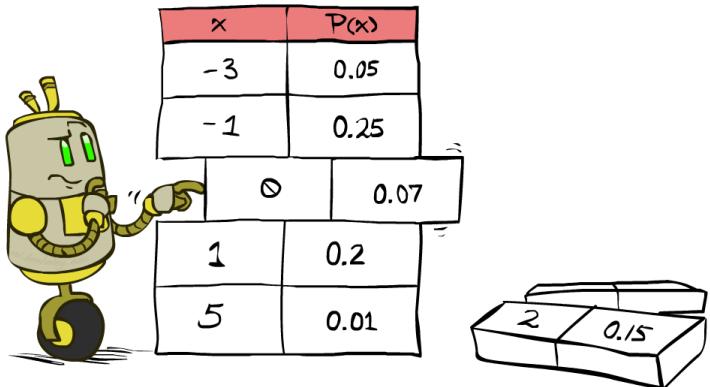
$$\times \frac{1}{Z}$$

$$Z = \sum_q P(Q, e_1 \dots e_k)$$

$$P(Q|e_1 \dots e_k) = \frac{1}{Z} P(Q, e_1 \dots e_k)$$

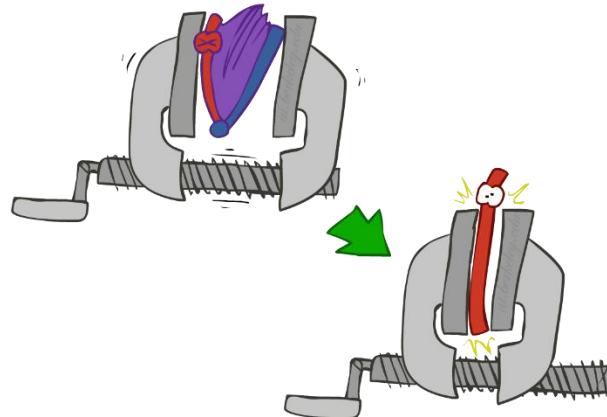
Variable Elimination

- General case:
 - Evidence variables: $E_1 \dots E_k = e_1 \dots e_k$
 - Query* variable: Q
 - Hidden variables: $H_1 \dots H_r$
- Step 1: Select the entries consistent with the evidence



X_1, X_2, \dots, X_n
 } All variables
 $E_1 \dots E_k = e_1 \dots e_k$
 $H_1 \dots H_r$

- Step 2: Sum out H to get joint of Query and evidence



$$P(Q, e_1 \dots e_k) = \sum_{h_1 \dots h_r} \underbrace{P(Q, h_1 \dots h_r, e_1 \dots e_k)}_{X_1, X_2, \dots, X_n}$$

- Interleave joining and summing out X_1, X_2, \dots, X_n

- We want:

* Works fine with multiple query variables, too

$$P(Q | e_1 \dots e_k)$$

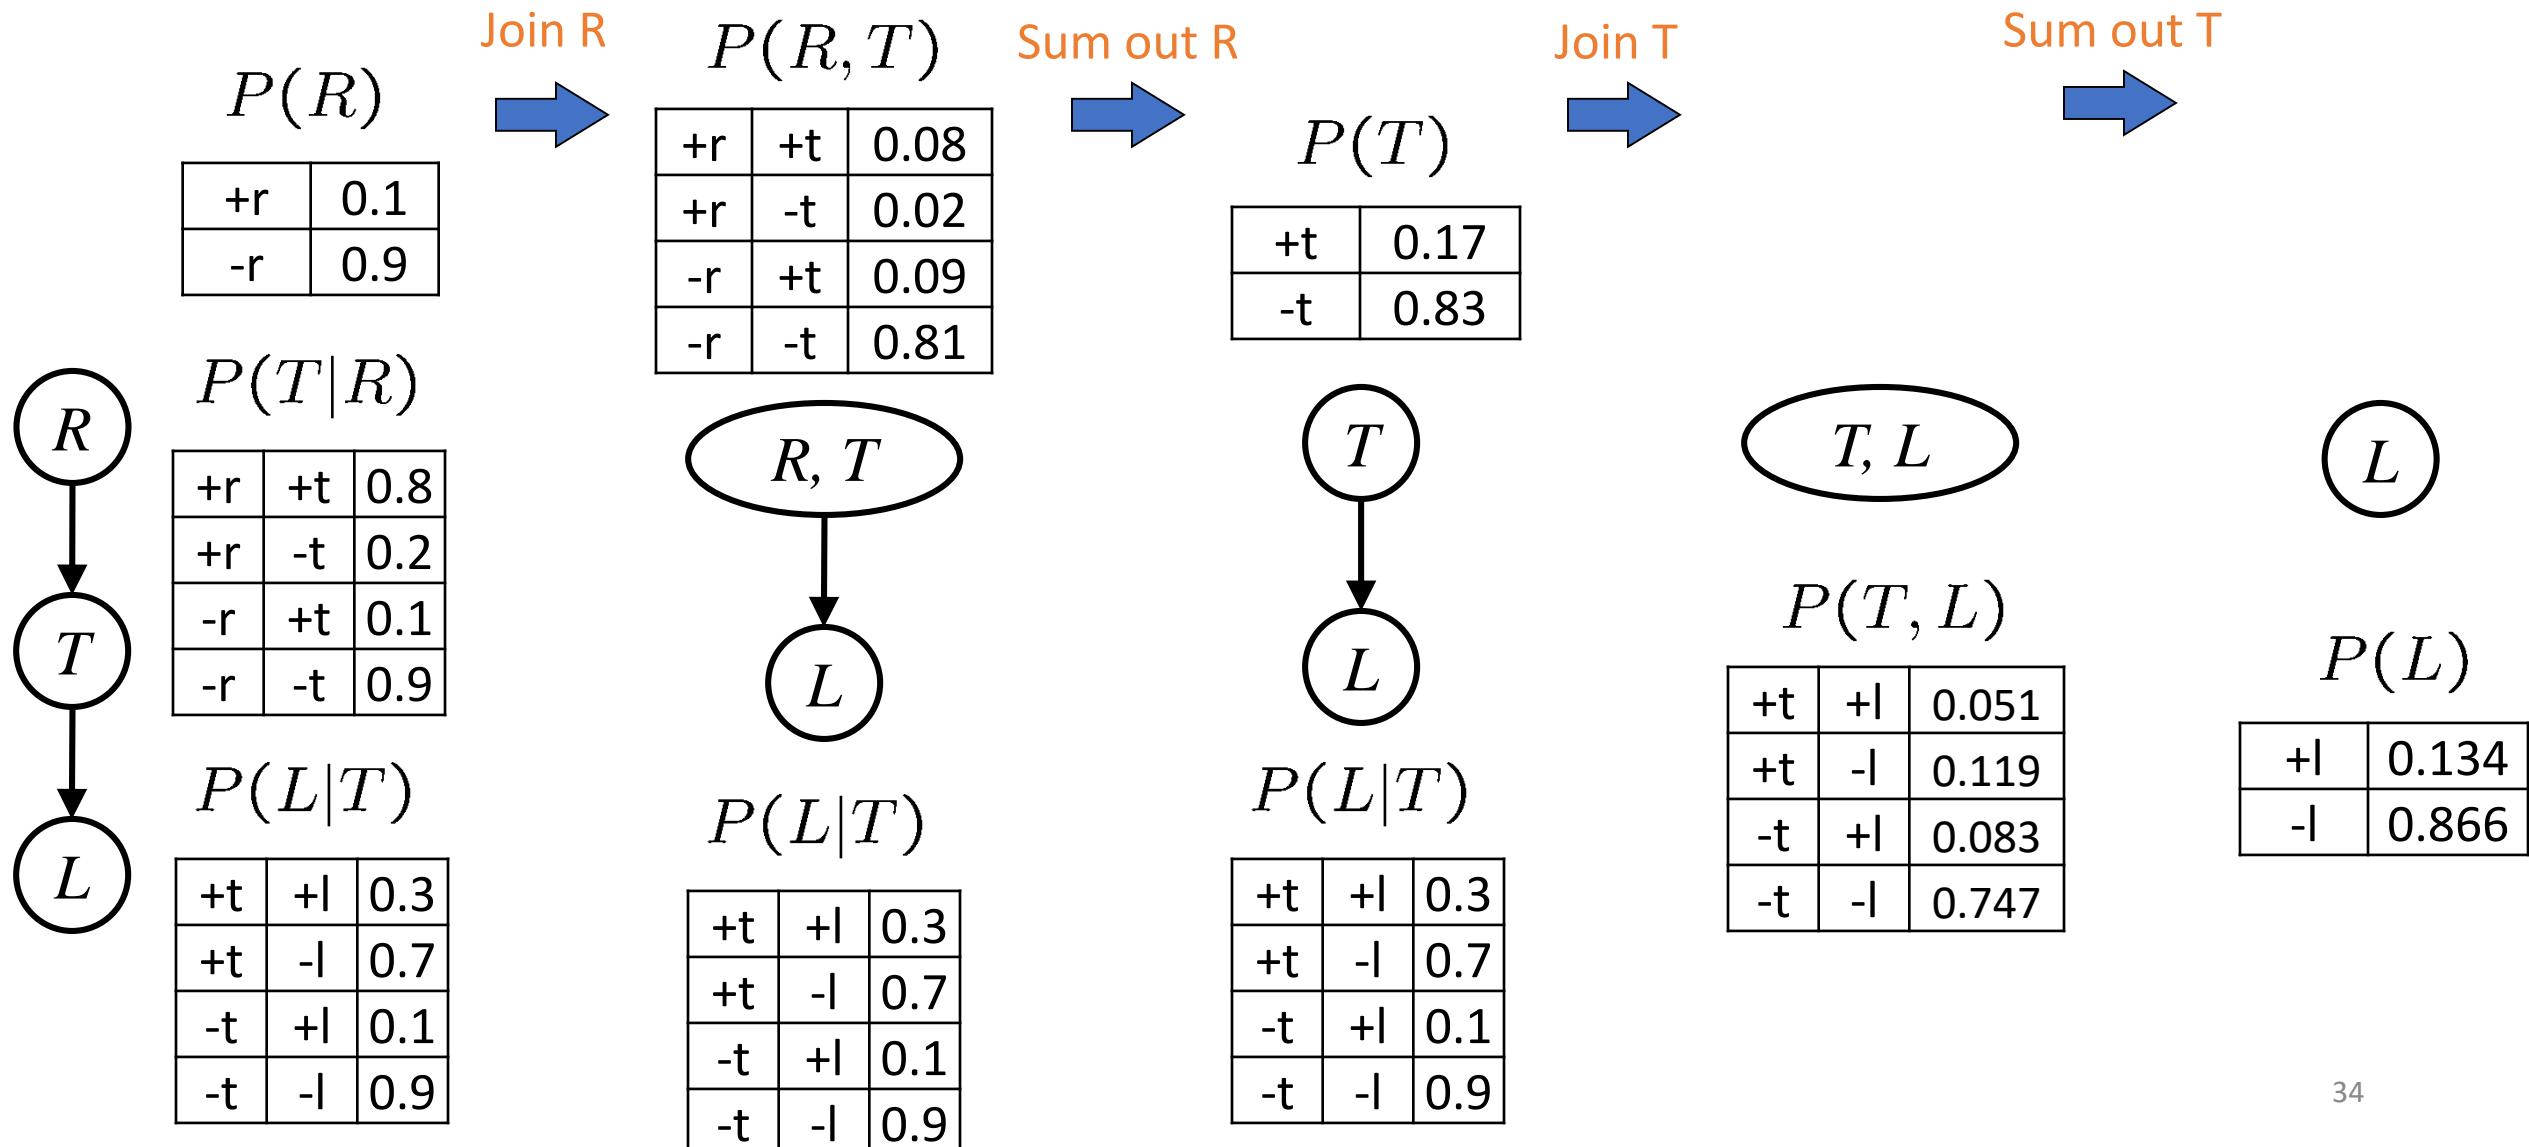
- Step 3: Normalize

$$\times \frac{1}{Z}$$

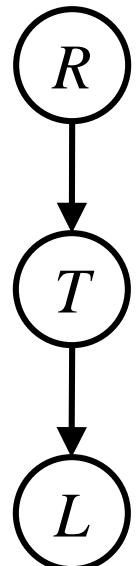
$$Z = \sum_q P(Q, e_1 \dots e_k)$$

$$P(Q | e_1 \dots e_k) = \frac{1}{Z} P(Q, e_1 \dots e_k)$$

Marginalizing Early! (aka VE)



Traffic Domain



$$P(L) = ?$$

- Inference by Enumeration

$$= \sum_t \sum_r P(L|t) P(r) P(t|r)$$

Join on r
Join on t
Eliminate r
Eliminate t

- Variable Elimination

$$= \sum_t P(L|t) \sum_r P(r) P(t|r)$$

Join on r
Join on t
Eliminate r
Eliminate t

Inference Overview

- Given random variables Q, H, E (query, hidden, evidence)
- We know how to do inference on a joint distribution

$$\begin{aligned} P(q|e) &= \alpha P(q, e) \\ &= \alpha \sum_{h \in \{h_1, h_2\}} P(q, h, e) \end{aligned}$$

- We know Bayes nets can break down joint in to CPT factors

$$\begin{aligned} P(q|e) &= \alpha \sum_{h \in \{h_1, h_2\}} P(h) P(q|h) P(e|q) \\ &= \alpha [P(h_1) P(q|h_1) P(e|q) + P(h_2) P(q|h_2) P(e|q)] \end{aligned}$$

- But we can be more efficient

$$\begin{aligned} P(q|e) &= \alpha P(e|q) \sum_{h \in \{h_1, h_2\}} P(h) P(q|h) \\ &= \alpha P(e|q) [P(h_1) P(q|h_1) + P(h_2) P(q|h_2)] \\ &= \alpha P(e|q) P(q) \end{aligned}$$

- Now just extend to larger Bayes nets and a variety of queries



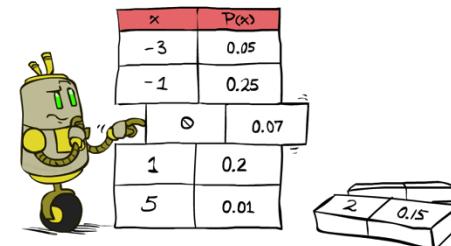
Enumeration
Variable
Elimination

Variable Elimination: The basic ideas

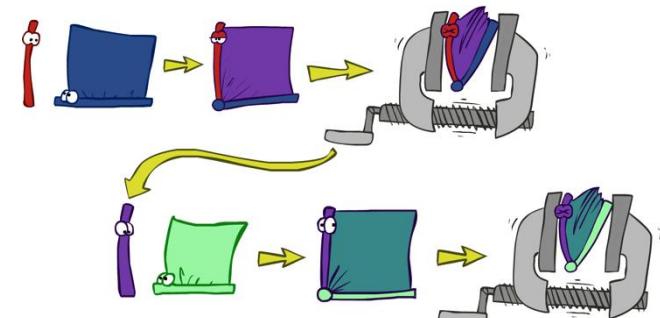
- Move summations inwards as far as possible
 - $$\begin{aligned} P(B \mid j, m) &= \alpha \sum_e \sum_a P(B) P(e) P(a \mid B, e) P(j \mid a) P(m \mid a) \\ &= \alpha P(B) \sum_e P(e) \sum_a P(a \mid B, e) P(j \mid a) P(m \mid a) \end{aligned}$$
 - Do the calculation from the inside out
 - I.e., sum over a first, then sum over e
 - Problem: $P(a \mid B, e)$ isn't a single number, it's a bunch of different numbers depending on the values of B and e
 - Solution: use arrays of numbers (of various dimensions) with appropriate operations on them; these are **factors**
-

General Variable Elimination

- Query: $P(Q|E_1 = e_1, \dots, E_k = e_k)$
- Start with initial factors:
 - Local CPTs (but instantiated by evidence)
- While there are still hidden variables (not Q or evidence):
 - Pick a hidden variable H
 - Join all factors mentioning H
 - Eliminate (sum out) H
- Join all remaining factors and normalize



x	$P(x)$
-3	0.05
-1	0.25
0	0.07
1	0.2
5	0.01
2	0.15



$$f \times \text{[blue flag]} = \text{[purple flag]} \times \frac{1}{Z}$$

Variable Elimination

```
function VariableElimination( $Q, e, bn$ ) returns a distribution over  $Q$ 
   $factors \leftarrow []$ 
  for each  $var$  in ORDER( $bn.vars$ ) do
     $factors \leftarrow [MAKE-FACTOR(var, e) | factors]$ 
    if  $var$  is a hidden variable then
       $factors \leftarrow SUM-OUT(var, factors)$ 
  return NORMALIZE(POINTWISE-PRODUCT(factors))
```

Evidence

- If evidence, start with factors that select that evidence

- No evidence, uses these initial factors:

$$P(R)$$

+r	0.1
-r	0.9

$$P(T|R)$$

+r	+t	0.8
+r	-t	0.2
-r	+t	0.1
-r	-t	0.9

$$P(L|T)$$

+t	+l	0.3
+t	-l	0.7
-t	+l	0.1
-t	-l	0.9

- Computing $P(L|+r)$, the initial factors become:

$$P(+r)$$

+r	0.1
----	-----

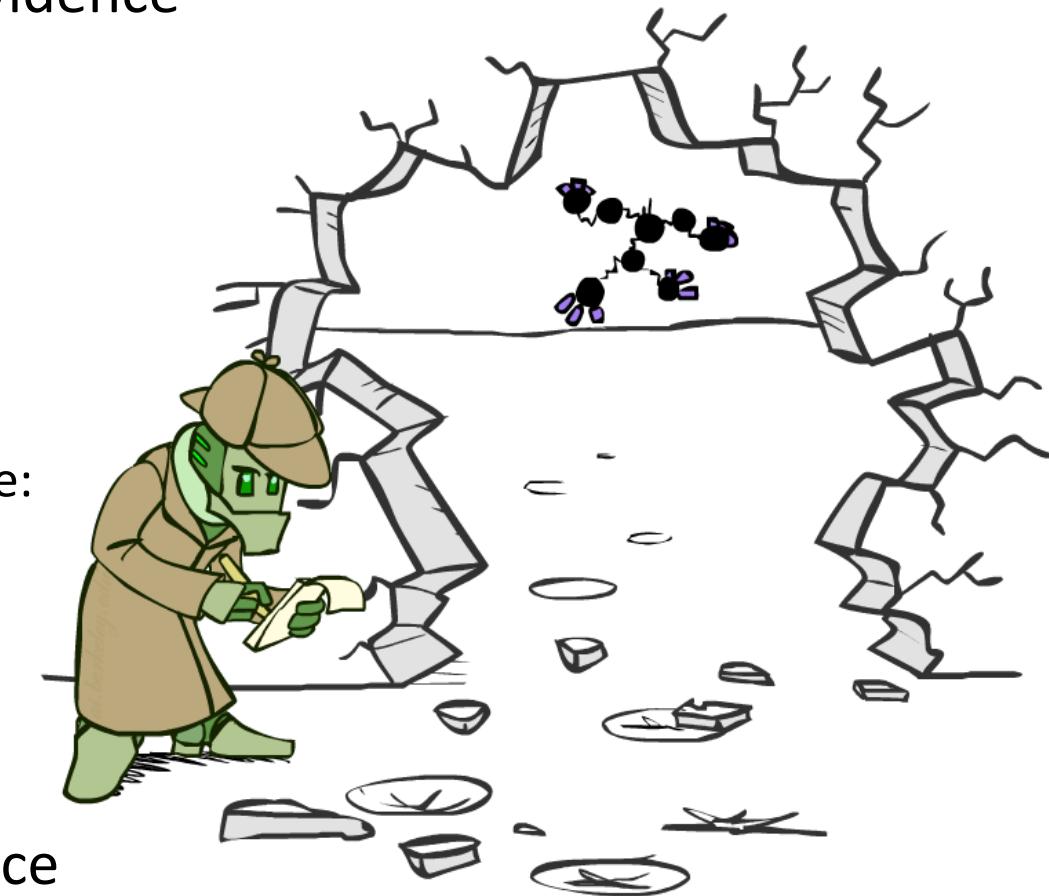
$$P(T|+r)$$

+r	+t	0.8
+r	-t	0.2

$$P(L|T)$$

+t	+l	0.3
+t	-l	0.7
-t	+l	0.1
-t	-l	0.9

- We eliminate all vars other than query + evidence



Evidence II

- Result will be a selected joint of query and evidence
 - E.g. for $P(L | +r)$, we would end up with:

$$P(+r, L)$$

+r	+l	0.026
+r	-l	0.074

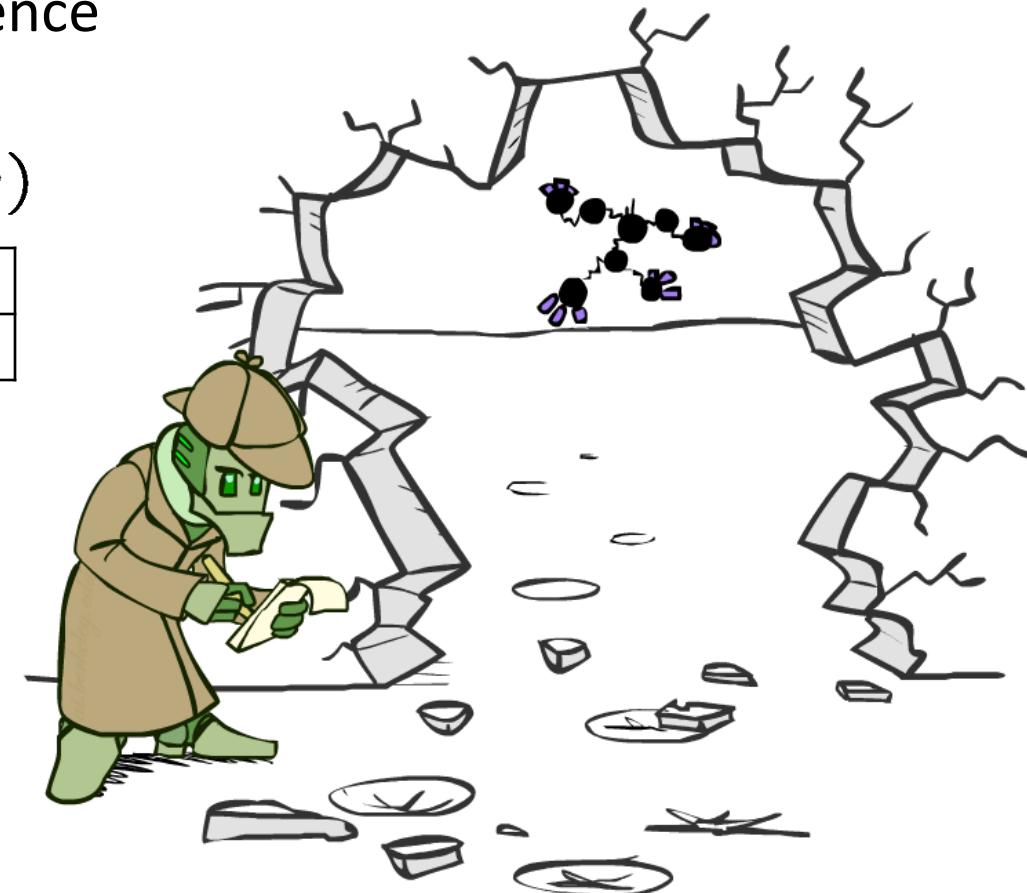
Normalize



$$P(L | +r)$$

+l	0.26
-l	0.74

- To get our answer, just normalize this!
- That's it!



Example

$$P(B|j, m) \propto P(B, j, m)$$

$P(B)$	$P(E)$	$P(A B, E)$	$P(j A)$	$P(m A)$
--------	--------	-------------	----------	----------

$$P(B|j, m) \propto P(B, j, m)$$

$$= \sum_{e,a} P(B, j, m, e, a)$$

$$= \sum_{e,a} P(B)P(e)P(a|B, e)P(j|a)P(m|a)$$

$$= \sum_e P(B)P(e) \sum_a P(a|B, e)P(j|a)P(m|a)$$

$$= \sum_e P(B)P(e)f_1(j, m|B, e)$$

$$= P(B) \sum_e P(e)f_1(j, m|B, e)$$

$$= P(B)f_2(j, m|B)$$

marginal can be obtained from joint by summing out

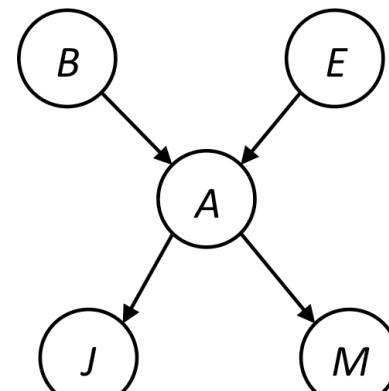
use Bayes' net joint distribution expression

use $x^*(y+z) = xy + xz$

joining on a, and then summing out gives f_1

use $x^*(y+z) = xy + xz$

joining on e, and then summing out gives f_2

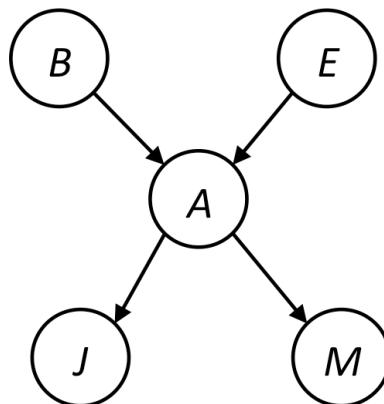


All we are doing is exploiting $uw\bar{y} + uw\bar{z} + u\bar{x}\bar{y} + u\bar{x}\bar{z} + v\bar{w}\bar{y} + v\bar{w}\bar{z} + v\bar{x}\bar{y} + v\bar{x}\bar{z} = (u+v)(w+x)(y+z)$ to improve computational efficiency!

Example (cont'd)

$$P(B|j, m) \propto P(B, j, m)$$

$P(B)$	$P(E)$	$P(A B, E)$	$P(j A)$	$P(m A)$
--------	--------	-------------	----------	----------

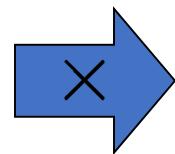


Choose A

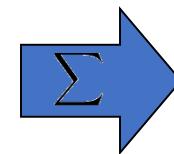
$$P(A|B, E)$$

$$P(j|A)$$

$$P(m|A)$$



$$P(j, m, A|B, E)$$



$$P(j, m|B, E)$$

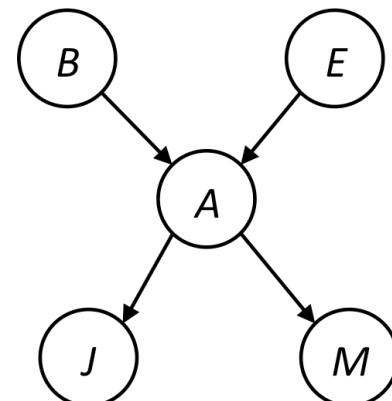
$P(B)$	$P(E)$	$P(j, m B, E)$
--------	--------	----------------

Example (cont'd)

$P(B)$	$P(E)$	$P(j, m B, E)$
--------	--------	----------------

Choose E

$$\begin{array}{ccccc} P(E) & \xrightarrow{\times} & P(j, m, E|B) & \xrightarrow{\sum} & P(j, m|B) \\ P(j, m|B, E) & & & & \end{array}$$



$P(B)$	$P(j, m B)$
--------	-------------

Finish with B

$$\begin{array}{ccccc} P(B) & \xrightarrow{\times} & P(j, m, B) & \xrightarrow{\text{Normalize}} & P(B|j, m) \\ P(j, m|B) & & & & \end{array}$$

Another Variable Elimination Example

Query: $P(X_3|Y_1 = y_1, Y_2 = y_2, Y_3 = y_3)$

Start by inserting evidence, which gives the following initial factors:

$$P(Z), P(X_1|Z), P(X_2|Z), P(X_3|Z), P(y_1|X_1), P(y_2|X_2), P(y_3|X_3)$$

Eliminate X_1 , this introduces the factor $f_1(y_1|Z) = \sum_{x_1} P(x_1|Z)P(y_1|x_1)$,
and we are left with:

$$P(Z), P(X_2|Z), P(X_3|Z), P(y_2|X_2), P(y_3|X_3), f_1(y_1|Z)$$

Eliminate X_2 , this introduces the factor $f_2(y_2|Z) = \sum_{x_2} P(x_2|Z)P(y_2|x_2)$,
and we are left with:

$$P(Z), P(X_3|Z), P(y_3|X_3), f_1(y_1|Z), f_2(y_2|Z)$$

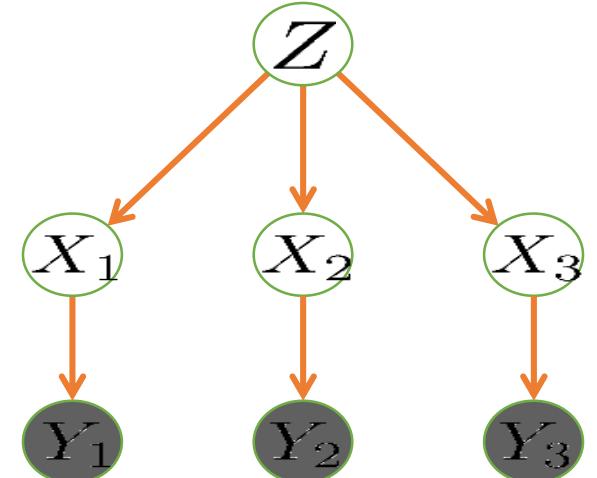
Eliminate Z , this introduces the factor $f_3(y_1, y_2, X_3) = \sum_z P(z)P(X_3|z)f_1(y_1|Z)f_2(y_2|Z)$,
and we are left with:

$$P(y_3|X_3), f_3(y_1, y_2, X_3)$$

No hidden variables left. Join the remaining factors to get:

$$f_4(y_1, y_2, y_3, X_3) = P(y_3|X_3) f_3(y_1, y_2, X_3)$$

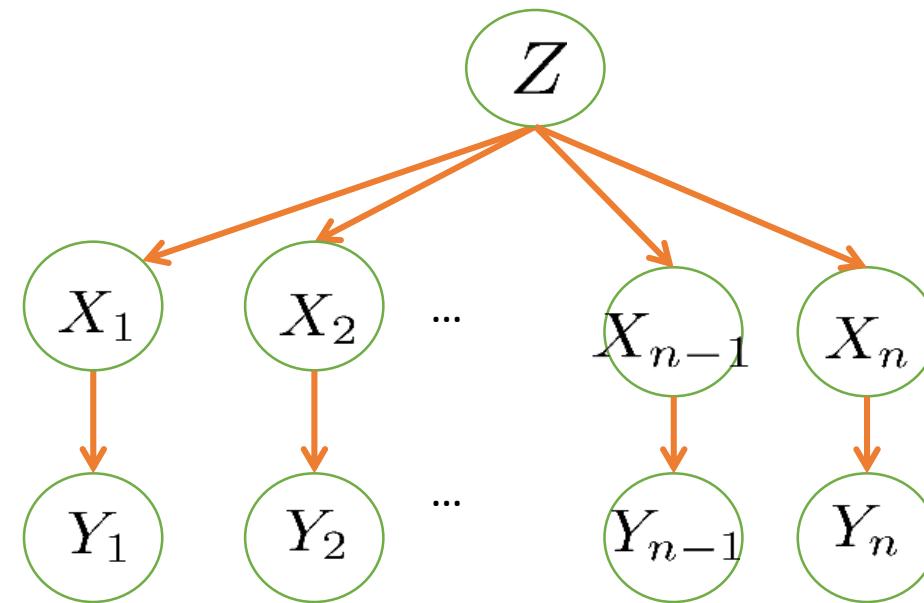
Normalizing over X_3 gives $P(X_3|y_1, y_2, y_3) = f_4(y_1, y_2, y_3, X_3) / \sum_{x_3} f_4(y_1, y_2, y_3, x_3)$



- Computational complexity critically depends on the largest factor being generated in this process
- Size of factor = number of entries in table
- In example above (assuming binary) all factors generated are of size 2 --- as they all only have one variable (Z, Z, and X3 respectively)

Variable Elimination Ordering

- For the query $P(X_n | y_1, \dots, y_n)$ work through the following two different orderings as done in previous slide: Z, X_1, \dots, X_{n-1} and X_1, \dots, X_{n-1}, Z . What is the size of the maximum factor generated for each of the orderings?



- Answer: 2^n versus 2 (assuming binary)
- In general: the ordering can greatly affect efficiency

Detail of size 4

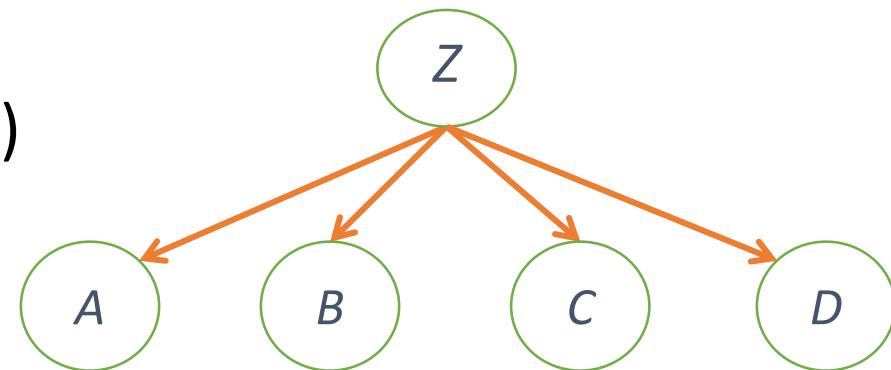
- Elimination order: C, B, A, Z

- $P(D) = \alpha \sum_{z,a,b,c} P(D|z) P(z) P(a|z) P(b|z) P(c|z)$
- $= \alpha \sum_z P(D|z) P(z) \sum_a P(a|z) \sum_b P(b|z) \sum_c P(c|z)$
- Largest factor has 2 variables (D,Z)

- Elimination order: Z, C, B, A

- $P(D) = \alpha \sum_{a,b,c,z} P(a|z) P(b|z) P(c|z) P(D|z) P(z)$
- $= \alpha \sum_a \sum_b \sum_c \sum_z P(a|z) P(b|z) P(c|z) P(D|z) P(z)$
- Largest factor has 4 variables (A,B,C,D)

- In general, with n leaves, factor of size 2^n



VE: Computational and Space Complexity

- The computational and space complexity of variable elimination is determined by the largest factor
- The elimination ordering can greatly affect the size of the largest factor
 - E.g., previous slide's example 2^n vs. 2
- Does there always exist an ordering that only results in small factors?
 - No!

Worst Case Complexity?

- 3-SAT:

$$(x_1 \vee x_2 \vee \neg x_3) \wedge (\neg x_1 \vee x_3 \vee \neg x_4) \wedge (x_2 \vee \neg x_2 \vee x_4) \wedge (\neg x_3 \vee \neg x_4 \vee \neg x_5) \wedge (x_2 \vee x_5 \vee x_7) \wedge (x_4 \vee x_5 \vee x_6) \wedge (\neg x_5 \vee x_6 \vee \neg x_7) \wedge (\neg x_5 \vee \neg x_6 \vee x_7)$$

$$P(X_i = 0) = P(X_i = 1) = 0.5$$

$$Y_1 = X_1 \vee X_2 \vee \neg X_3$$

$$Y_8 = \neg X_5 \vee X_6 \vee X_7$$

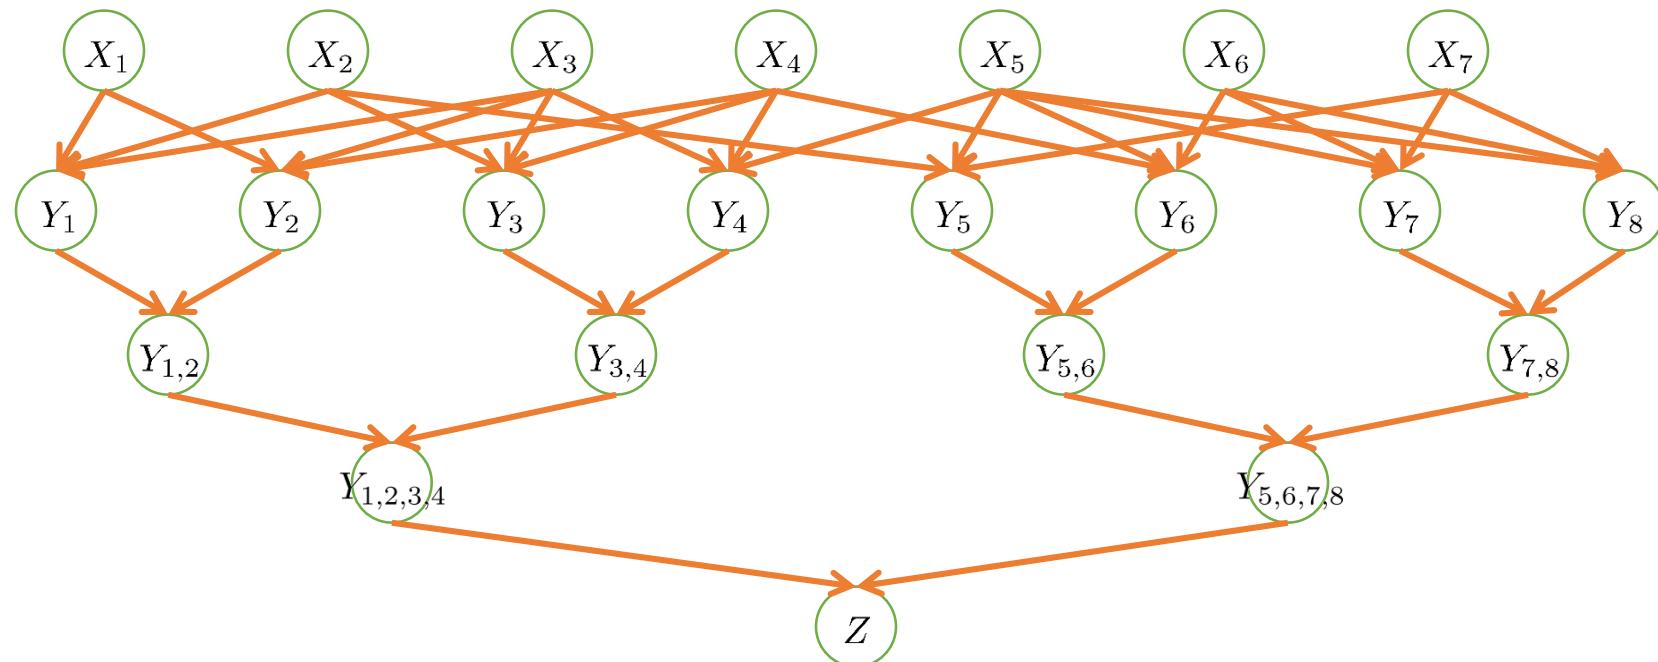
$$Y_{1,2} = Y_1 \wedge Y_2$$

$$Y_{7,8} = Y_7 \wedge Y_8$$

$$Y_{1,2,3,4} = Y_{1,2} \wedge Y_{3,4}$$

$$Y_{5,6,7,8} = Y_{5,6} \wedge Y_{7,8}$$

$$Z = Y_{1,2,3,4} \wedge Y_{5,6,7,8}$$



- If we can answer $P(z)$ equal to zero or not, we answered whether the 3-SAT problem has a solution
- Hence inference in Bayes' nets is NP-hard. No known efficient probabilistic inference in general

“Easy” Structures: Polytrees

- A polytree is a directed graph with no undirected cycles
- For poly-trees you can always find an ordering that is efficient
 - Try it!!
- Cut-set conditioning for Bayes’ net inference
 - Choose set of variables such that if removed only a polytree remains
 - (Exercise) Think about how the specifics would work out!

Sampling

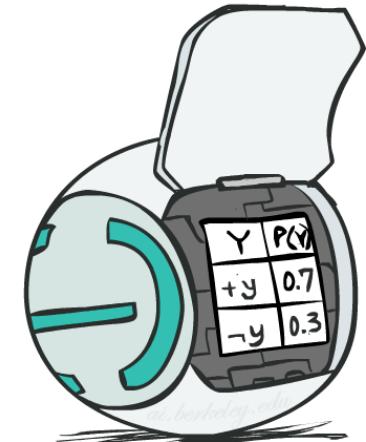
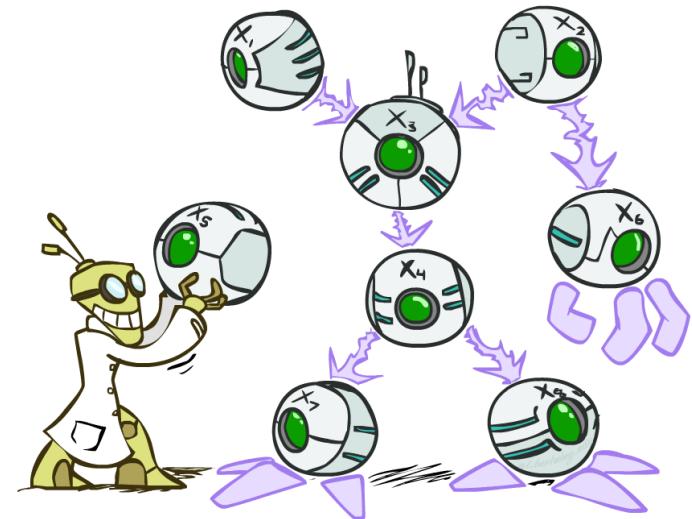
Recall: Bayes' Net Representation

- A directed, acyclic graph, one node per random variable
- A conditional probability table (CPT) for each node
 - A collection of distributions over X, one for each combination of parents' values

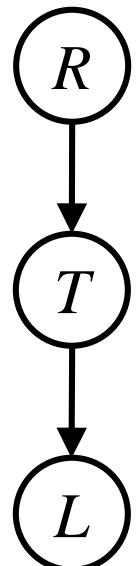
$$P(X|a_1 \dots a_n)$$

- Bayes' nets implicitly encode joint distributions
 - As a product of local conditional distributions
 - To see what probability a BN gives to a full assignment, multiply all the relevant conditionals together:

$$P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i | \text{parents}(X_i))$$



Recap: Bayesian Inference (Exact)



$$P(L) = ?$$

- Inference by Enumeration

$$= \sum_t \sum_r P(L|t) P(r) P(t|r)$$

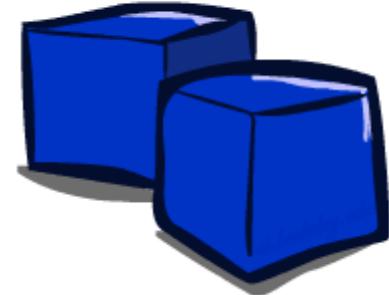
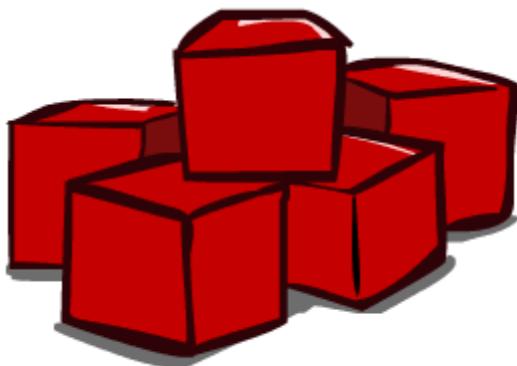
Join on r
Join on t
Eliminate r
Eliminate t

- Variable Elimination

$$= \sum_t P(L|t) \sum_r P(r) P(t|r)$$

Join on r
Eliminate r
Join on t
Eliminate t

Approximate Inference: Sampling

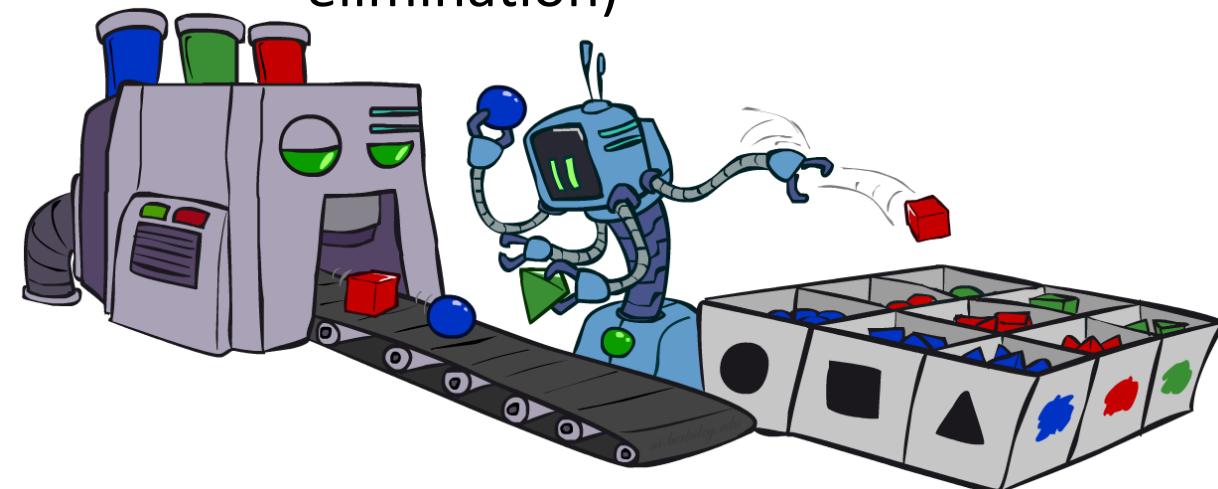


Sampling

- Sampling is a lot like repeated simulation
 - Predicting the weather, basketball games, ...
- Basic idea
 - Draw N samples from a sampling distribution S
 - Compute an approximate posterior probability
 - Show this converges to the true probability P

- Why sample?

- Learning: get samples from a distribution you don't know
- Inference: getting a sample is faster than computing the right answer (e.g. with variable elimination)



Sampling 2

- Sampling from given distribution
 - Step 1: Get sample u from uniform distribution over $[0, 1]$
 - E.g. `random()` in python
 - Step 2: Convert this sample u into an outcome for the given distribution by having each target outcome associated with a sub-interval of $[0,1]$ with sub-interval size equal to probability of the outcome

- Example

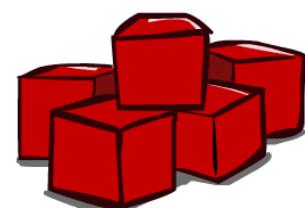
C	P(C)
red	0.6
green	0.1
blue	0.3

$0 \leq u < 0.6, \rightarrow C = \text{red}$

$0.6 \leq u < 0.7, \rightarrow C = \text{green}$

$0.7 \leq u < 1, \rightarrow C = \text{blue}$

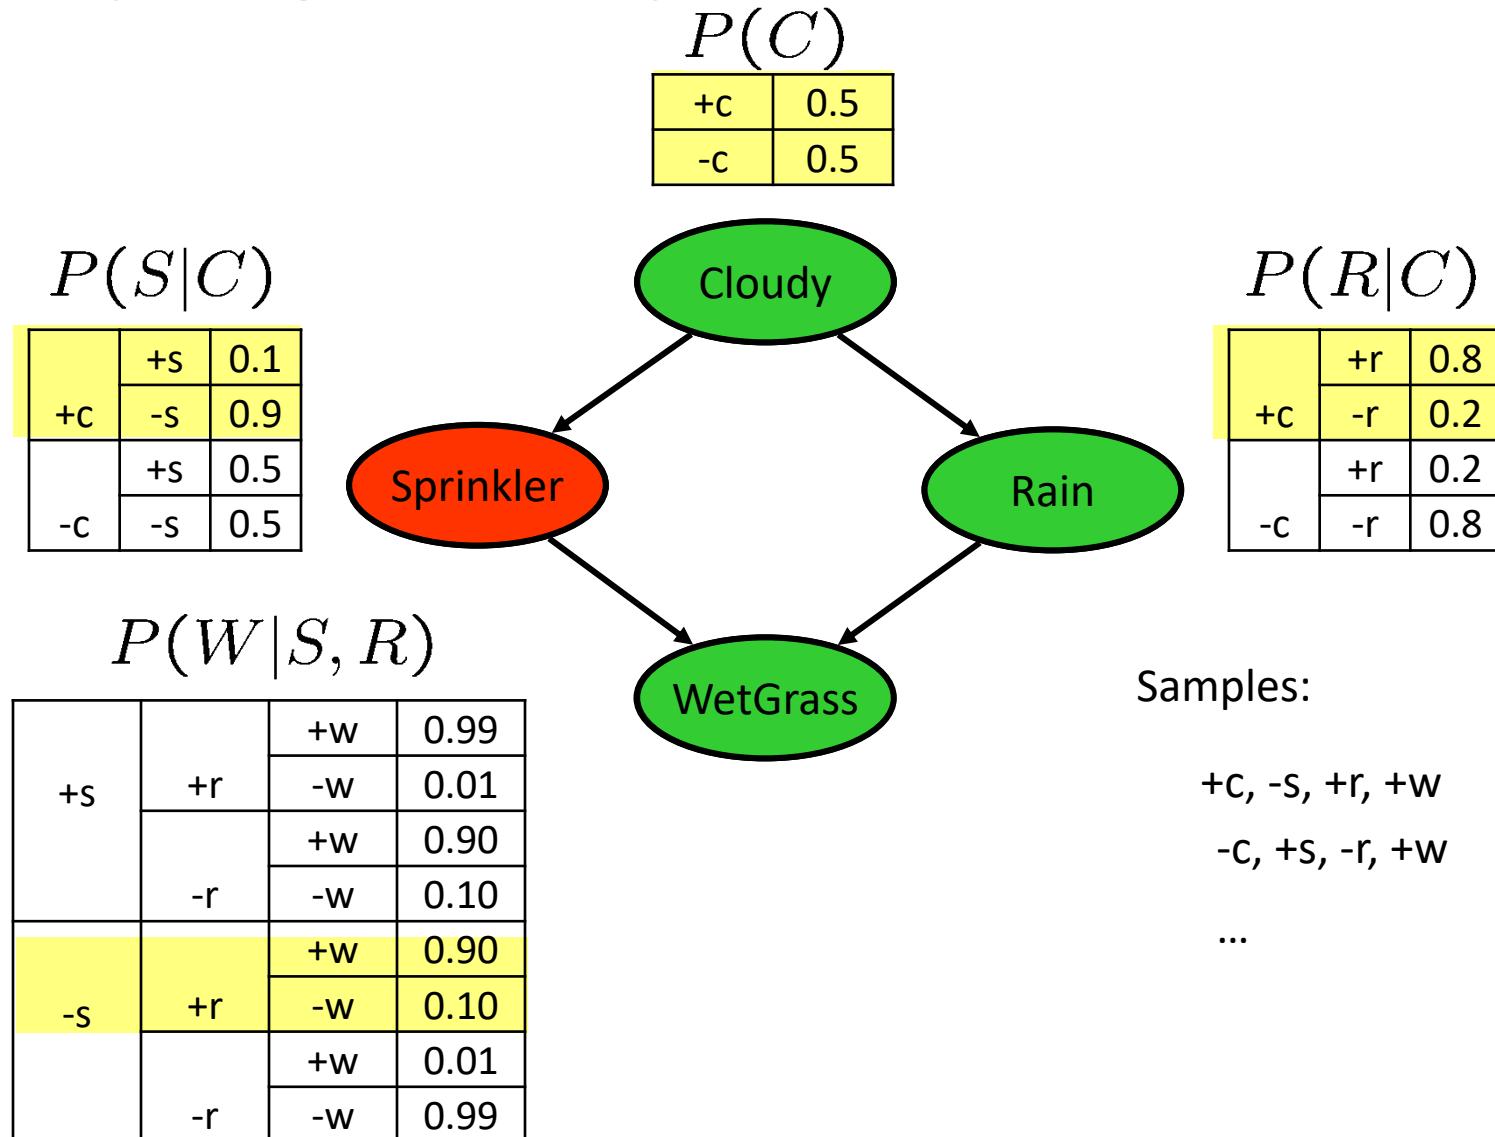
- If `random()` returns $u = 0.83$, then our sample is $C = \text{blue}$
- E.g, after sampling 8 times:



Sampling in Bayes' Nets

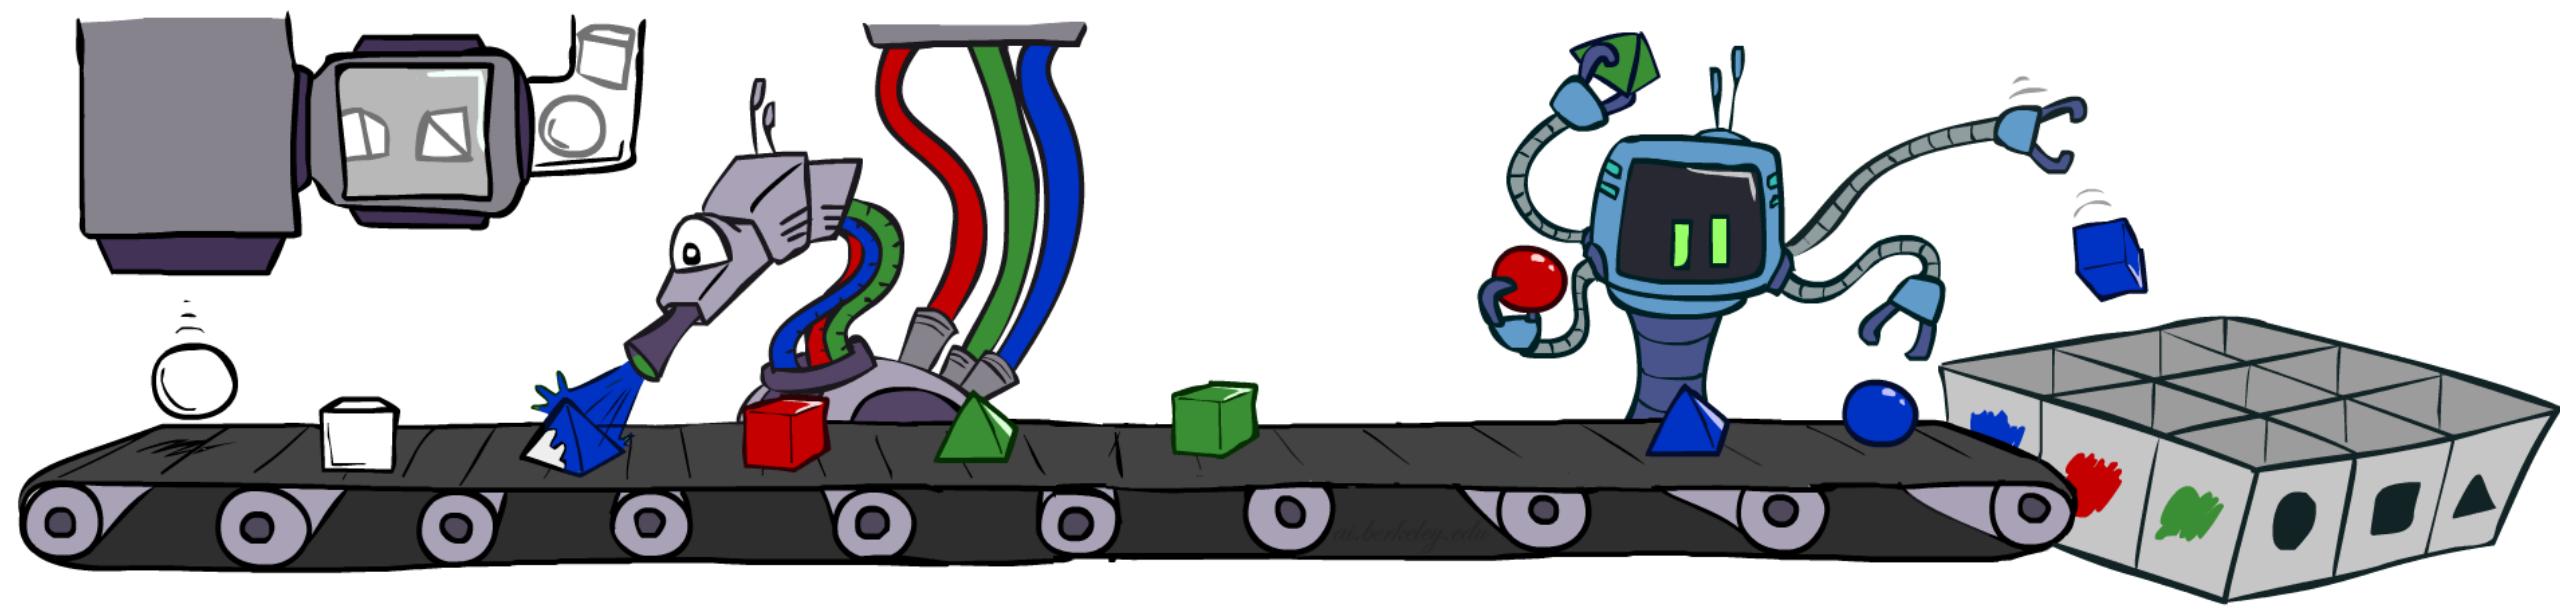
- Prior Sampling
- Rejection Sampling
- Likelihood Weighting
- Gibbs Sampling

Prior Sampling: Example



Prior Sampling: Algorithm

- For $i = 1, 2, \dots, n$ in topological order
 - Sample x_i from $P(X_i \mid \text{Parents}(X_i))$
- Return (x_1, x_2, \dots, x_n)



Prior Sampling

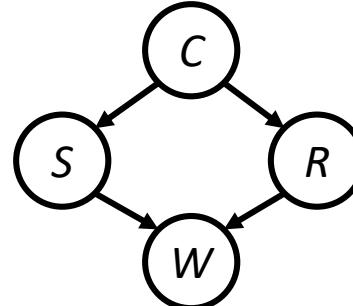
- This process generates samples with probability:

$$S_{PS}(x_1 \dots x_n) = \prod_{i=1}^n P(x_i | \text{Parents}(X_i)) = P(x_1 \dots x_n)$$

- ...i.e. the BN's joint probability
- Let the number of samples of an event be $N_{PS}(x_1 \dots x_n)$
- Then $\lim_{N \rightarrow \infty} \hat{P}(x_1, \dots, x_n) = \lim_{N \rightarrow \infty} N_{PS}(x_1, \dots, x_n)/N$
 $= S_{PS}(x_1, \dots, x_n)$
 $= P(x_1 \dots x_n)$
- i.e., the sampling procedure is consistent

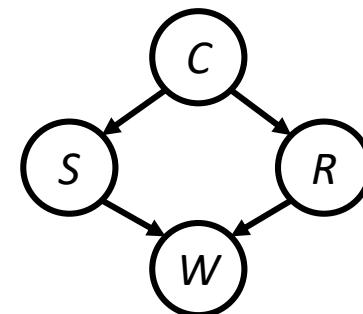
Example

- We'll get a bunch of samples from the BN:
 - $+c, -s, +r, +w$
 - $+c, +s, +r, +w$
 - $-c, +s, +r, -w$
 - $+c, -s, +r, +w$
 - $-c, -s, -r, +w$
- If we want to know $P(W)$
 - We have counts $\langle +w:4, -w:1 \rangle$
 - Normalize to get $P(W) = \langle +w:0.8, -w:0.2 \rangle$
 - This will get closer to the true distribution with more samples
 - Can estimate anything else, too
 - $P(C | +w)? P(C | +r, +w)?$
 - Can also use this to estimate expected value of $f(X)$ - Monte Carlo Estimation
 - What about $P(C | -r, -w)?$



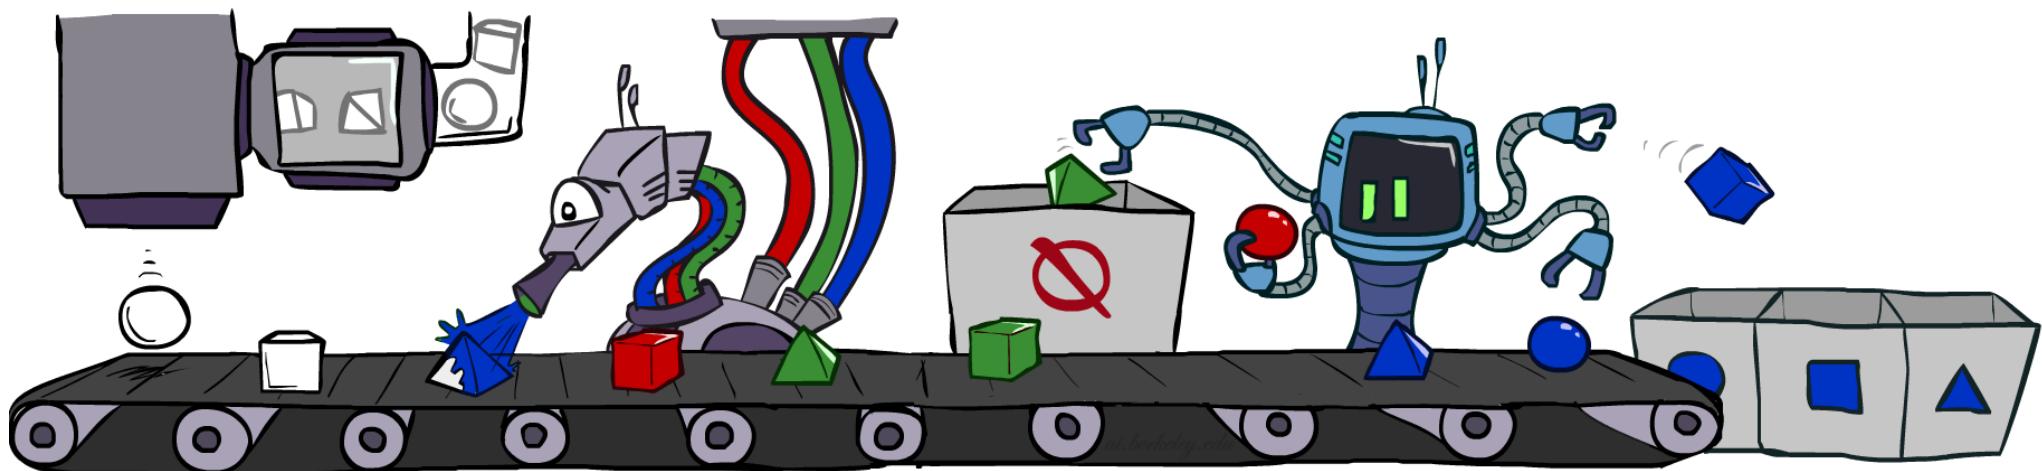
Rejection Sampling

- Let's say we want $P(C)$
 - Just tally counts of C as we go
- Let's say we want $P(C \mid +s)$
 - Same thing: tally C outcomes, but ignore (reject) samples which don't have $S=+s$
 - This is called rejection sampling
 - We can toss out samples early!
 - It is also consistent for conditional probabilities (i.e., correct in the limit)



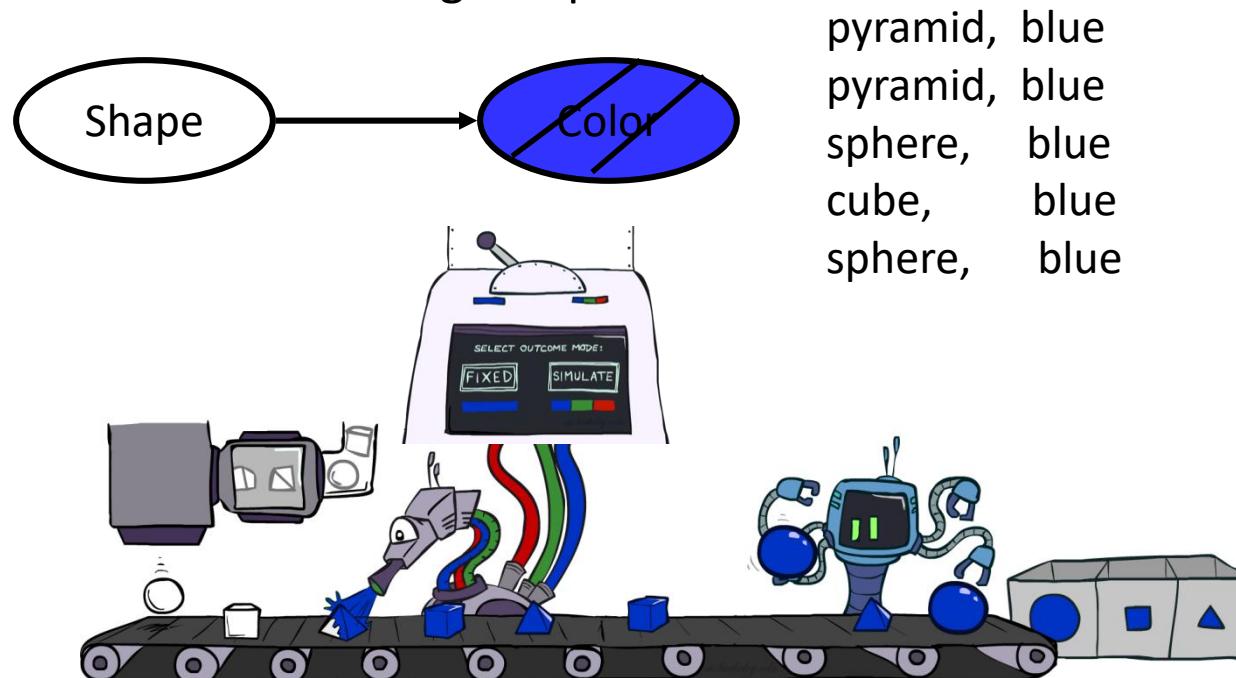
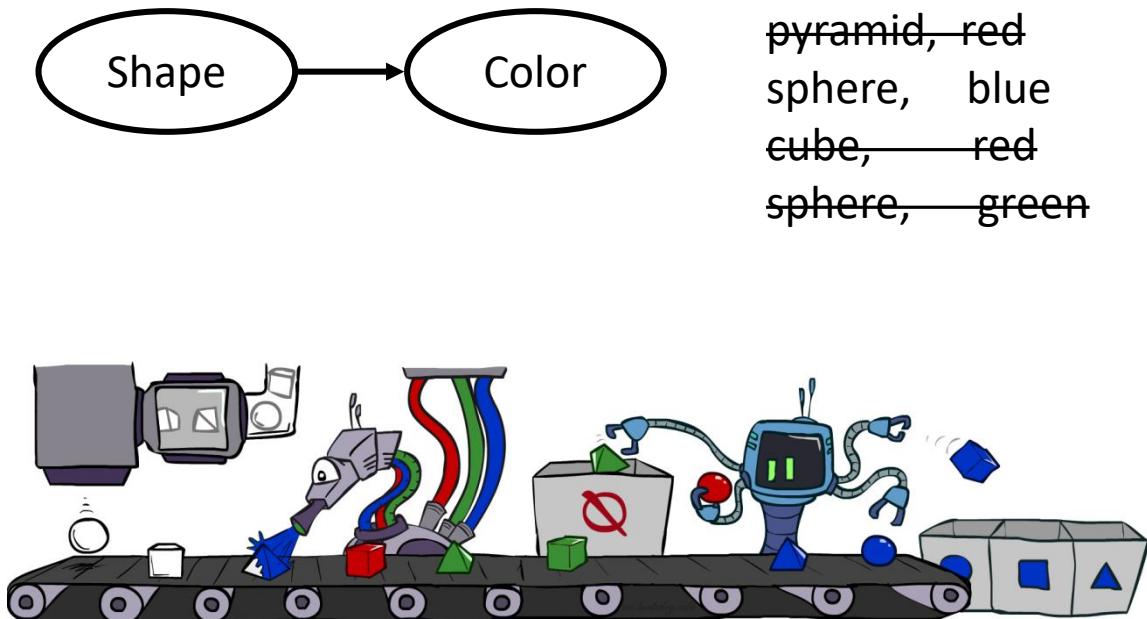
Rejection Sampling: Algorithm

- Input: evidence instantiation
- For $i = 1, 2, \dots, n$ in topological order
 - Sample x_i from $P(X_i \mid \text{Parents}(X_i))$
 - If x_i not consistent with evidence
 - Reject: return – no sample is generated in this cycle
- Return (x_1, x_2, \dots, x_n)

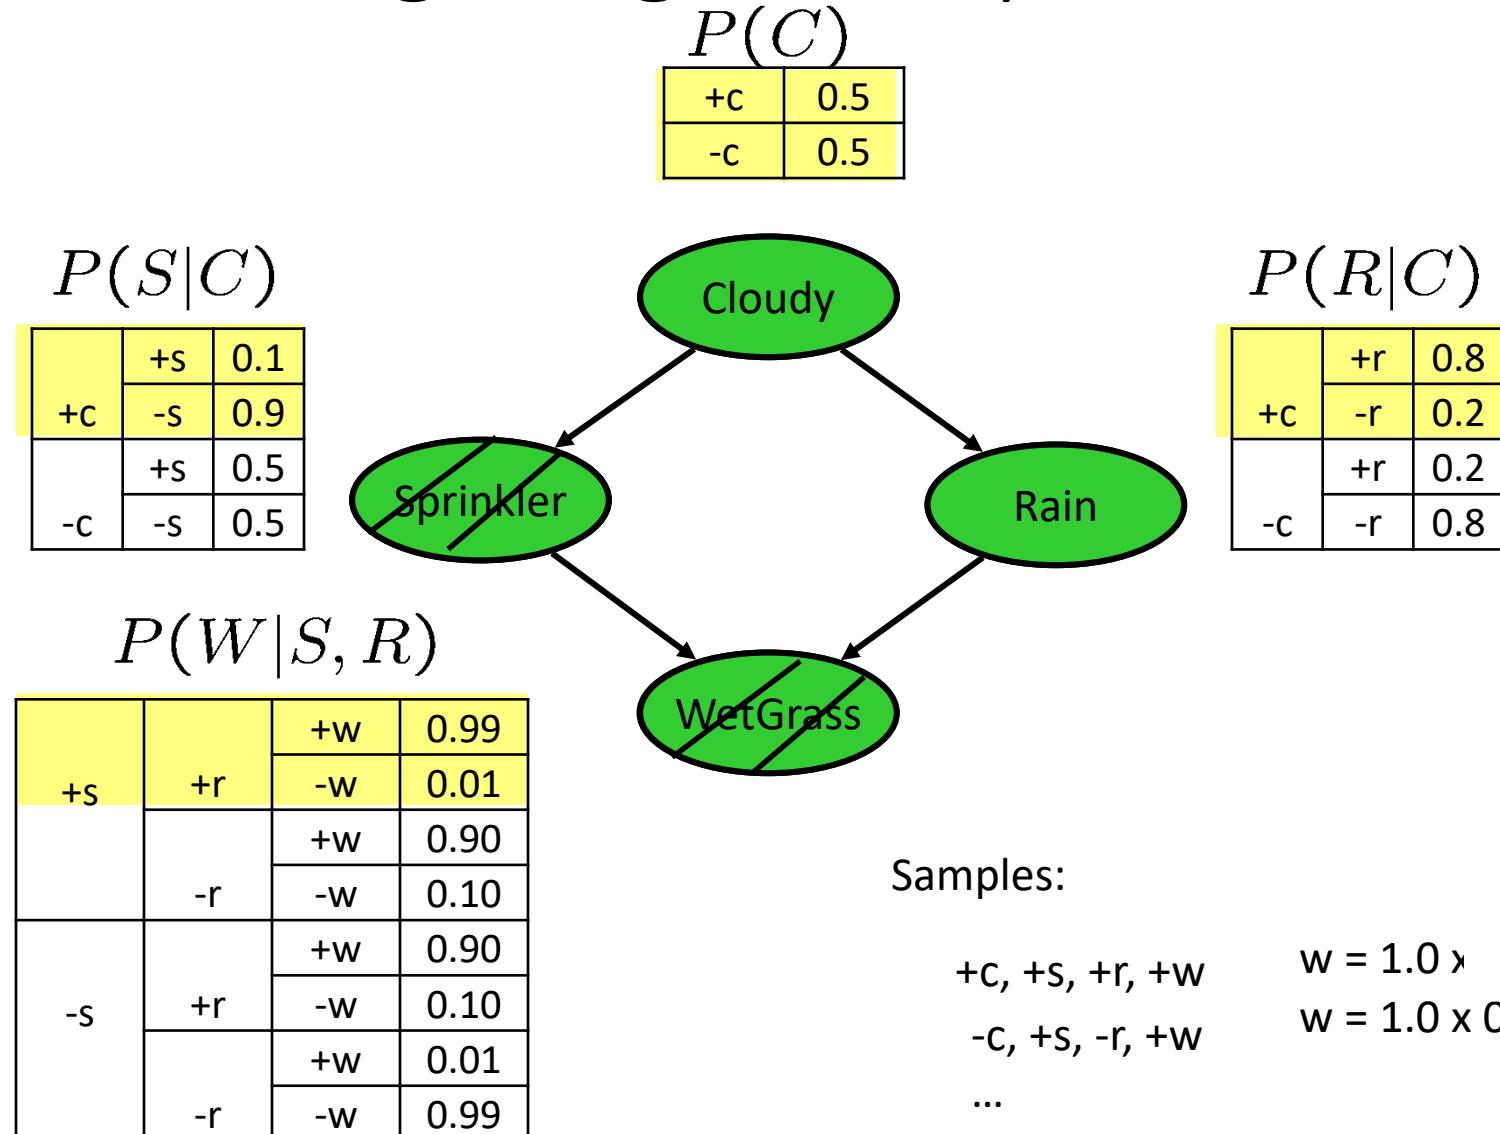


Likelihood Weighting

- Problem with rejection sampling:
 - If evidence is unlikely, rejects lots of samples
 - Consider $P(\text{Shape} \mid \text{blue})$
- Idea: fix evidence variables and sample the rest
 - Problem: sample distribution not consistent!
 - Solution: weight by probability of evidence given parents

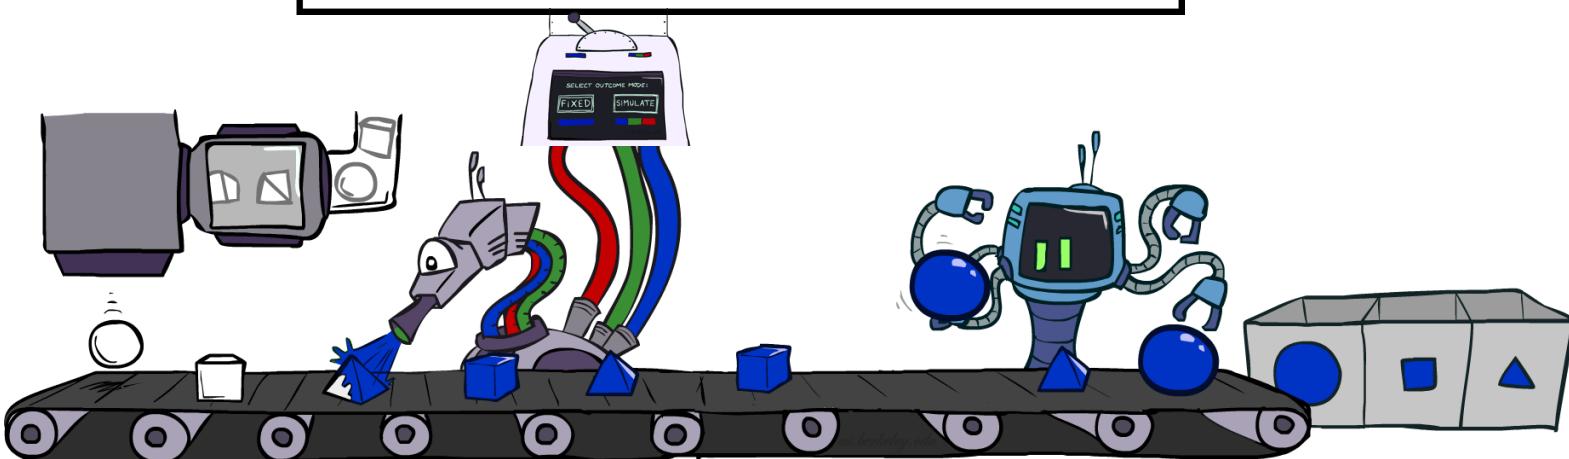


Likelihood Weighting: Example



Likelihood Weighting: Algorithm

- Input: evidence instantiation
- $w = 1.0$
- for $i = 1, 2, \dots, n$ in topological order
 - if X_i is an evidence variable
 - $X_i = \text{observation } x_i \text{ for } X_i$
 - Set $w = w * P(x_i | \text{Parents}(X_i))$
 - else
 - Sample x_i from $P(X_i | \text{Parents}(X_i))$
- return $(x_1, x_2, \dots, x_n), w$



Likelihood Weighting

- Sampling distribution if \mathbf{z} sampled and \mathbf{e} fixed evidence

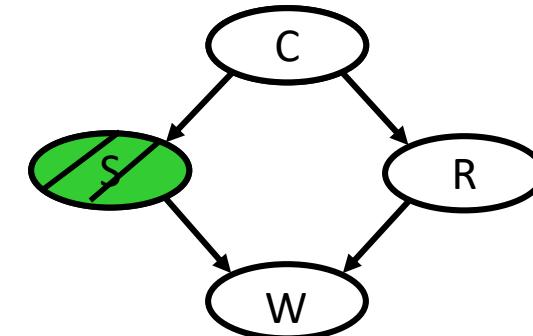
$$S_{WS}(\mathbf{z}, \mathbf{e}) = \prod_{i=1}^l P(z_i | \text{Parents}(Z_i))$$

- Now, samples have weights

$$w(\mathbf{z}, \mathbf{e}) = \prod_{i=1}^m P(e_i | \text{Parents}(E_i))$$

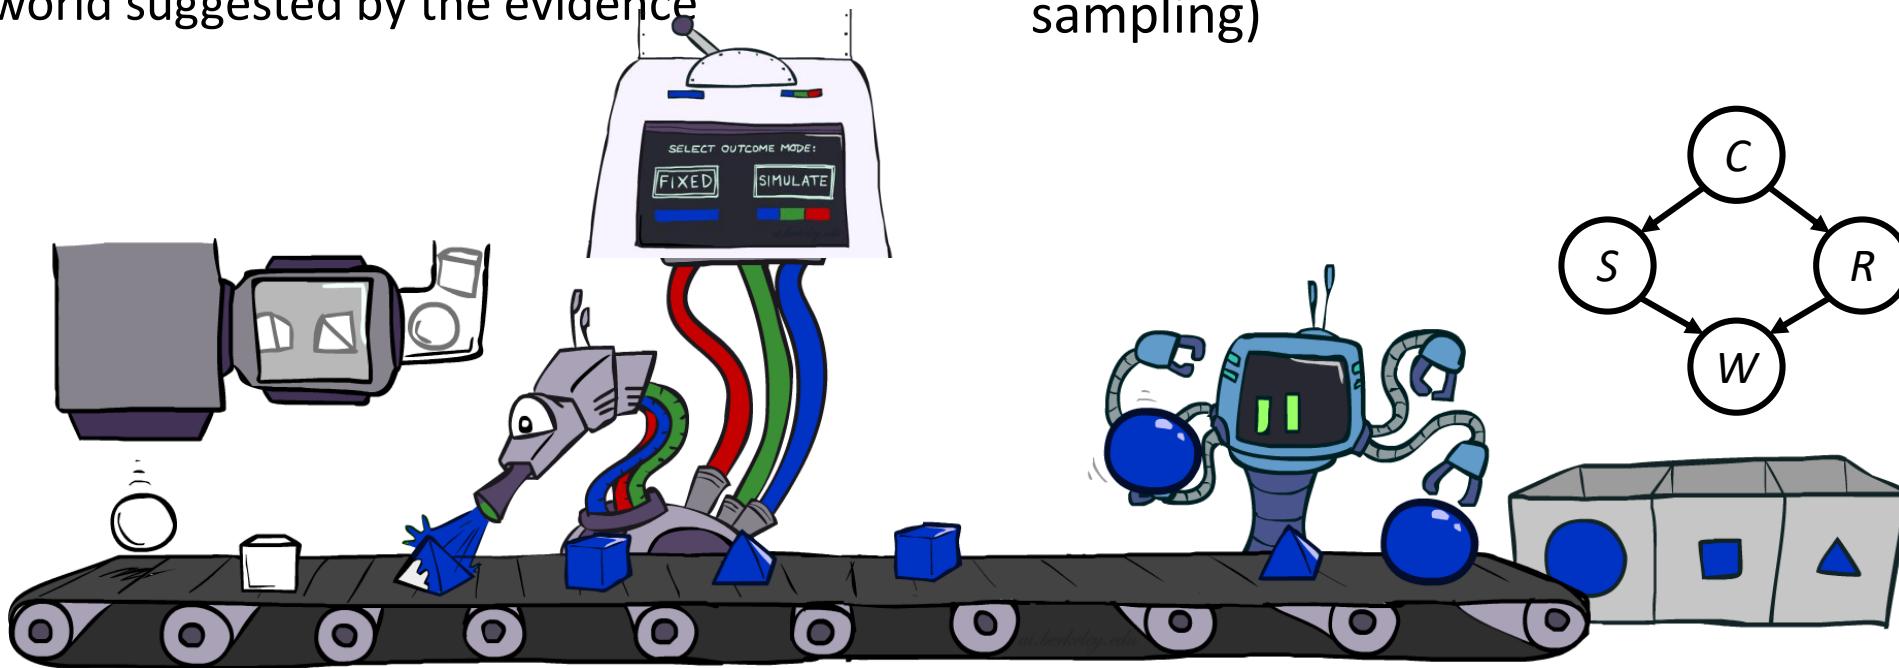
- Together, weighted sampling distribution is consistent

$$\begin{aligned} S_{WS}(\mathbf{z}, \mathbf{e}) \cdot w(\mathbf{z}, \mathbf{e}) &= \prod_{i=1}^l P(z_i | \text{Parents}(z_i)) \prod_{i=1}^m P(e_i | \text{Parents}(e_i)) \\ &= P(\mathbf{z}, \mathbf{e}) \end{aligned}$$



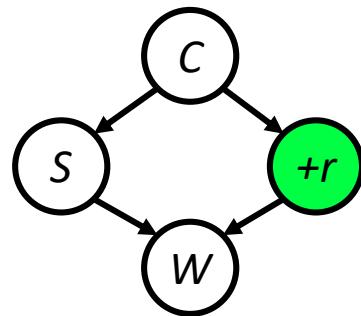
Likelihood Weighting

- Likelihood weighting is helpful
 - We have taken evidence into account as we generate the sample
 - E.g. here, W's value will get picked based on the evidence values of S, R
 - More of our samples will reflect the state of the world suggested by the evidence
- Likelihood weighting doesn't solve all our problems
 - Evidence influences the choice of downstream variables, but not upstream ones (C isn't more likely to get a value matching the evidence)
 - We would like to consider evidence when we sample every variable (leads to Gibbs sampling)

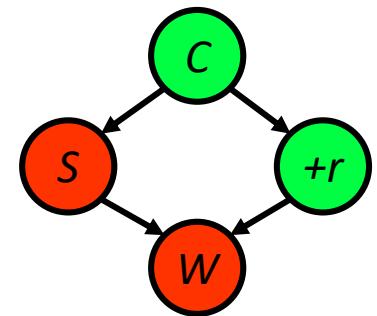


Gibbs Sampling: Example $P(S | +r)$

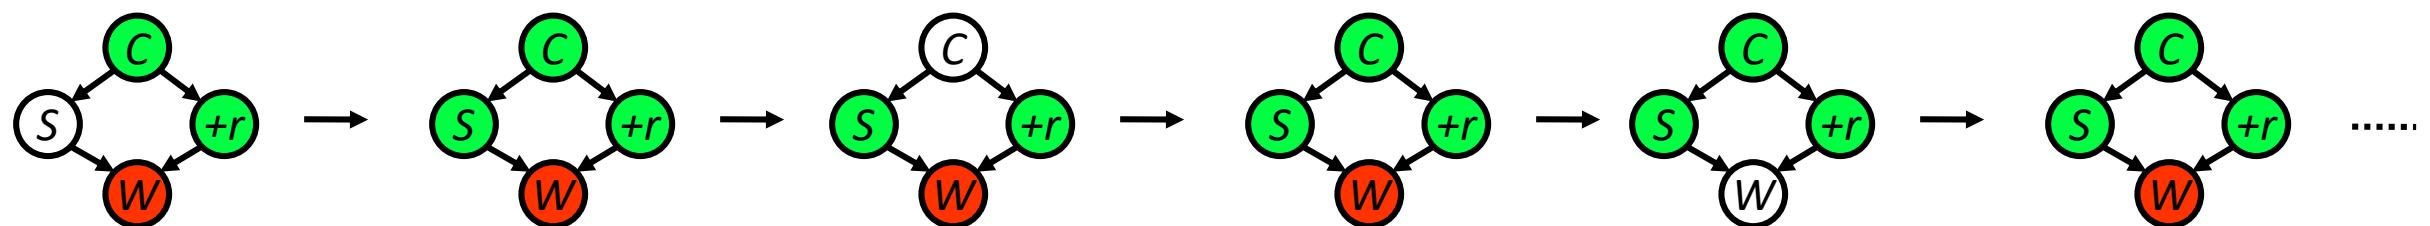
- Step 1: Fix evidence
 - $R = +r$



- Step 2: Initialize other variables
 - Randomly



- Steps 3: Repeat
 - Choose a non-evidence variable X
 - Resample X from $P(X | \text{all other variables})^*$



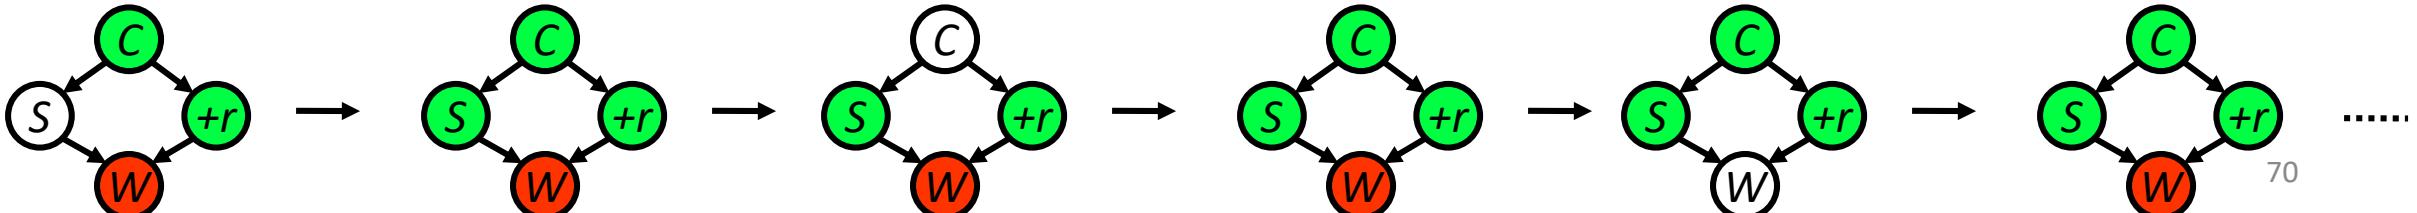
Sample from $P(S | +c, -w, +r)$

Sample from $P(C | +s, -w, +r)$

Sample from $P(W | +s, +c, +r)$

Gibbs Sampling

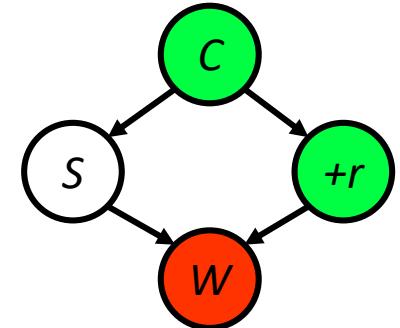
- **Procedure**
 - Keep track of a full instantiation x_1, \dots, x_n
 - Start with an arbitrary instantiation consistent with the evidence
 - Sample one variable at a time, conditioned on all the rest, but keep evidence fixed
 - Keep repeating this for a long time
- **Property**
 - In the limit of repeating this infinitely many times the resulting samples come from the correct distribution (i.e. conditioned on evidence)
- **Rationale**
 - Both upstream and downstream variables condition on evidence
- **In contrast:**
 - Likelihood weighting only conditions on upstream evidence, and hence weights obtained in likelihood weighting can sometimes be very small
 - Sum of weights over all samples is indicative of how many “effective” samples were obtained, so we want high weight



Resampling of One Variable

- Sample from $P(S | +c, +r, -w)$

$$\begin{aligned} P(S | +c, +r, -w) &= \frac{P(S, +c, +r, -w)}{P(+c, +r, -w)} \\ &= \frac{P(S, +c, +r, -w)}{\sum_s P(s, +c, +r, -w)} \\ &= \frac{P(+c)P(S | +c)P(+r | +c)P(-w | S, +r)}{\sum_s P(+c)P(s | +c)P(+r | +c)P(-w | s, +r)} \\ &= \frac{P(+c)P(S | +c)P(+r | +c)P(-w | S, +r)}{P(+c)P(+r | +c) \sum_s P(s | +c)P(-w | s, +r)} \\ &= \frac{P(S | +c)P(-w | S, +r)}{\sum_s P(s | +c)P(-w | s, +r)} \end{aligned}$$



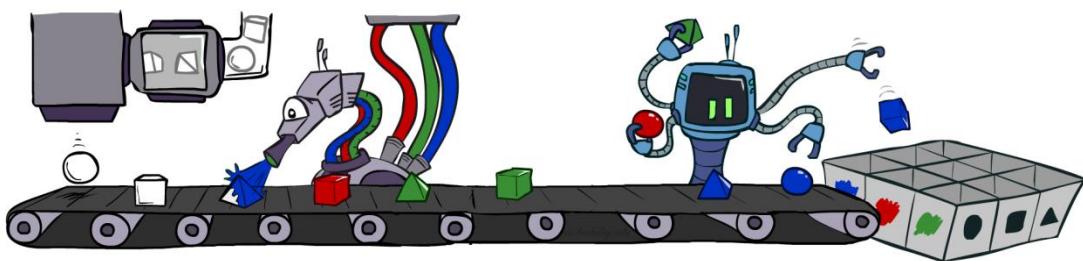
- Many things cancel out – only CPTs with S remain!
- More generally: only CPTs that have resampled variable need to be considered, and joined together

More Details on Gibbs Sampling*

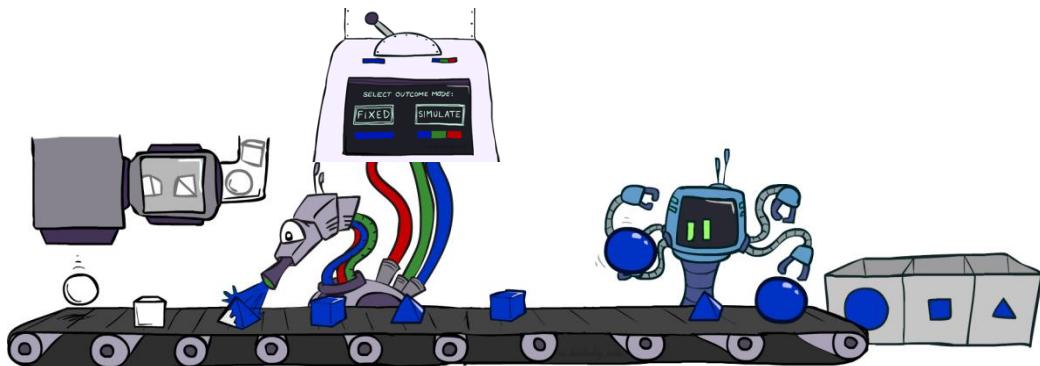
- Gibbs sampling belongs to a family of sampling methods called Markov chain Monte Carlo (MCMC)
 - Specifically, it is a special case of a subset of MCMC methods called Metropolis-Hastings
- You can read more about this here:
 - <https://ermongroup.github.io/cs228-notes/inference/sampling/>

Bayes' Net Sampling Summary

- Prior Sampling $P(Q)$



- Likelihood Weighting $P(Q|e)$



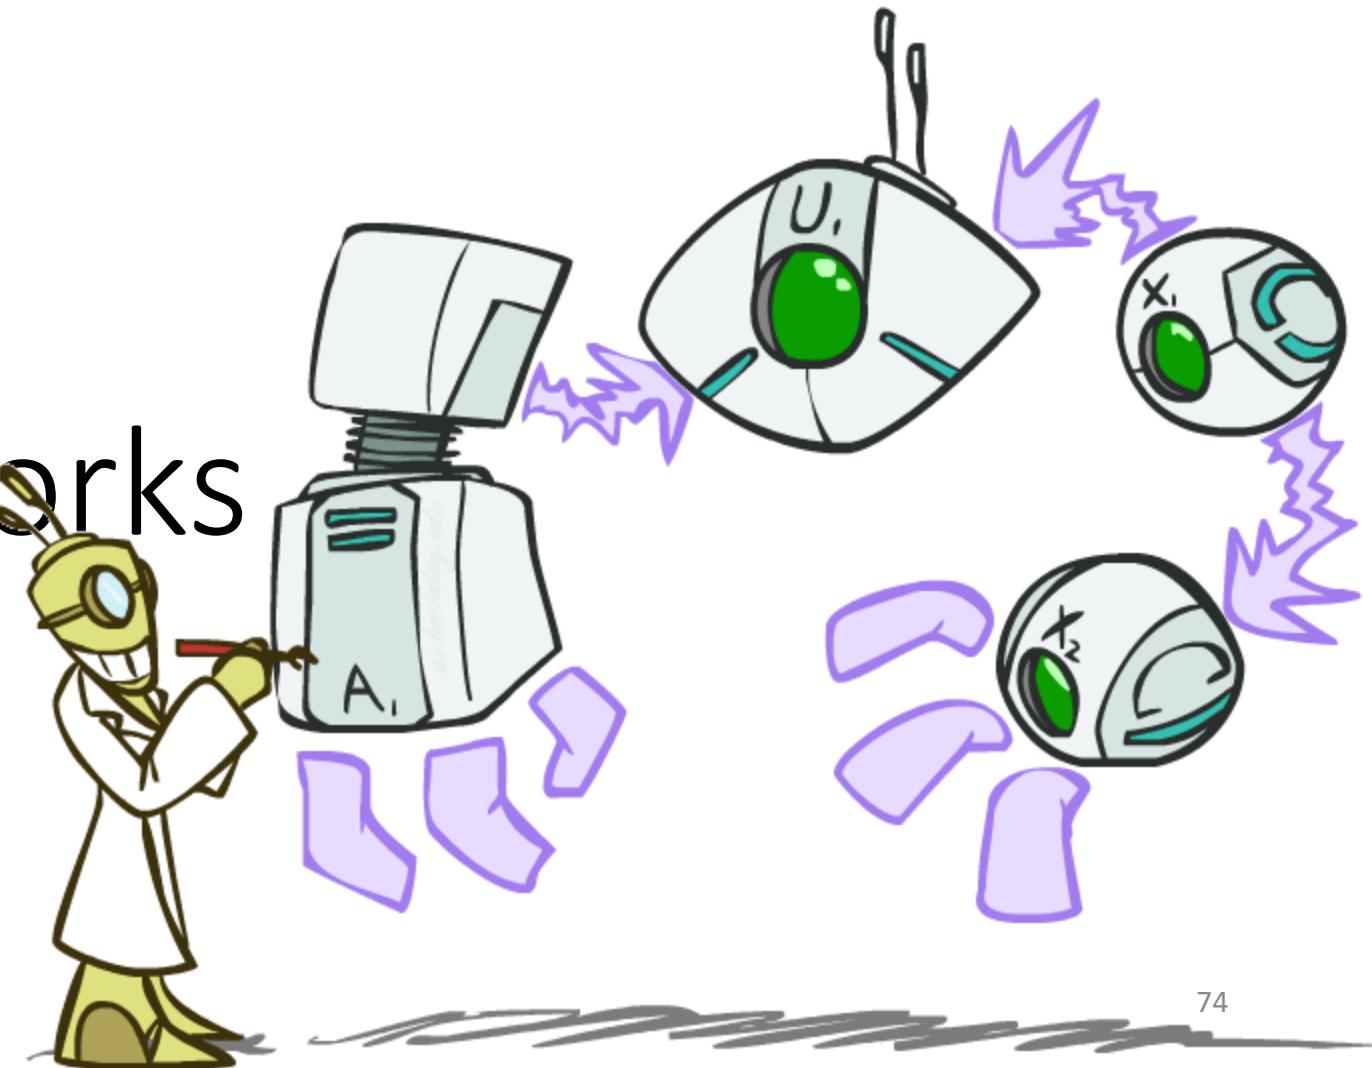
- Rejection Sampling $P(Q|e)$



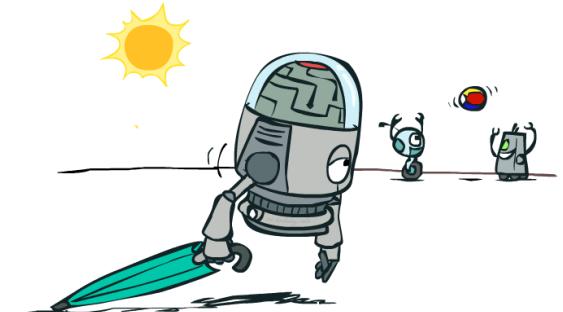
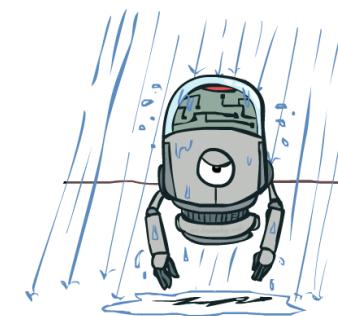
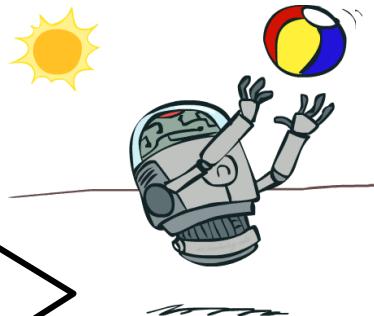
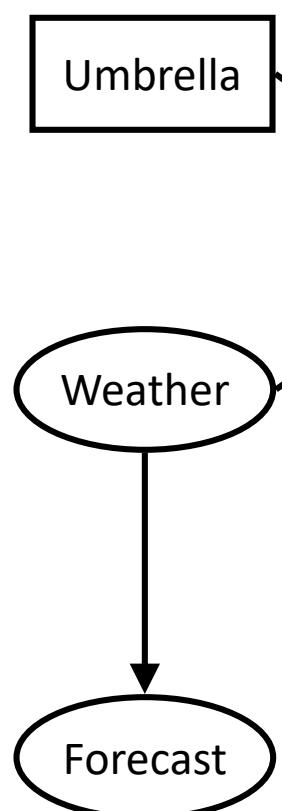
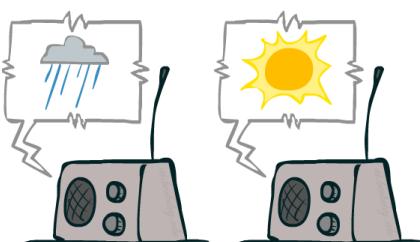
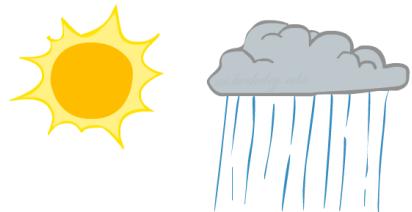
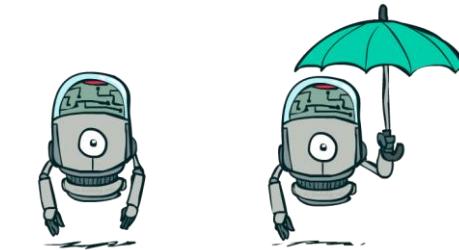
- Gibbs Sampling $P(Q|e)$



Decision Networks



Decision Networks



Decision Networks 2

- **MEU: choose the action which maximizes the expected utility given the evidence**

- Can directly operationalize this with decision networks

- Bayes nets with nodes for utility and actions
 - Lets us calculate the expected utility for each action

- New node types:



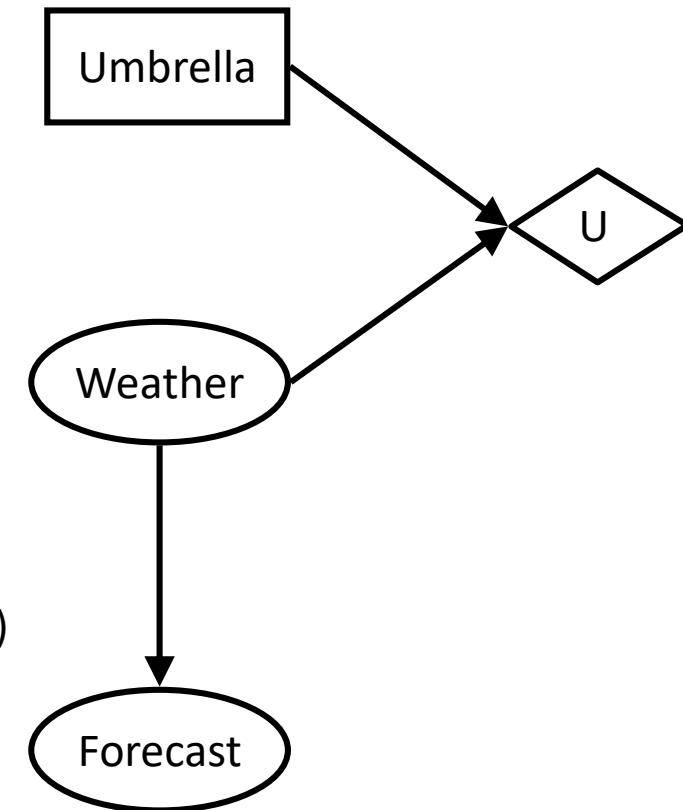
- Chance nodes (just like BNs)



- Actions (rectangles, cannot have parents, act as observed evidence)

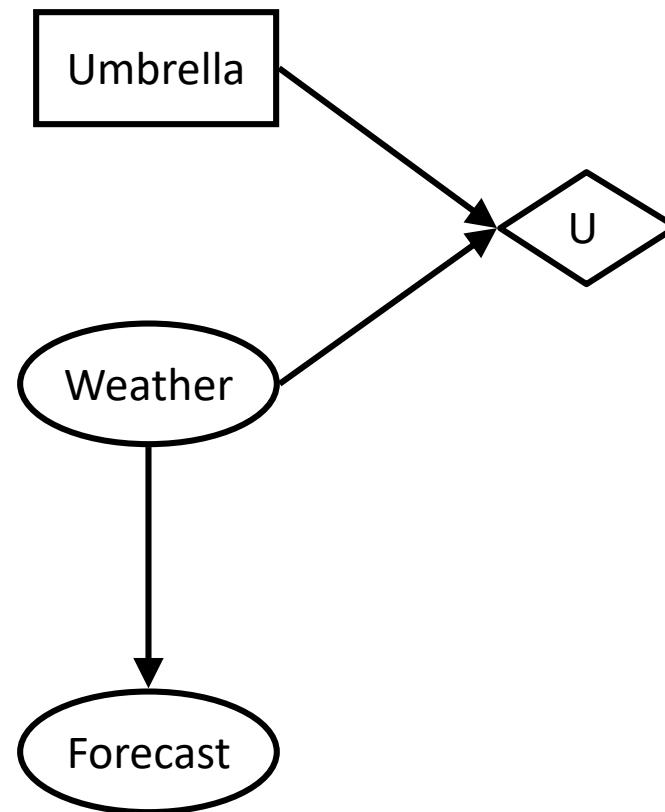


- Utility node (diamond, depends on action and chance nodes)



Decision Networks 3

- Action selection
 - Instantiate all evidence
 - Set action node(s) each possible way
 - Calculate posterior for all parents of utility node, given the evidence
 - Calculate expected utility for each action
 - Choose maximizing action



Maximum Expected Utility

Umbrella = leave

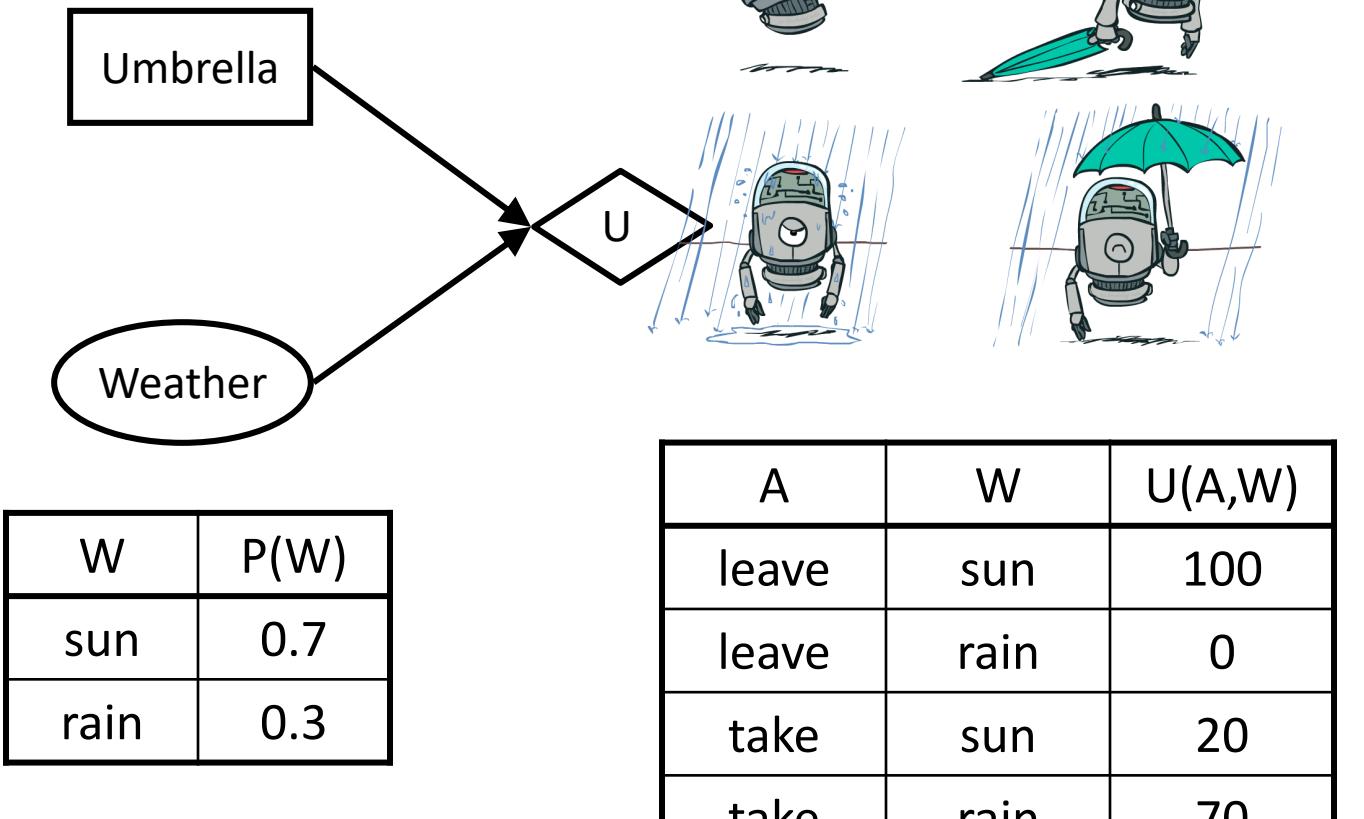
$$\begin{aligned} \text{EU(leave)} &= \sum_w P(w)U(\text{leave}, w) \\ &= 0.7 \cdot 100 + 0.3 \cdot 0 = 70 \end{aligned}$$

Umbrella = take

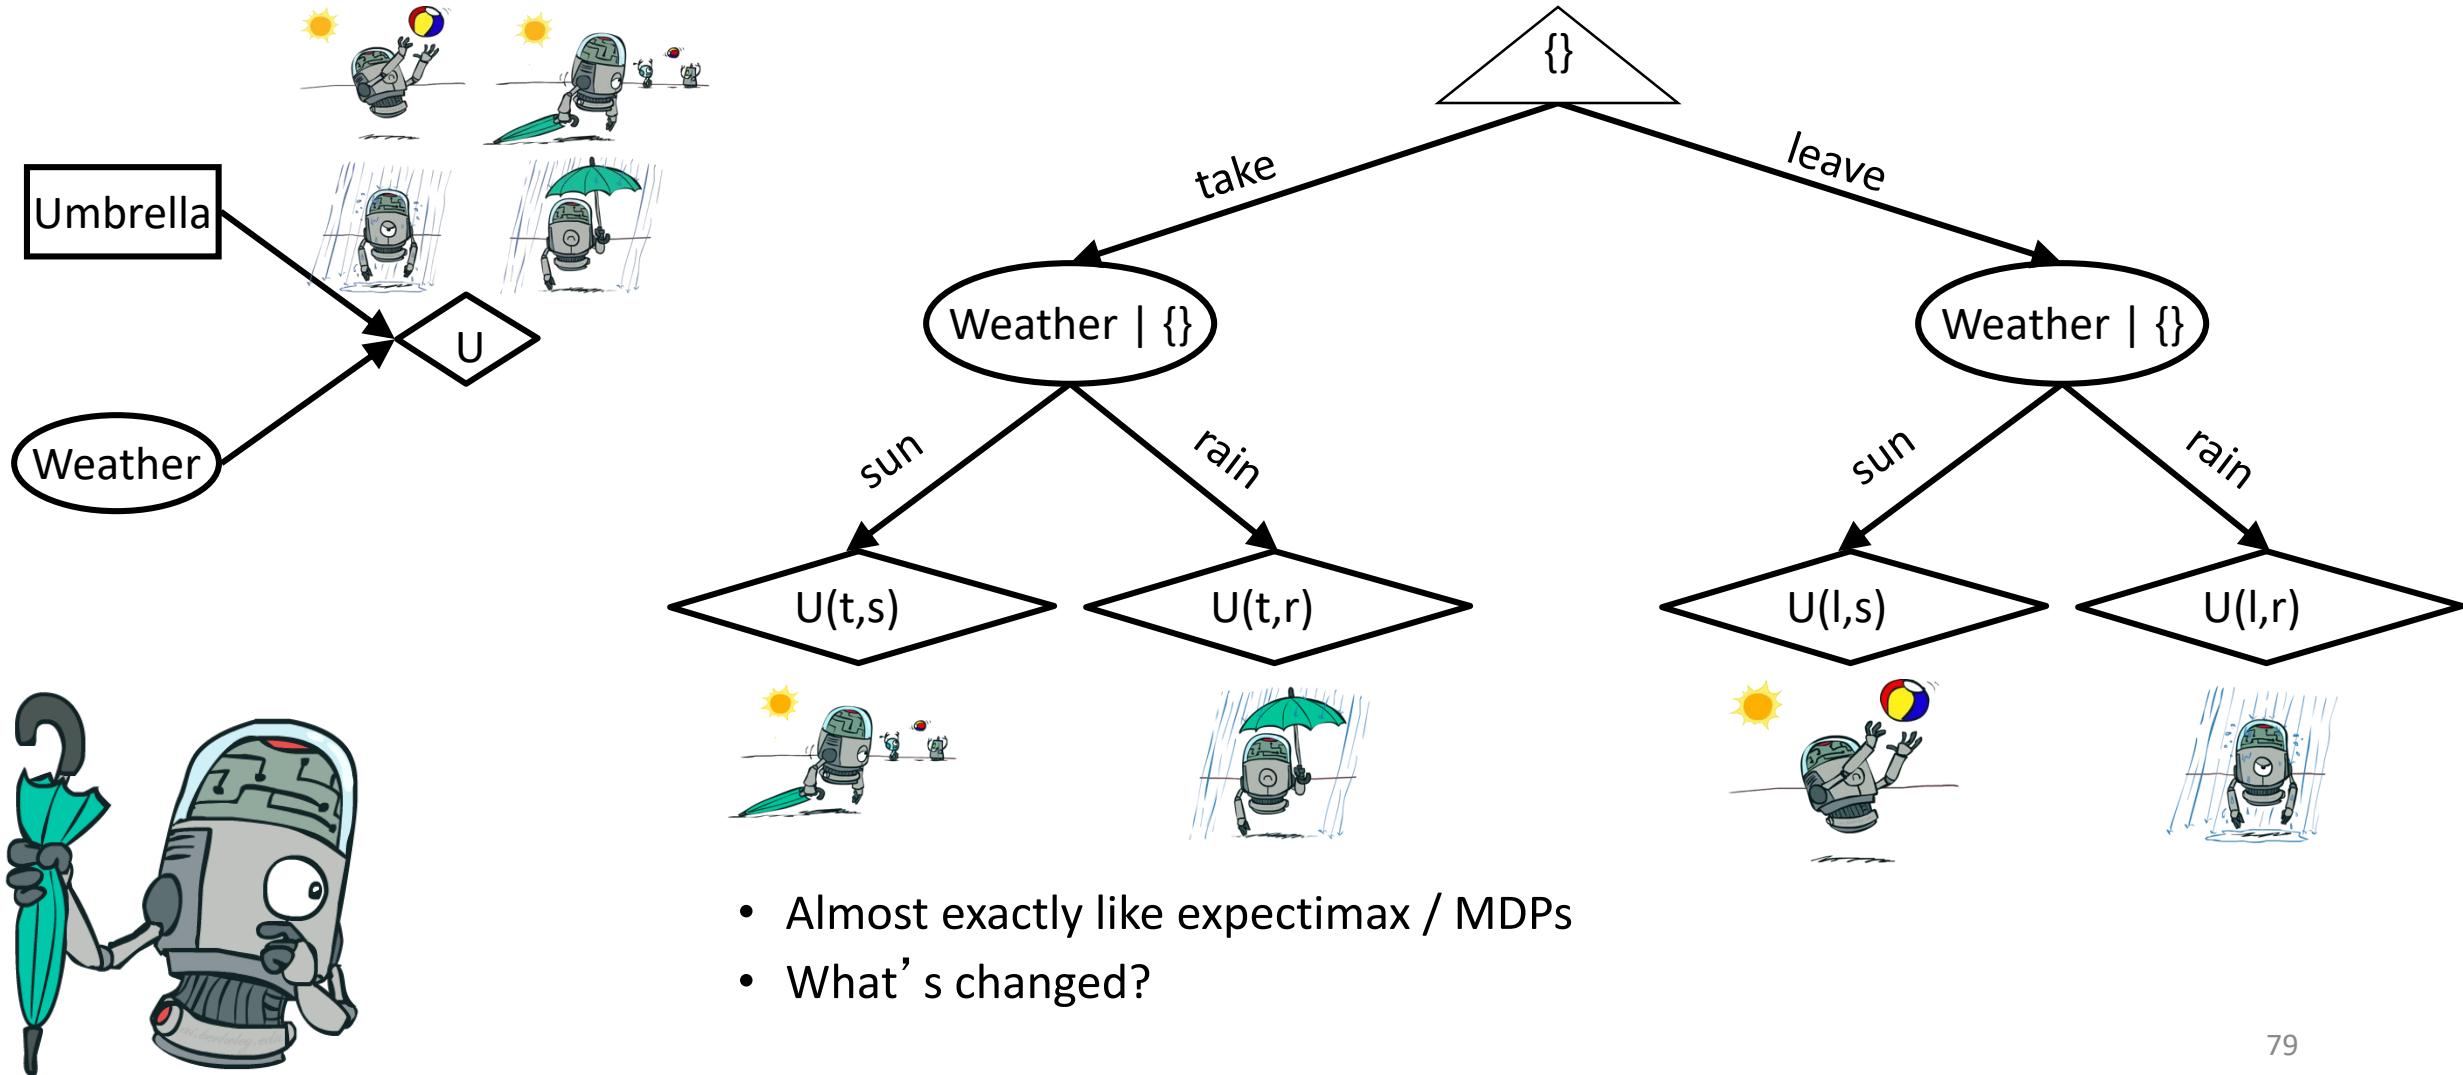
$$\begin{aligned} \text{EU(take)} &= \sum_w P(w)U(\text{take}, w) \\ &= 0.7 \cdot 20 + 0.3 \cdot 70 = 35 \end{aligned}$$

Optimal decision = leave

$$\text{MEU}(\emptyset) = \max_a \text{EU}(a) = 70$$



Decisions as Outcome Trees



Maximum Expected Utility

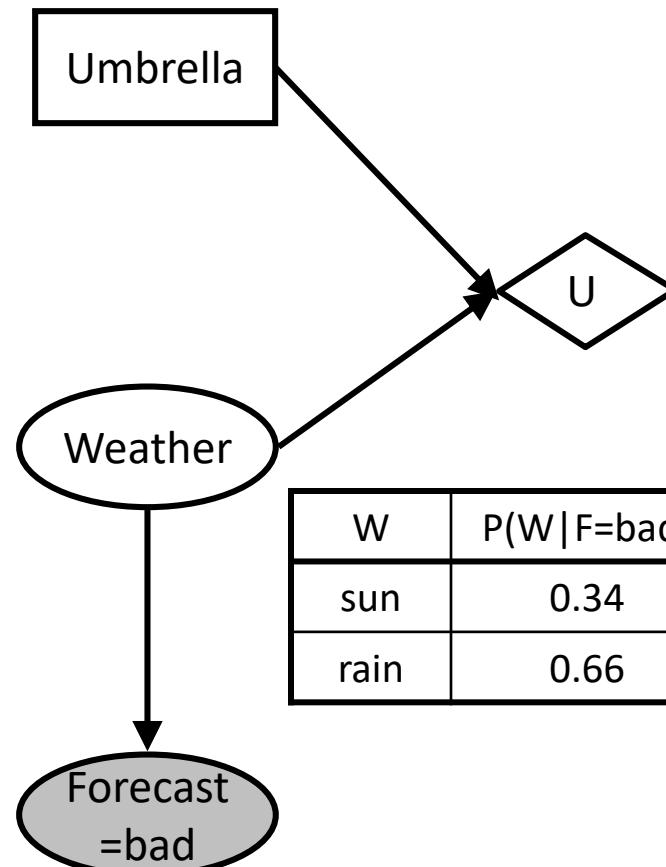
Umbrella = leave

$$EU(\text{leave}|\text{bad}) = \sum_w P(w|\text{bad})U(\text{leave}, w)$$

$$P(W) \quad P(F|W)$$

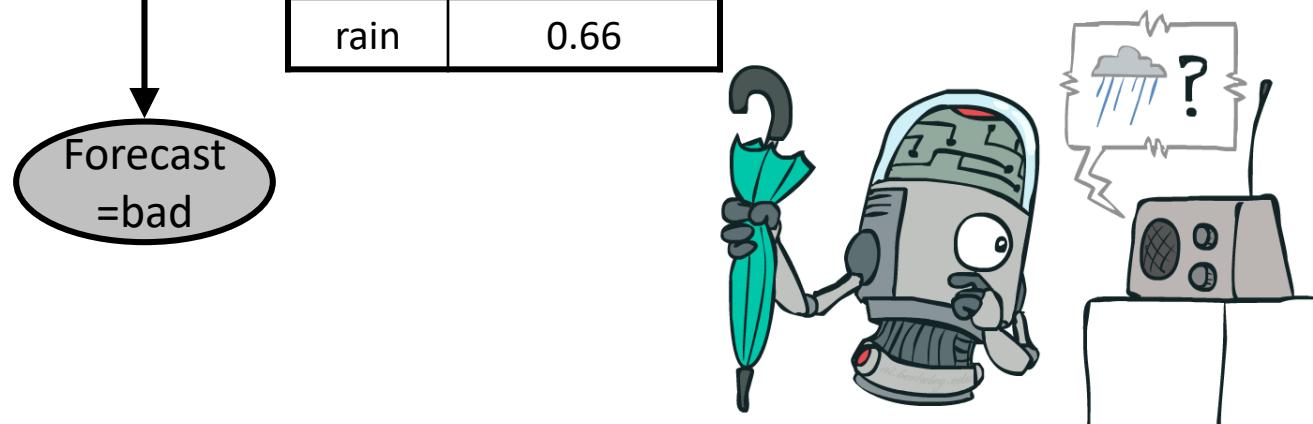
$$P(W|F) = \frac{P(W, F)}{\sum_w P(w, F)}$$

$$= \frac{P(F|W)P(W)}{\sum_w P(F|w)P(w)}$$



W	P(W F=bad)
sun	0.34
rain	0.66

A	W	U(A,W)
leave	sun	100
leave	rain	0
take	sun	20
take	rain	70



Maximum Expected Utility 2

Umbrella = leave

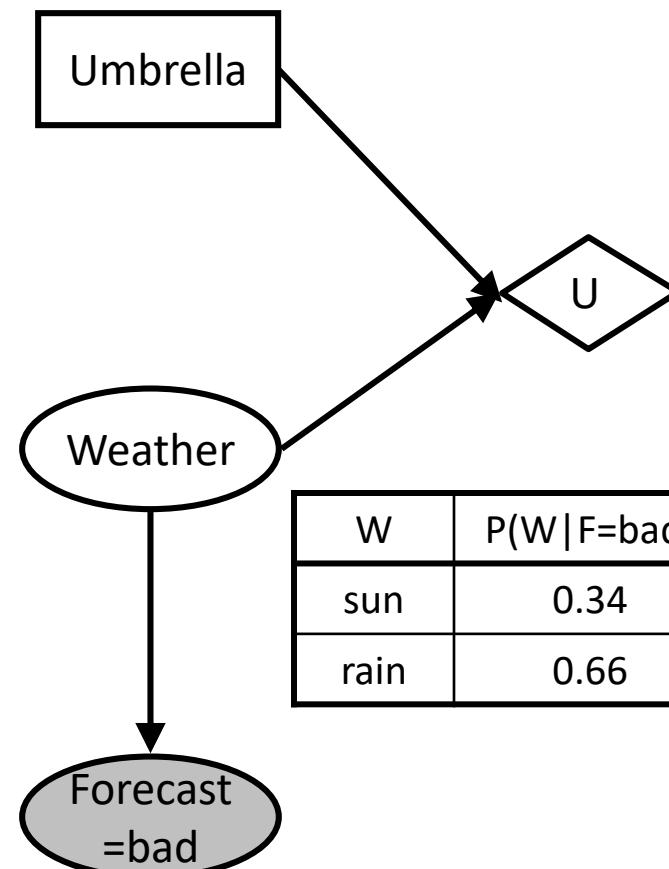
$$\begin{aligned} \text{EU}(\text{leave}|\text{bad}) &= \sum_w P(w|\text{bad})U(\text{leave}, w) \\ &= 0.34 \cdot 100 + 0.66 \cdot 0 = 34 \end{aligned}$$

Umbrella = take

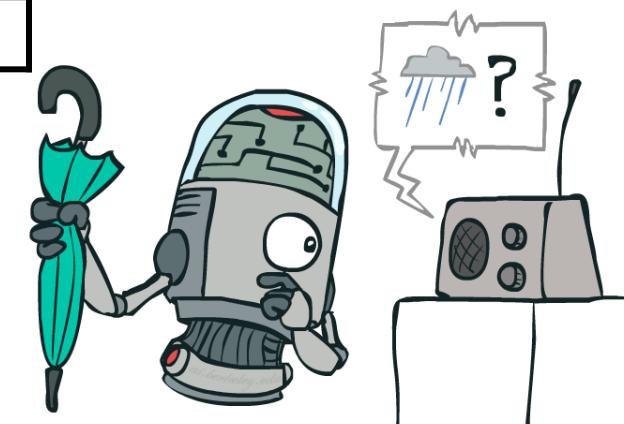
$$\begin{aligned} \text{EU}(\text{take}|\text{bad}) &= \sum_w P(w|\text{bad})U(\text{take}, w) \\ &= 0.34 \cdot 20 + 0.66 \cdot 70 = 53 \end{aligned}$$

Optimal decision = take

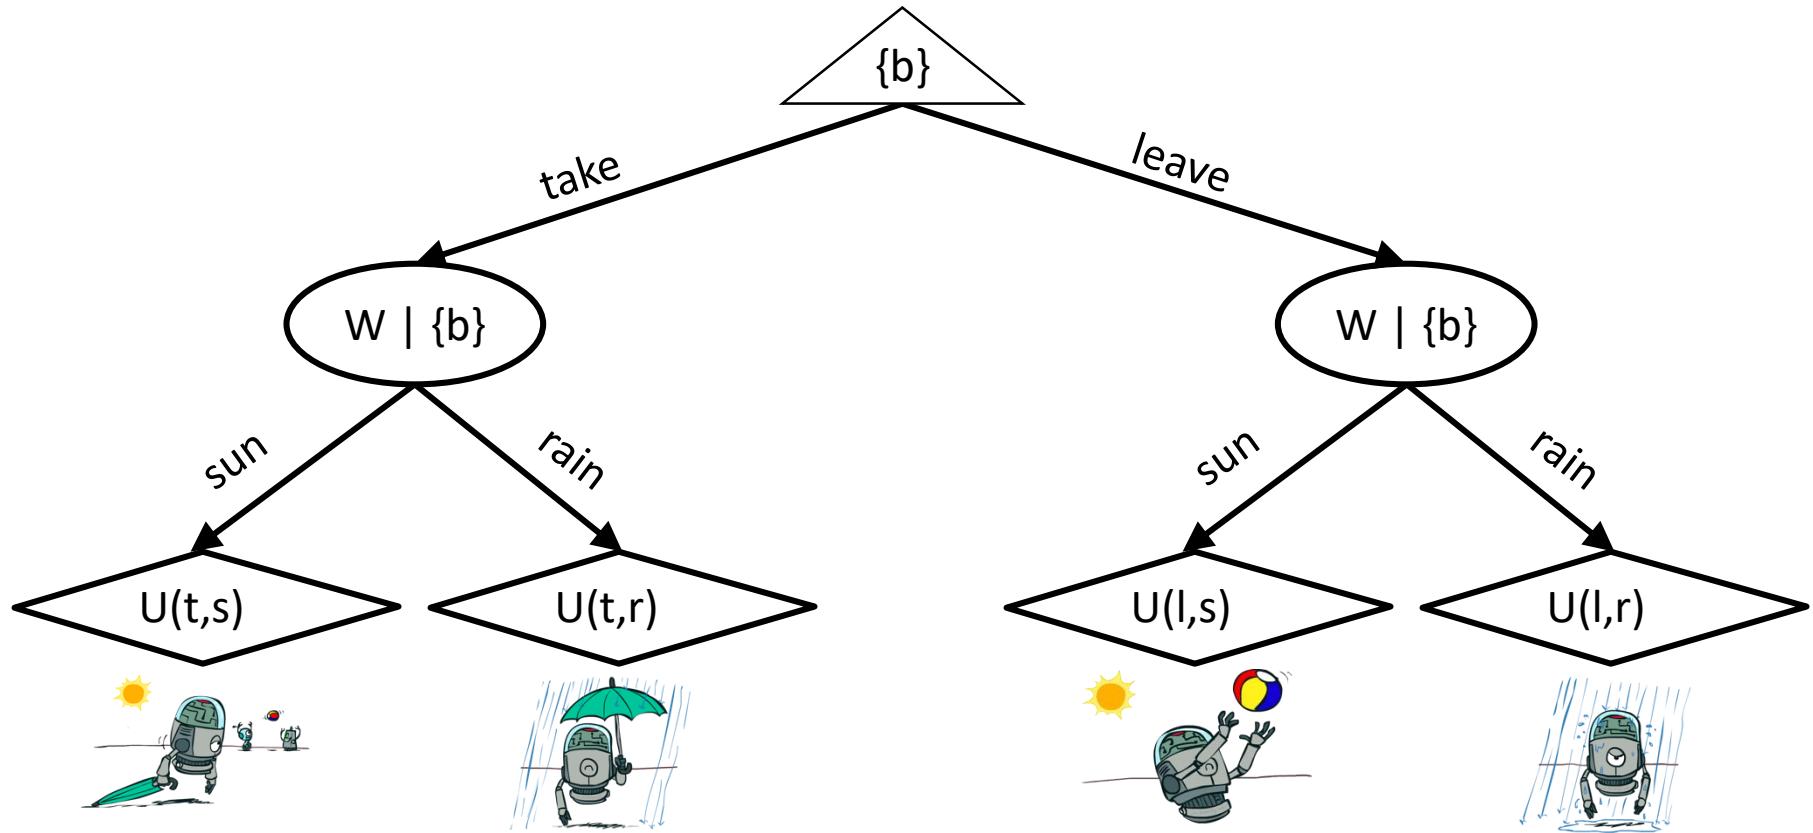
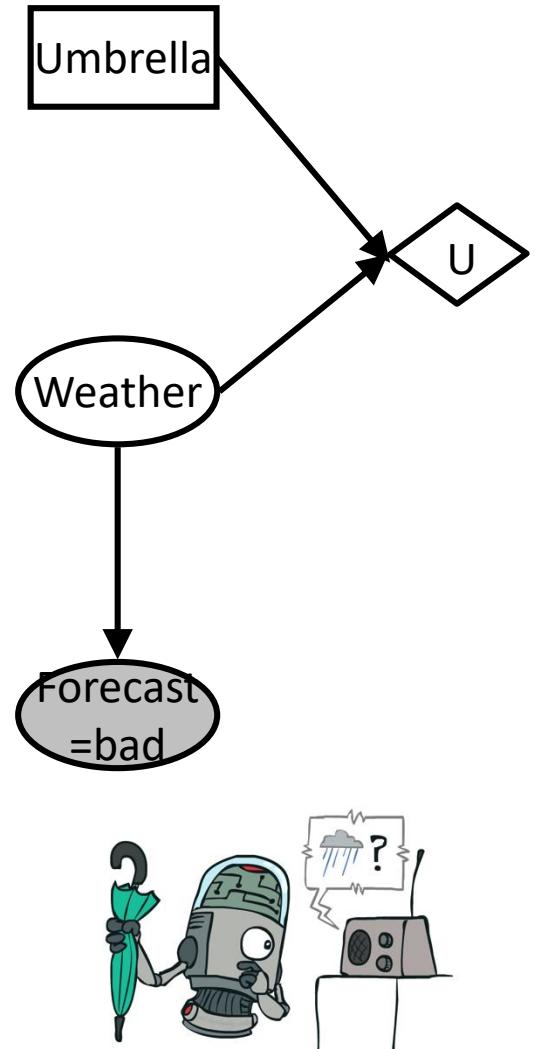
$$\text{MEU}(F = \text{bad}) = \max_a \text{EU}(a|\text{bad}) = 53$$



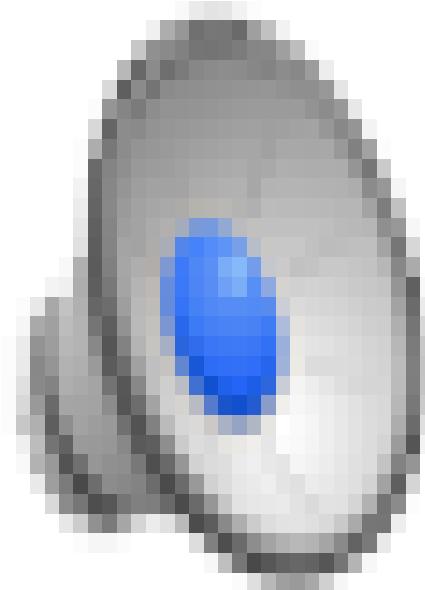
A	W	U(A,W)
leave	sun	100
leave	rain	0
take	sun	20
take	rain	70



Decisions as Outcome Trees

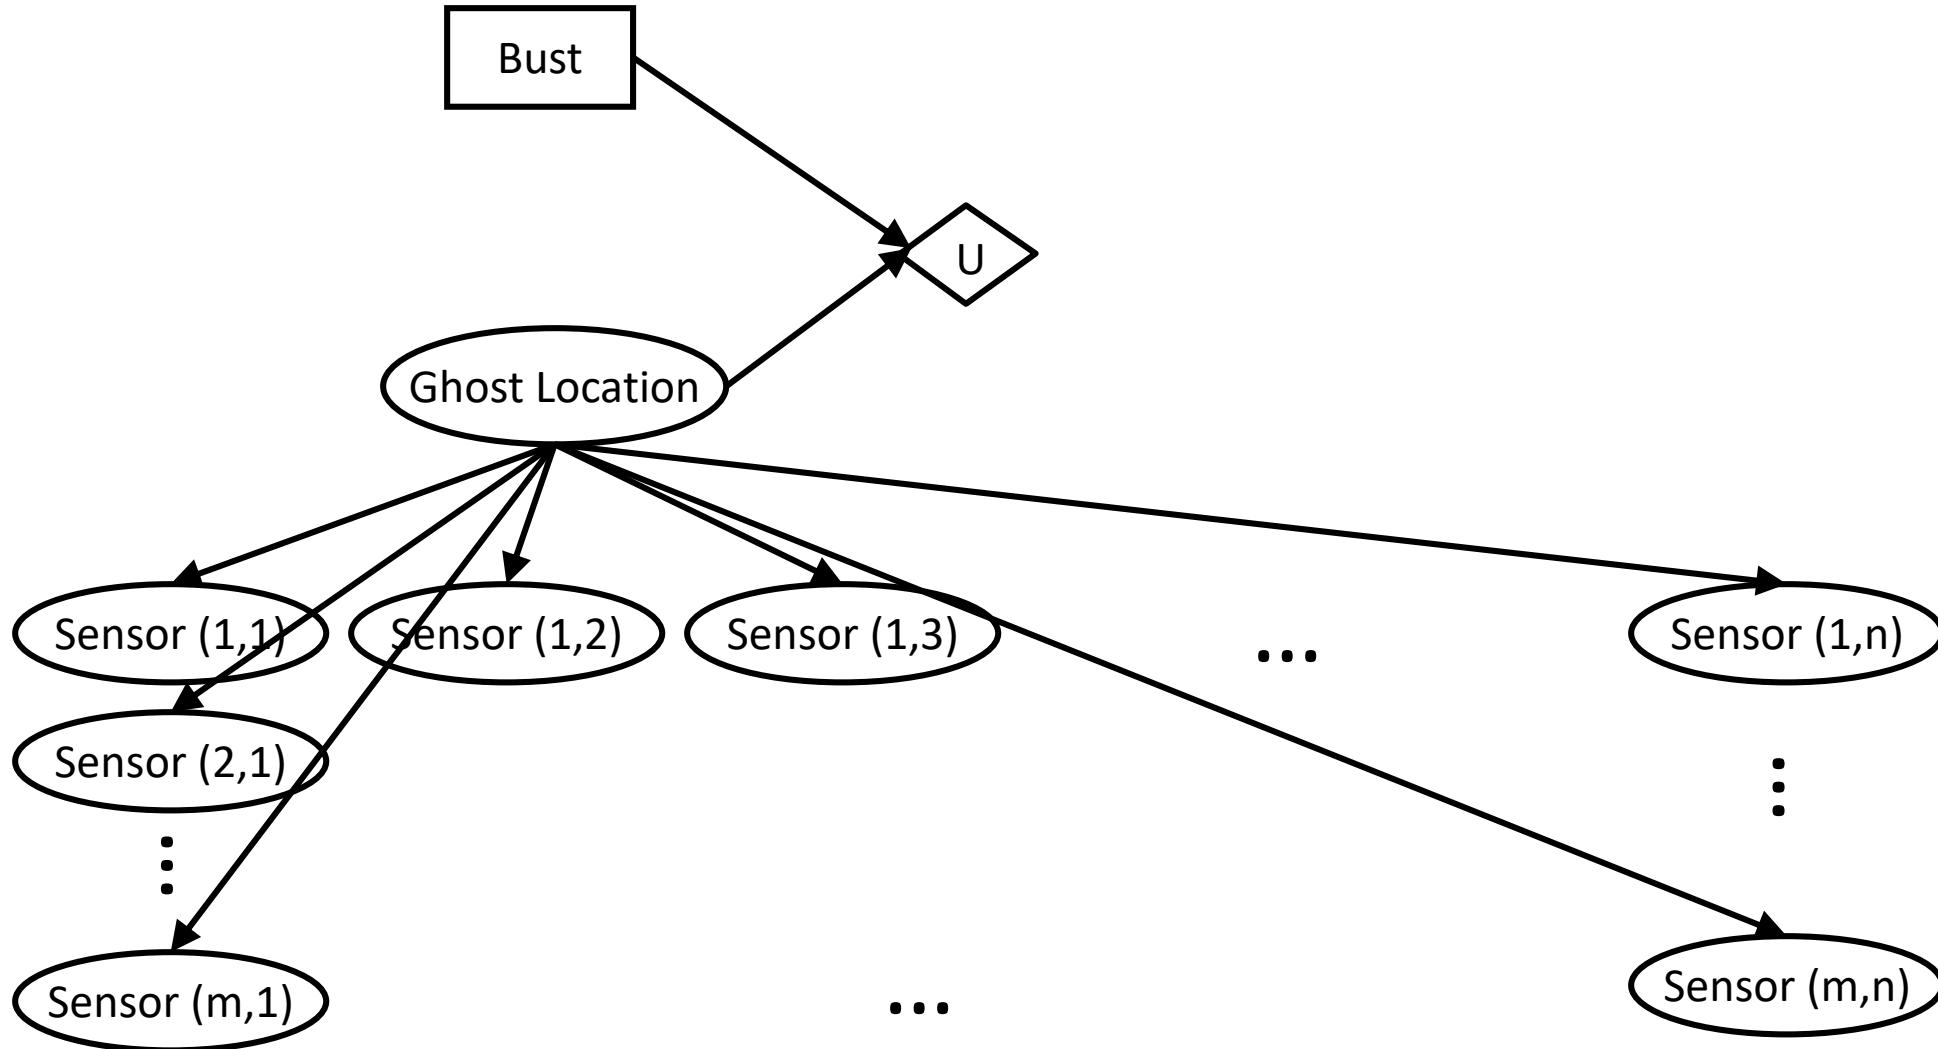


Video of Demo Ghostbusters with Probability

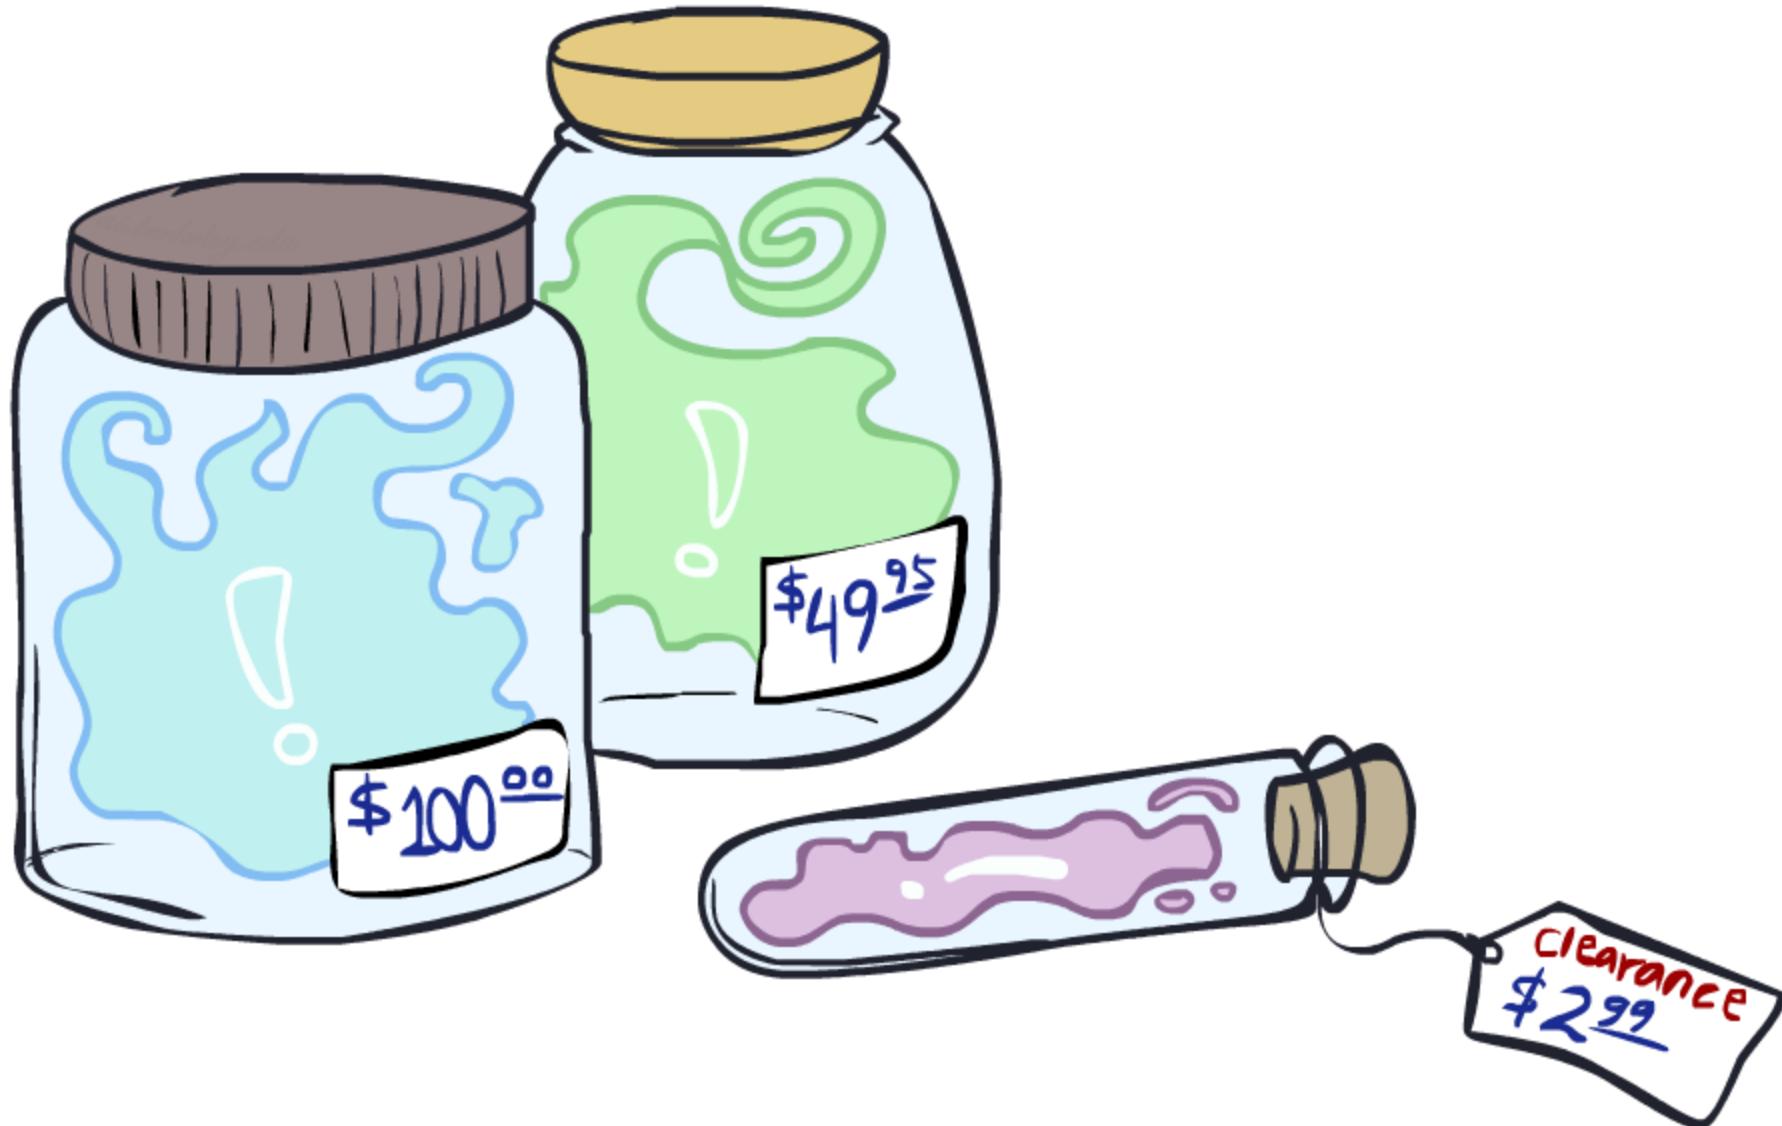


Ghostbusters Decision Network

Demo: Ghostbusters with probability

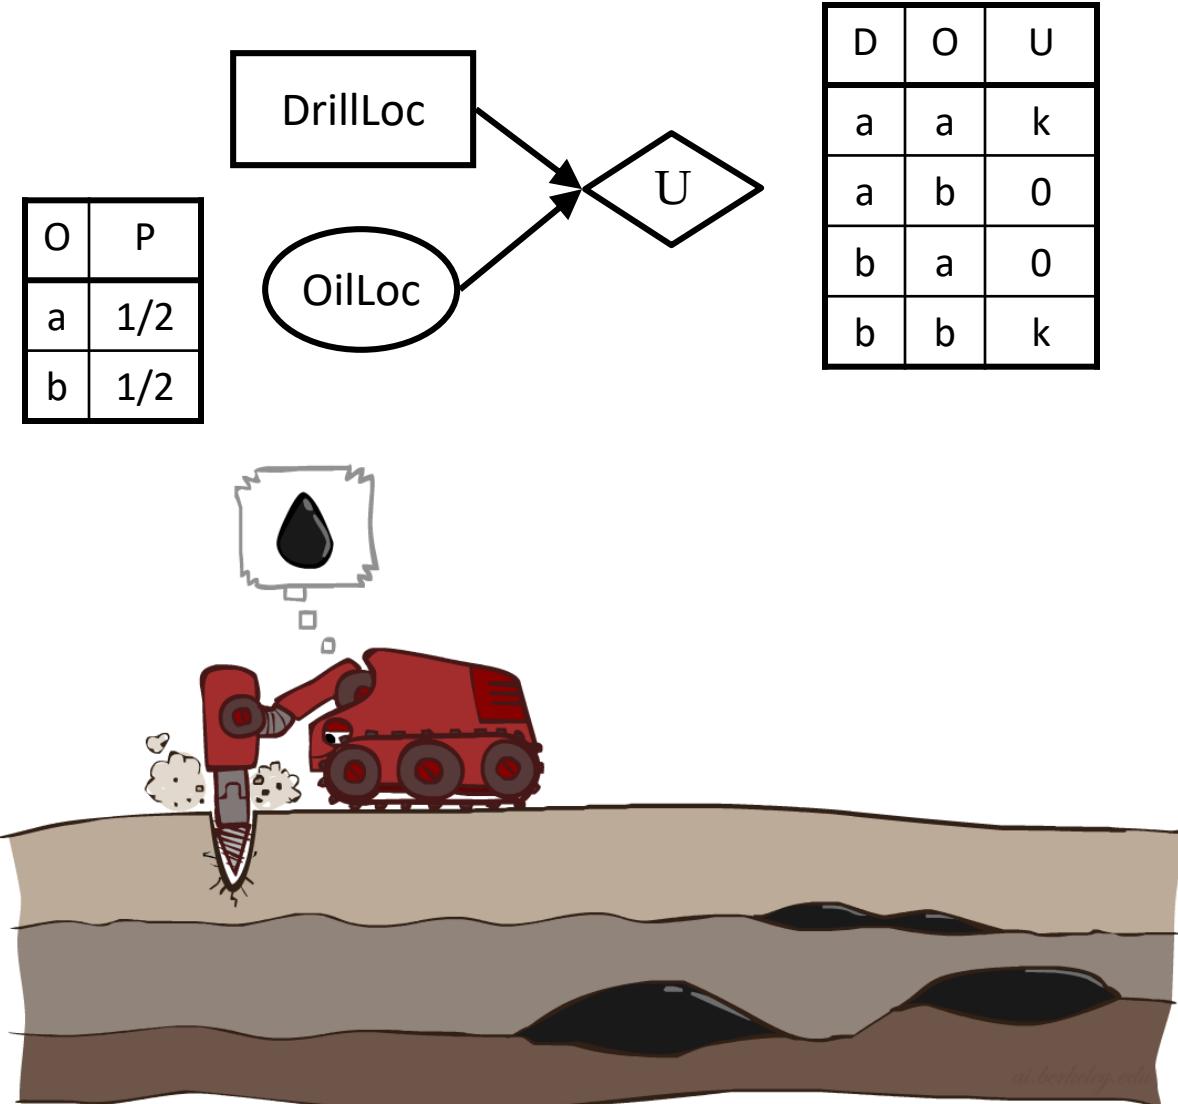


Value of Information



Value of Information

- Idea: compute value of acquiring evidence
 - Can be done directly from decision network
- Example: buying oil drilling rights
 - Two blocks A and B, exactly one has oil, worth k
 - You can drill in one location
 - Prior probabilities 0.5 each, & mutually exclusive
 - Drilling in either A or B has $EU = k/2$, $MEU = k/2$
- Question: what's the **value of information** of O ?
 - Value of knowing which of A or B has oil
 - Value is expected gain in MEU from new info
 - Survey may say "oil in a" or "oil in b," prob 0.5 each
 - If we know OilLoc, MEU is k (either way)
 - Gain in MEU from knowing OilLoc?
 - $VPI(OilLoc) = k/2$
 - Fair price of information: $k/2$



Value of Perfect Information

MEU with no evidence

$$\text{MEU}(\emptyset) = \max_a \text{EU}(a) = 70$$

MEU if forecast is bad

$$\text{MEU}(F = \text{bad}) = \max_a \text{EU}(a|\text{bad}) = 53$$

MEU if forecast is good

$$\text{MEU}(F = \text{good}) = \max_a \text{EU}(a|\text{good}) = 95$$

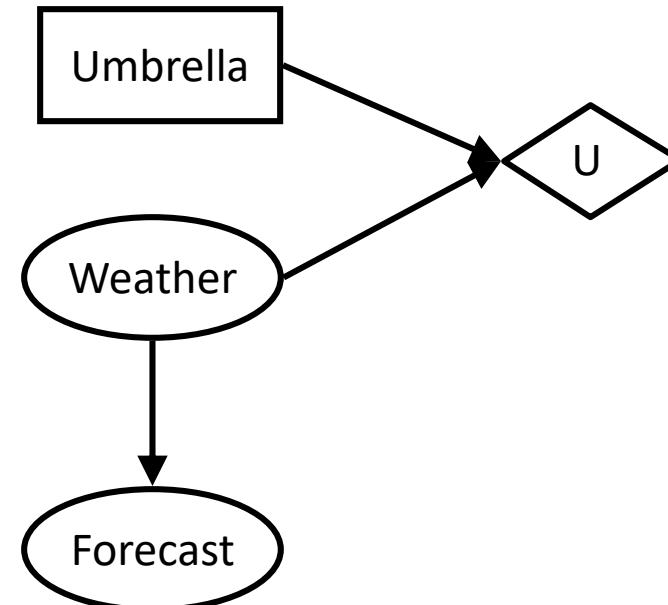
Forecast distribution

F	P(F)
good	0.59
bad	0.41

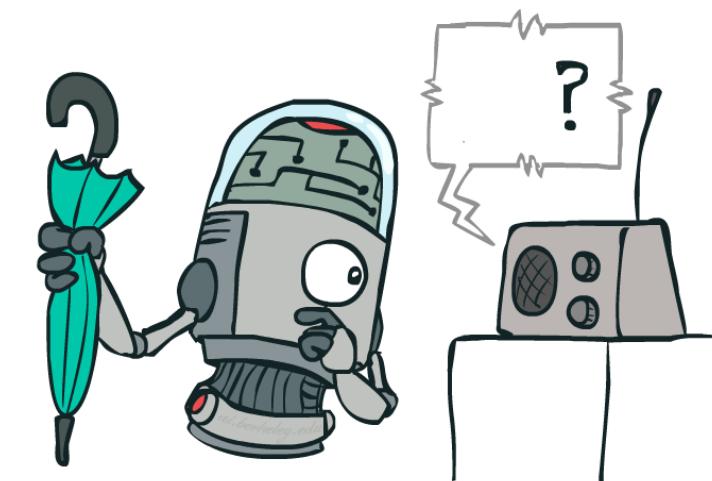


$$0.59 \cdot (95) + 0.41 \cdot (53) - 70 \\ 77.8 - 70 = 7.8$$

$$\text{VPI}(E'|e) = \left(\sum_{e'} P(e'|e) \text{MEU}(e, e') \right) - \text{MEU}(e)$$



A	W	U
leave	sun	100
leave	rain	0
take	sun	20
take	rain	70



Value of Information

- Assume we have evidence $E=e$. Value if we act now:

$$MEU(e) = \max_a \sum_s P(s|e) U(s, a)$$

- Assume we see that $E' = e'$. Value if we act then:

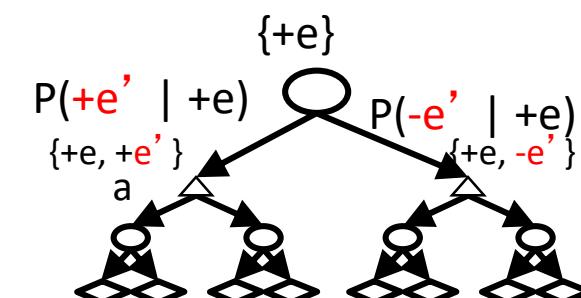
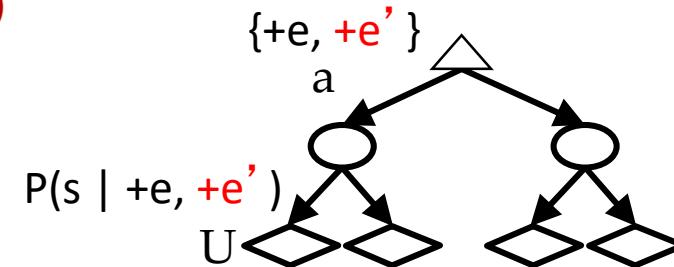
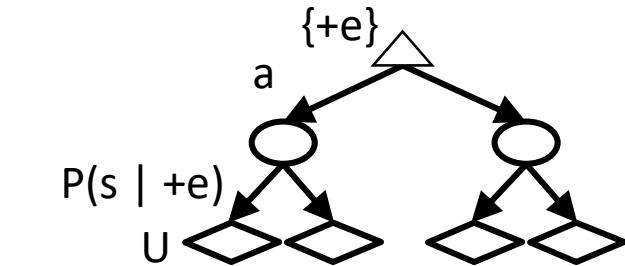
$$MEU(e, e') = \max_a \sum_s P(s|e, e') U(s, a)$$

- BUT E' is a random variable whose value is unknown, so we don't know what e' will be
- Expected value if E' is revealed and then we act:

$$MEU(e, E') = \sum_{e'} P(e'|e) MEU(e, e')$$

- Value of information: how much MEU goes up by revealing E' first then acting, over acting now:

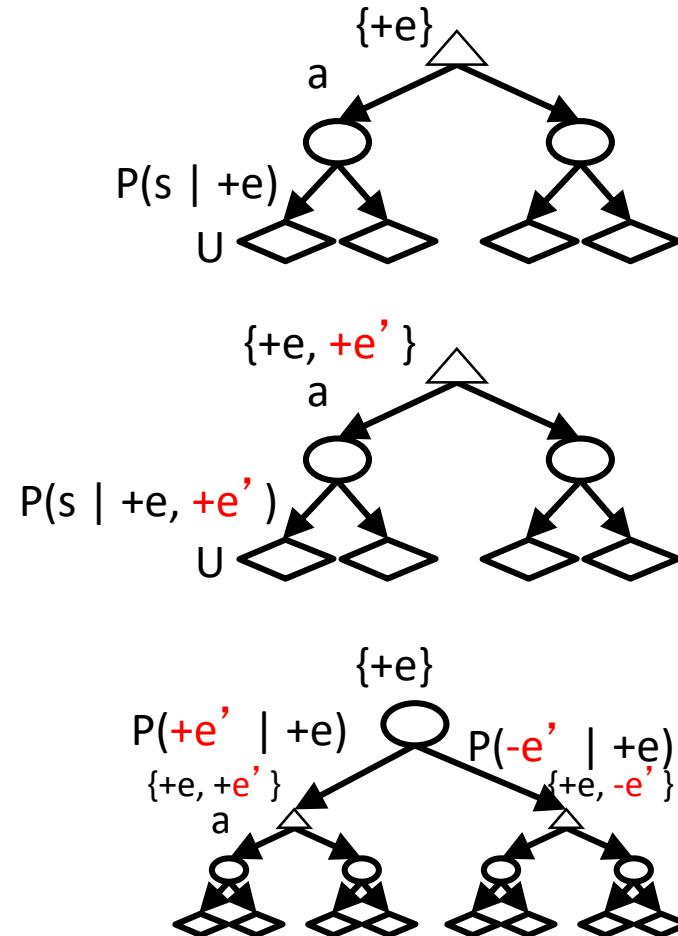
$$VPI(E'|e) = MEU(e, E') - MEU(e)$$



Value of Information 2

$$\begin{aligned} \text{MEU}(e, E') &= \sum_{e'} P(e'|e) \text{MEU}(e, e') \\ &= \sum_{e'} P(e'|e) \max_a \sum_s P(s|e, e') U(s, a) \end{aligned}$$

$$\begin{aligned} \text{MEU}(e) &= \max_a \sum_s P(s|e) U(s, a) \\ &= \max_a \sum_{e'} \sum_s P(s, e'|e) U(s, a) \\ &= \max_a \sum_{e'} P(e|e') \sum_s P(s|e, e') U(s, a) \end{aligned}$$



VPI Properties

- Nonnegative

$$\forall E', e : \text{VPI}(E'|e) \geq 0$$



- Nonadditive

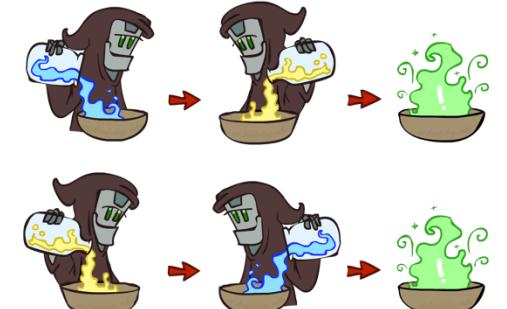
(think of observing E_i twice)

$$\text{VPI}(E_j, E_k|e) \neq \text{VPI}(E_j|e) + \text{VPI}(E_k|e)$$



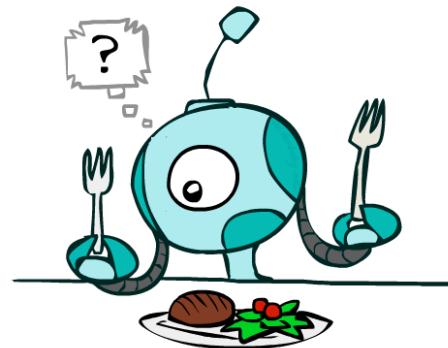
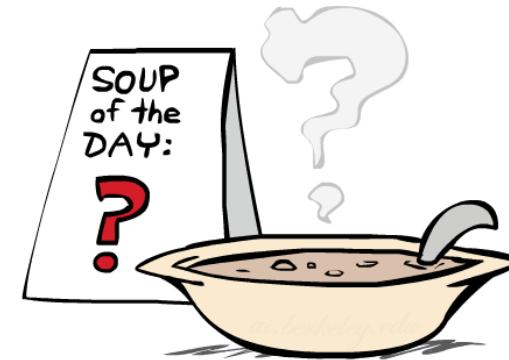
- Order-independent

$$\begin{aligned} \text{VPI}(E_j, E_k|e) &= \text{VPI}(E_j|e) + \text{VPI}(E_k|e, E_j) \\ &= \text{VPI}(E_k|e) + \text{VPI}(E_j|e, E_k) \end{aligned}$$



Quick VPI Questions

- The soup of the day is either clam chowder or split pea, but you wouldn't order either one. What's the value of knowing which it is?
- There are two kinds of plastic forks at a picnic. One kind is slightly sturdier. What's the value of knowing which?
- You're playing the lottery. The prize will be \$0 or \$100. You can play any number between 1 and 100 (chance of winning is 1%). What is the value of knowing the winning number?



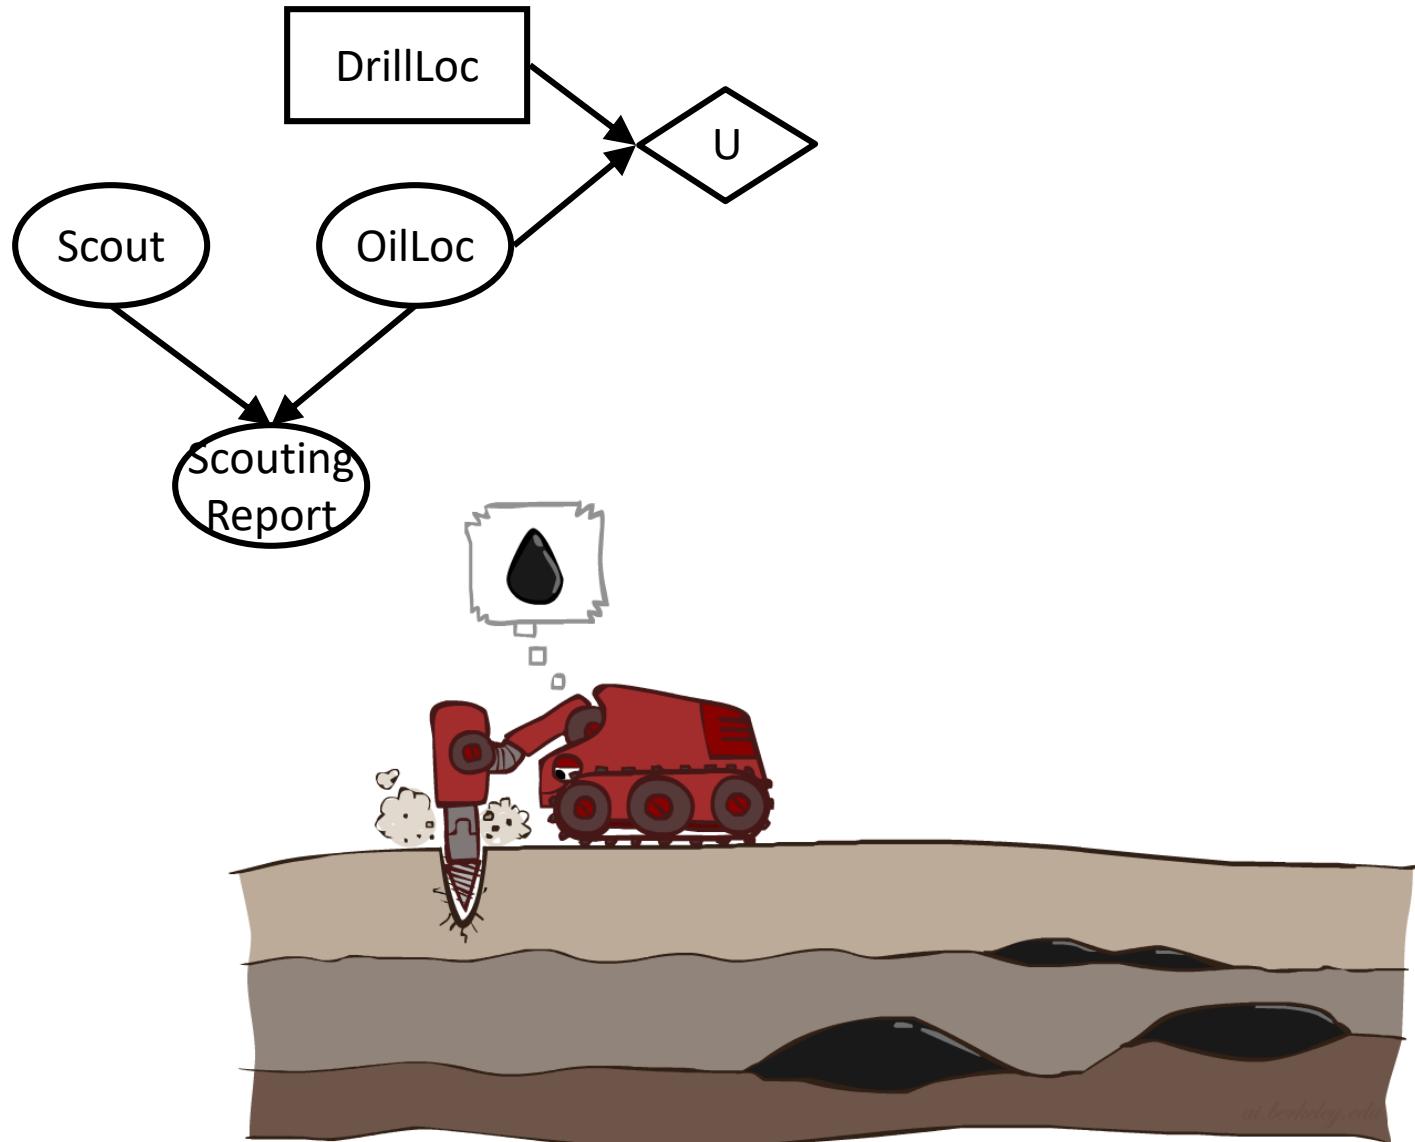
Value of Imperfect Information?

- No such thing
- Information corresponds to the observation of a node in the decision network
- If data is “noisy” that just means we don’t observe the original variable, but another variable which is a noisy version of the original one

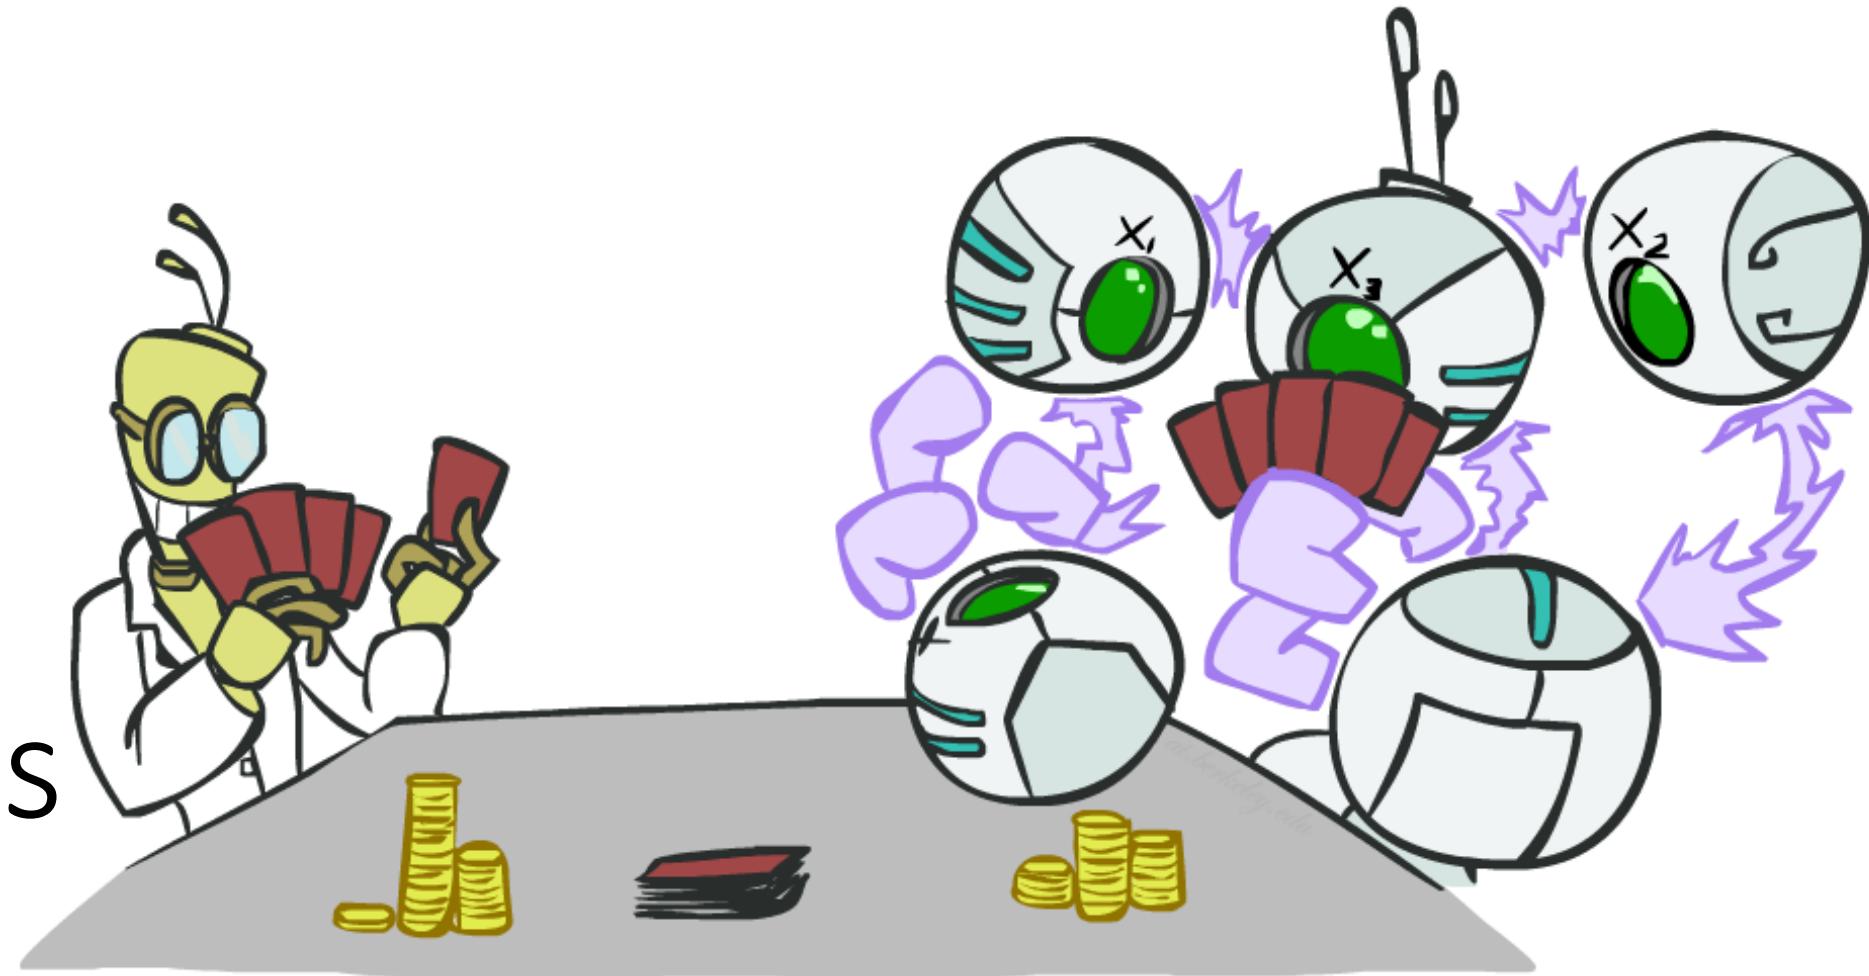


VPI Question

- VPI(OilLoc) ?
- VPI(ScoutingReport) ?
- VPI(Scout) ?
- VPI(Scout | ScoutingReport) ?
- Generally:
If $\text{Parents}(U) \perp\!\!\!\perp Z \mid \text{CurrentEvidence}$
Then $\text{VPI}(Z \mid \text{CurrentEvidence}) = 0$

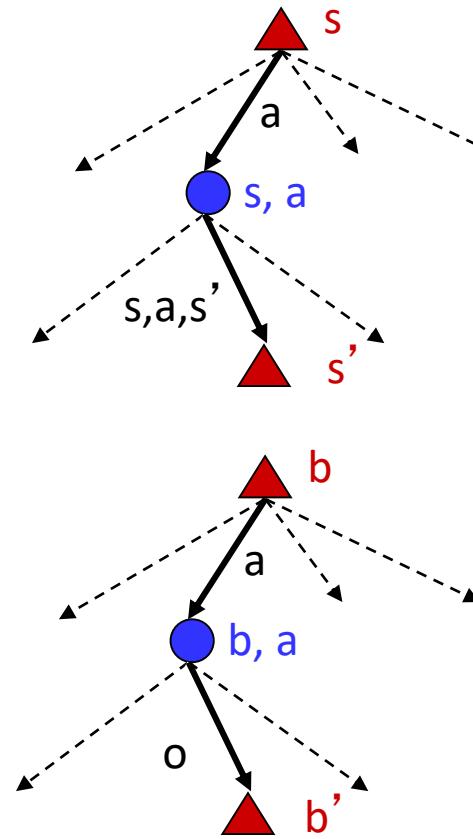


POMDPs



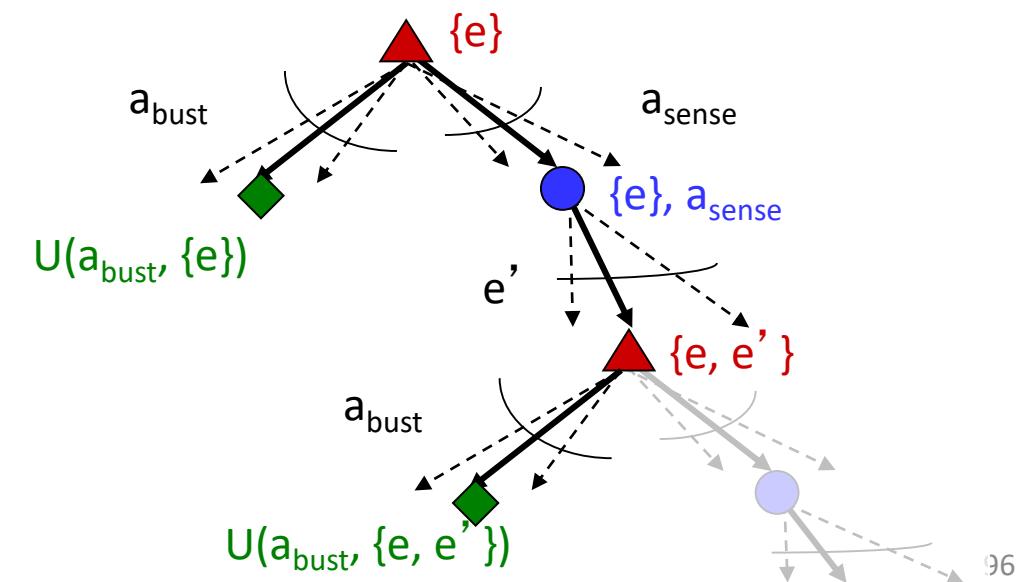
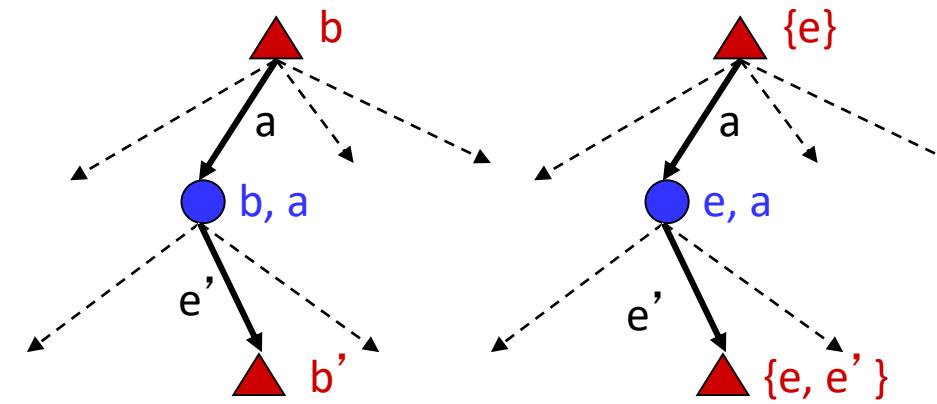
POMDPs

- MDPs have:
 - States S
 - Actions A
 - Transition function $P(s' | s, a)$ (or $T(s, a, s')$)
 - Rewards $R(s, a, s')$
- POMDPs add:
 - Observations O
 - Observation function $P(o | s)$ (or $O(s, o)$)
- POMDPs are MDPs over belief states b (distributions over S)

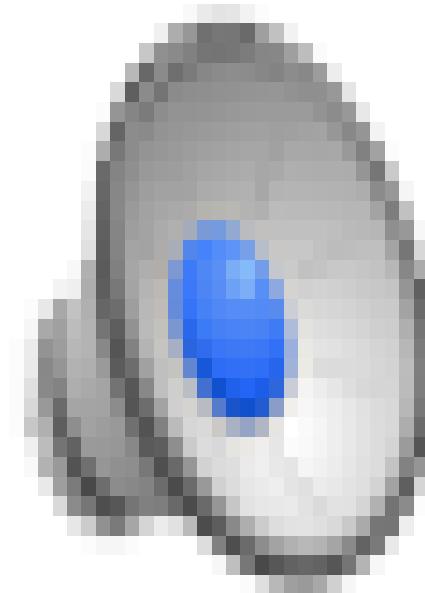


Example: Ghostbusters

- In (static) Ghostbusters:
 - Belief state determined by evidence to date $\{e\}$
 - Tree really over evidence sets
 - Probabilistic reasoning needed to predict new evidence given past evidence
- Solving POMDPs
 - One way: use truncated expectimax to compute approximate value of actions
 - What if you only considered busting or one sense followed by a bust?
 - You get a VPI-based agent!

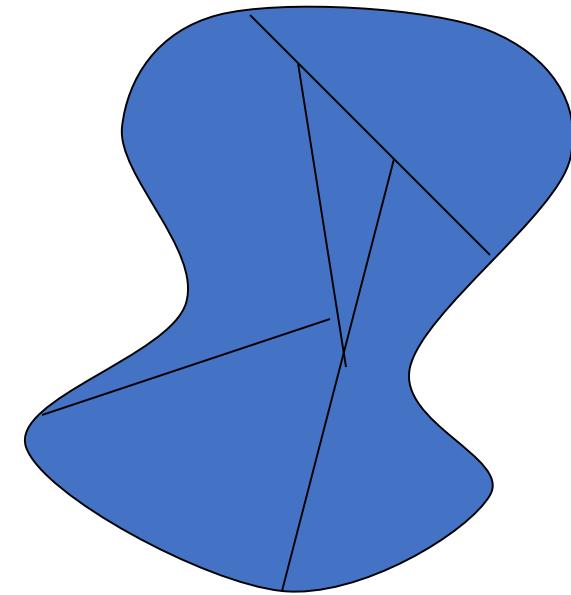


Video of Demo Ghostbusters with VPI



More Generally*

- General solutions map belief functions to actions
 - Can divide regions of belief space (set of belief functions) into policy regions (gets complex quickly)
 - Can build approximate policies using discretization methods
 - Can factor belief functions in various ways
- Overall, POMDPs are very hard
- Most real problems are POMDPs, but we can rarely solve them in general!



Summary

- Bayes rule
- Inference
- Variable Elimination
- Sampling
- Decision Networks

Shuai Li

<https://shuaili8.github.io>

Questions?